



KubeCon

CloudNativeCon

North America 2018

Airflow + Kubernetes

Daniel Imberman, Bloomberg

Barni Seetharaman, Google

Bio: Daniel + Barni

Raw Data → **Actionable Data**



KubeCon



CloudNativeCon

North America 2018

Pipelines are hard

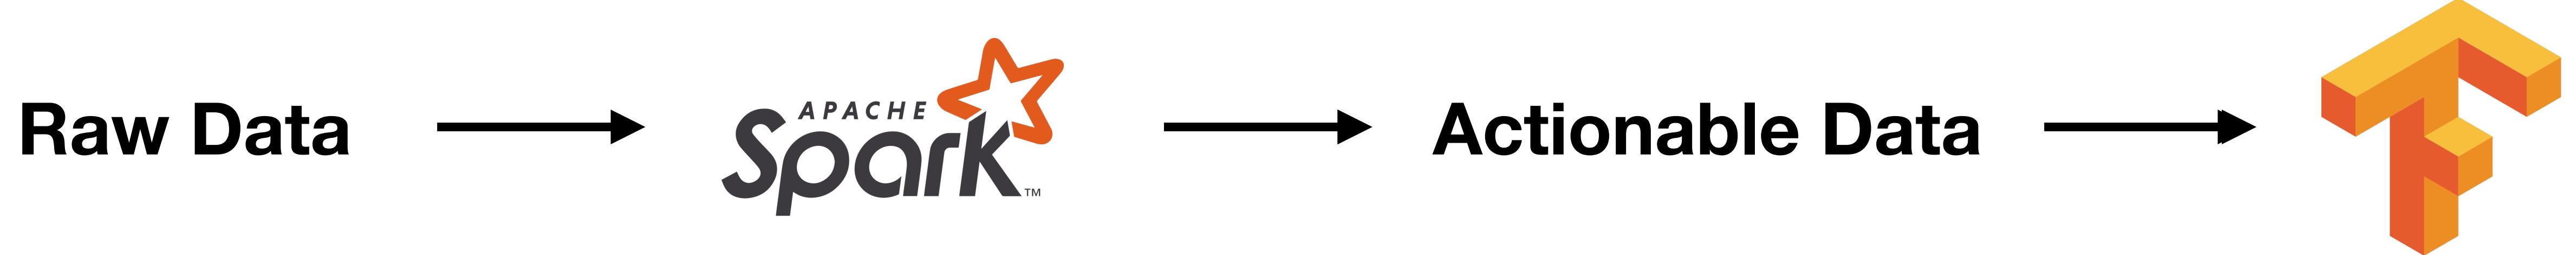
Pipelines are hard

Raw Data → **Actionable Data**

Pipelines are hard



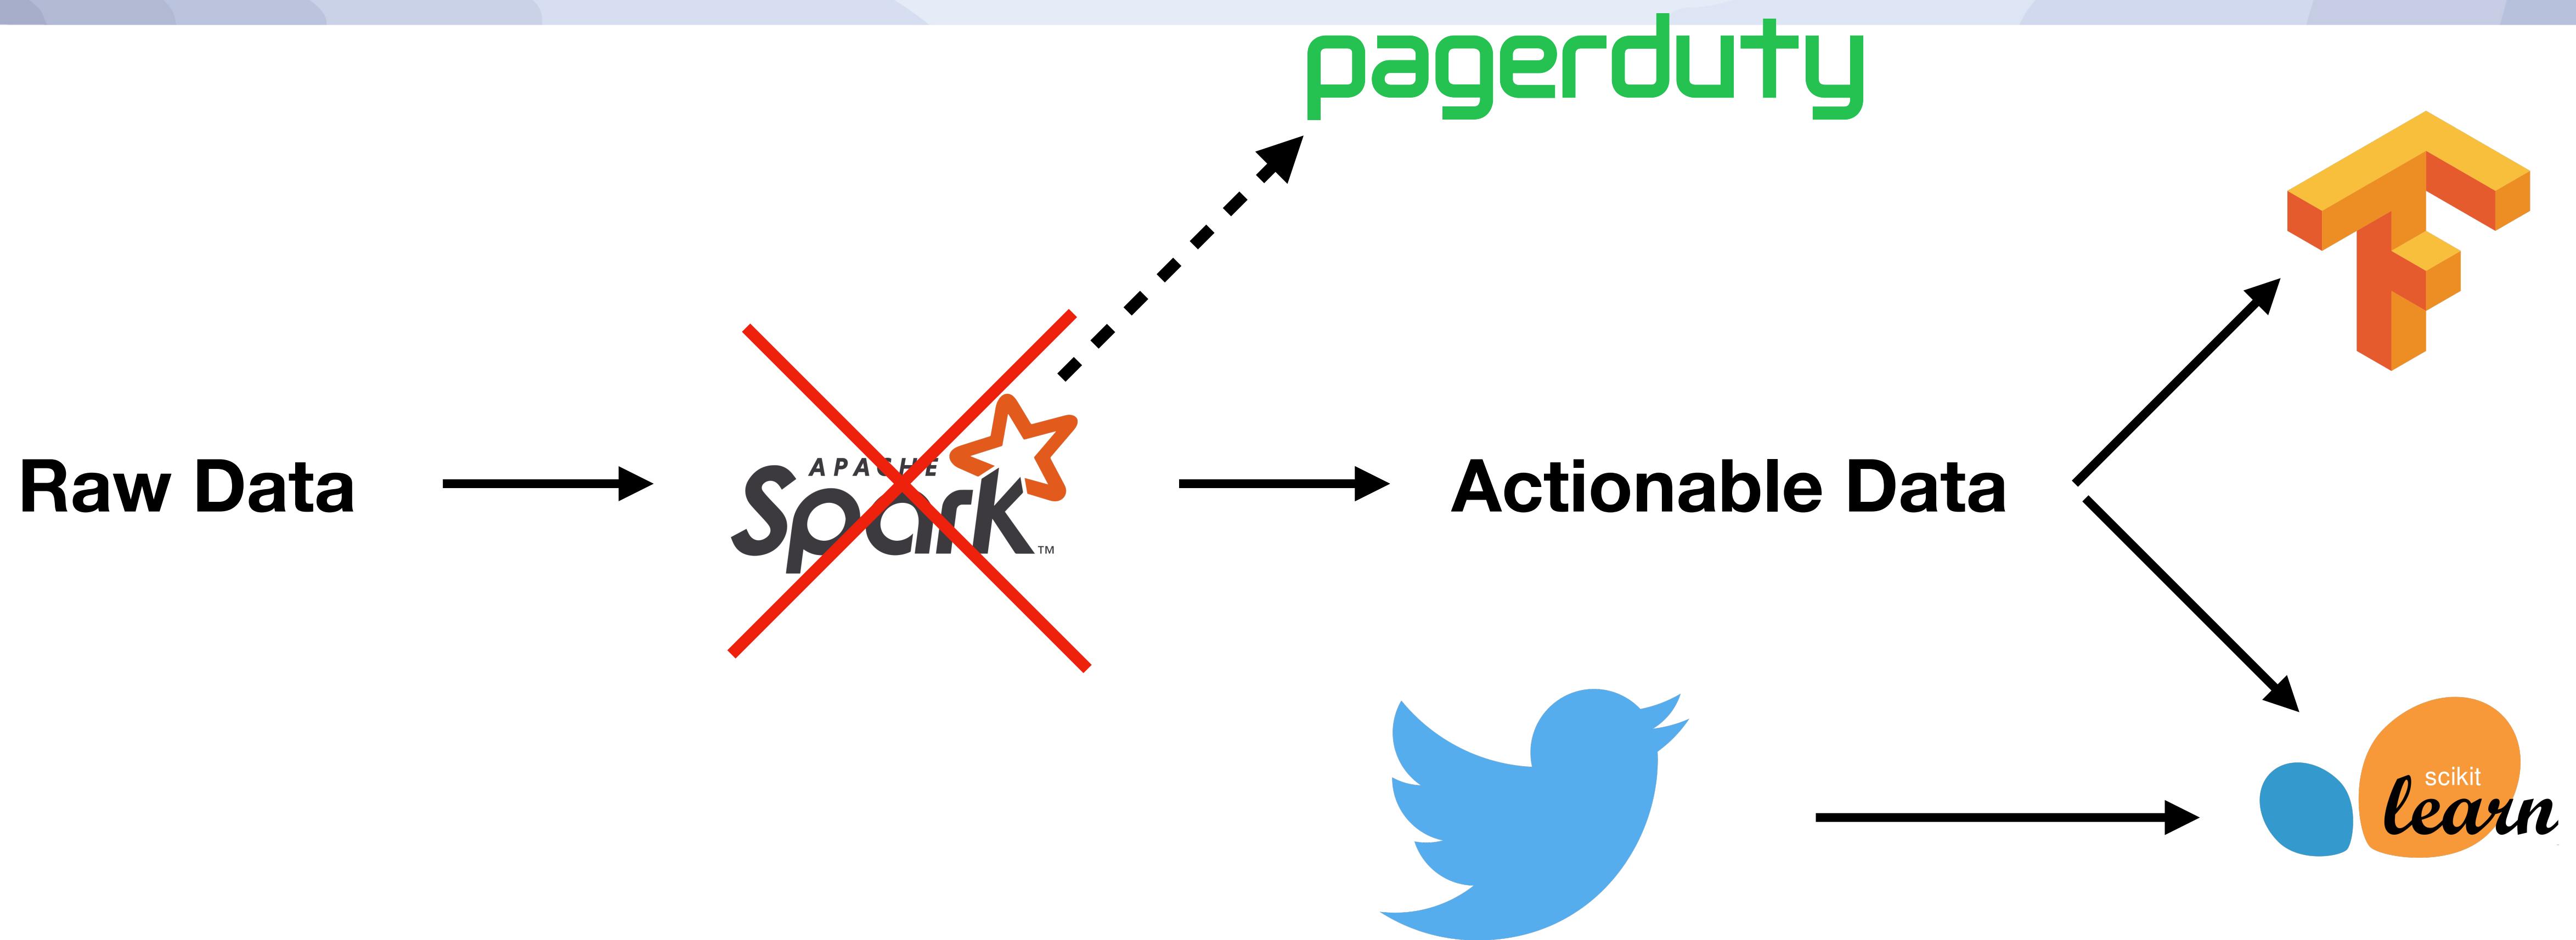
Pipelines are hard



Pipelines are hard



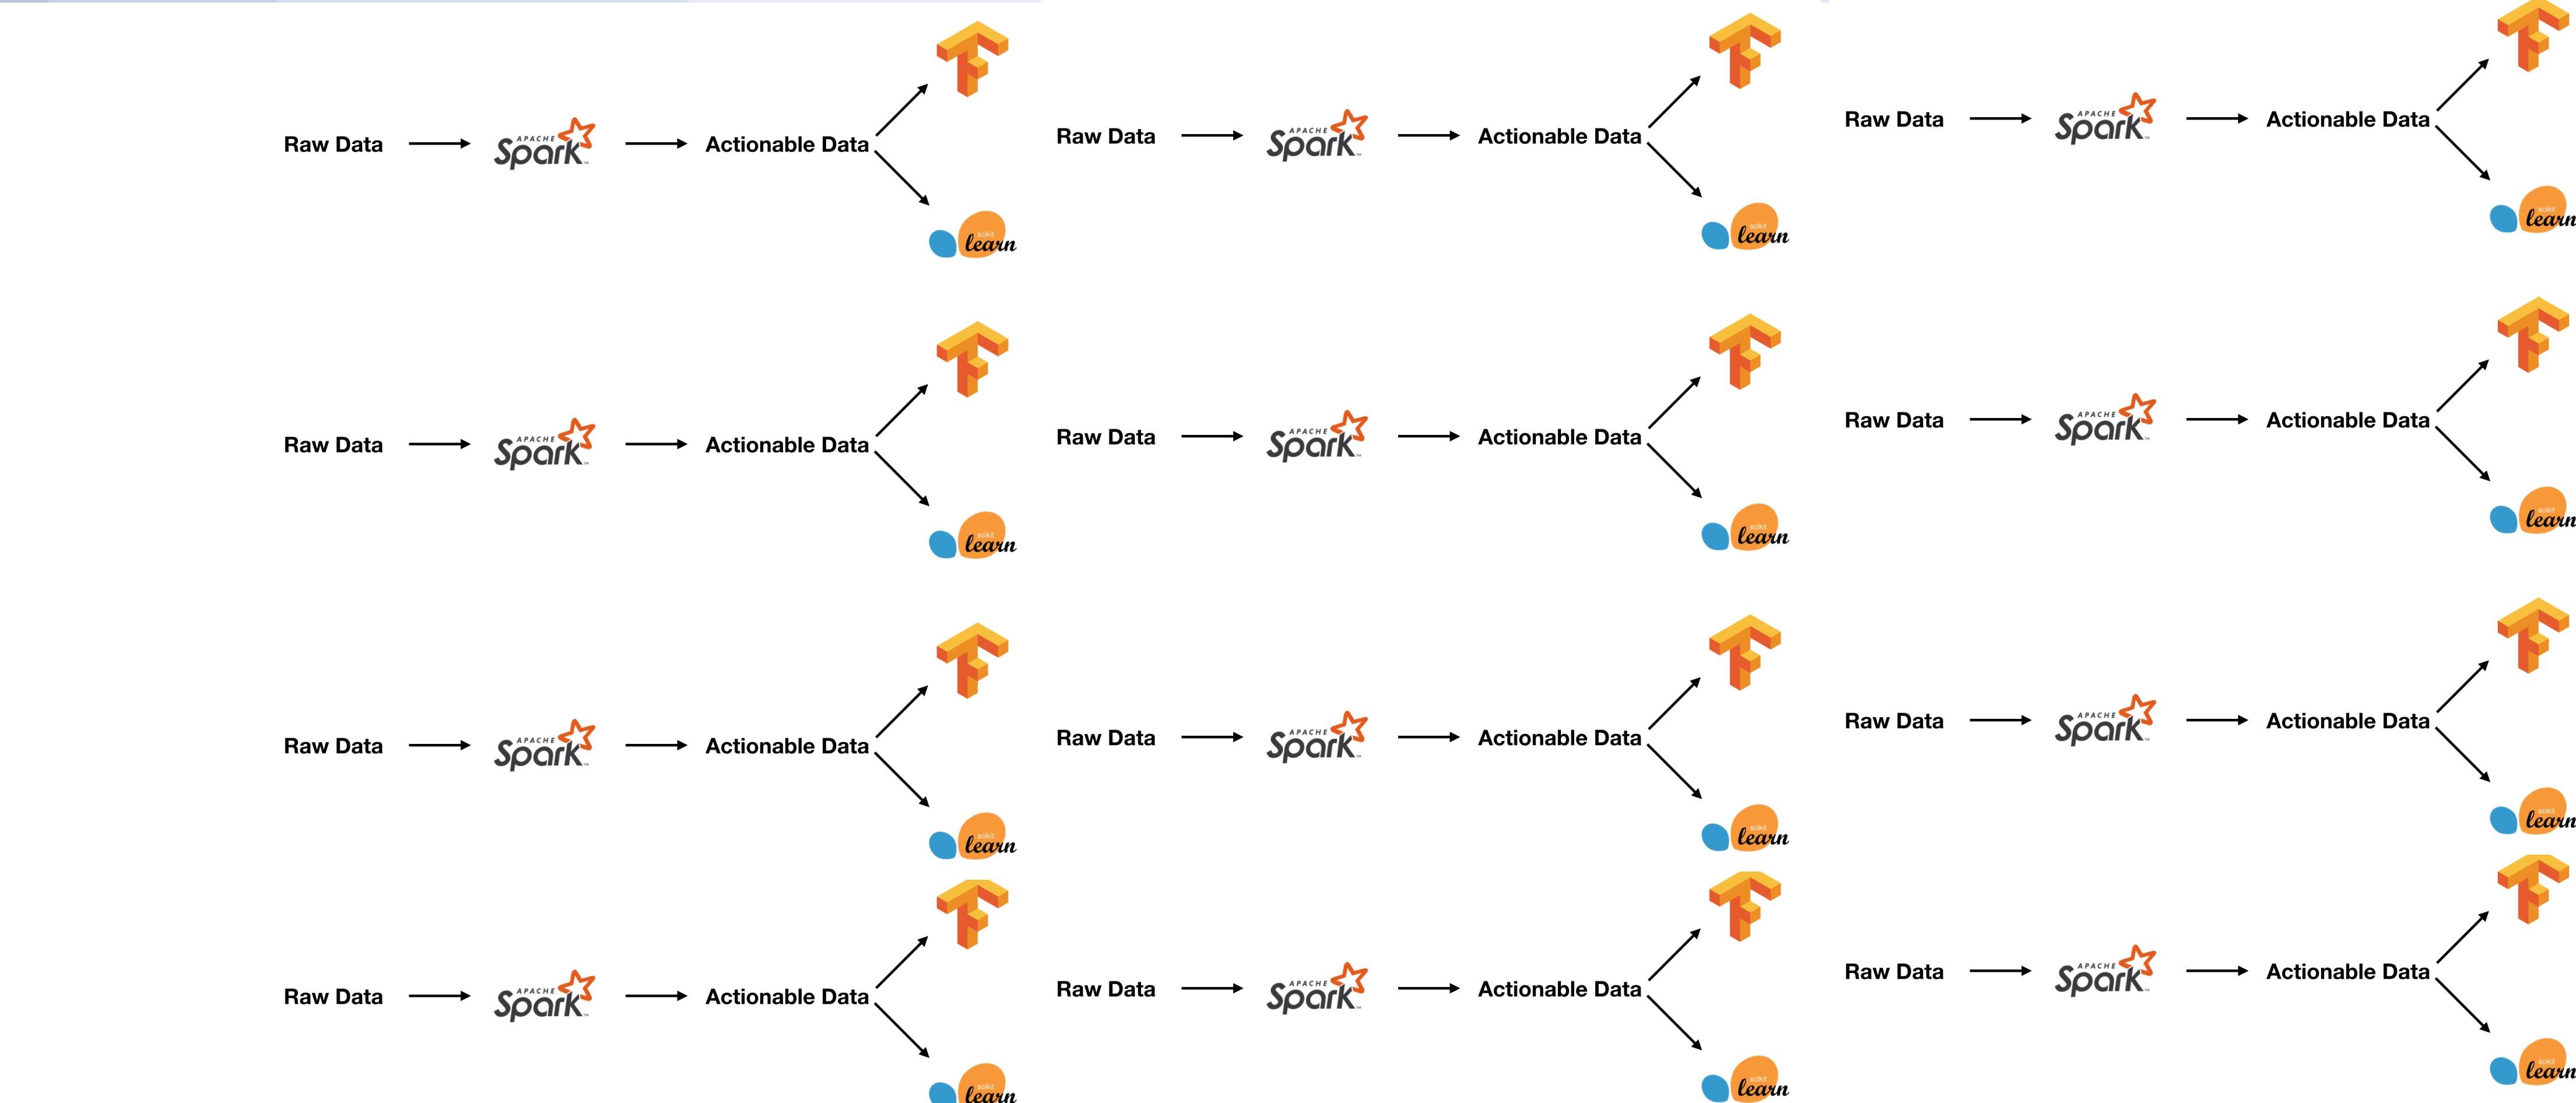
Pipelines are hard



Lots of pipelines are really hard



North America 2018



Enter Apache Airflow



Apache Airflow

- Workflow Scheduler developed @ Airbnb
- Converts Python code into DAGs
- Has large number of operators/hooks (HDFS, Spark, Bash, Hive, etc...)



Apache Airflow

Airflow - DAGs Dilbert-2005091 +
http://localhost:8081/admin/airflow/tree?dag_id=batch_postgresql_v1

Airflow DAGs Data Profiling Browse Admin Docs About 11:27 UTC ⚡

DAGs

Show entries Search:

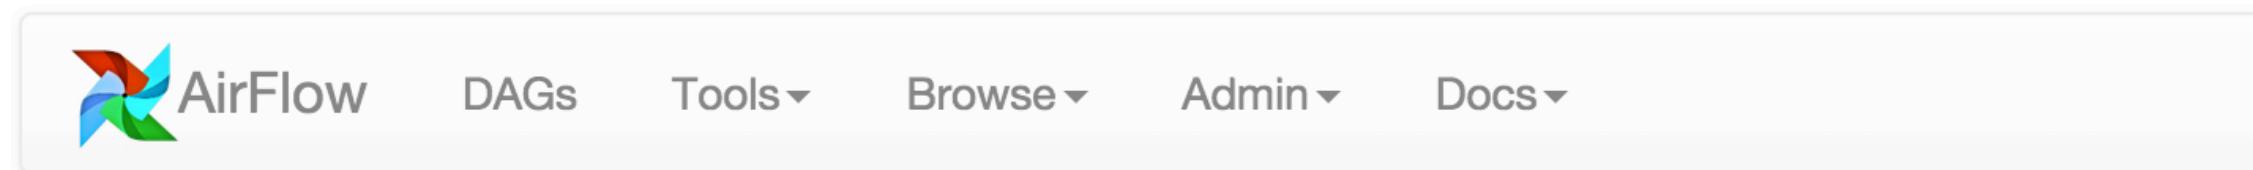
	i	DAG	Schedule	Owner	Recent Tasks i	Last Run i	DAG Runs i	Links
<input checked="" type="checkbox"/>	On	batch_mysql_v2	0 1 * * *	airflow	2 ○ ○ ○ ○ ○	2017-08-08 01:00 i	28 ○ ○ ○	@ # * !! ✖ ⚡ ☰
<input checked="" type="checkbox"/>	On	batch_postgresql_v1	30 0 * * *	airflow	○ ○ 1 ○ ○ ○	2017-08-08 00:30 i	18 ○ 1 ○	@ # * !! ✖ ⚡ ☰
<input checked="" type="checkbox"/>	On	batch_sqlserver_v2	30 1 * * *	airflow	2 ○ ○ ○ ○ ○	2017-08-08 01:30 i	28 ○ ○ ○	@ # * !! ✖ ⚡ ☰
<input checked="" type="checkbox"/>	On	sales_ftp_v1	0 1 * * *	airflow	2 ○ ○ ○ ○ ○	2017-08-08 01:00 i	28 ○ ○ ○	@ # * !! ✖ ⚡ ☰

Showing 1 to 4 of 4 entries Previous **1** Next

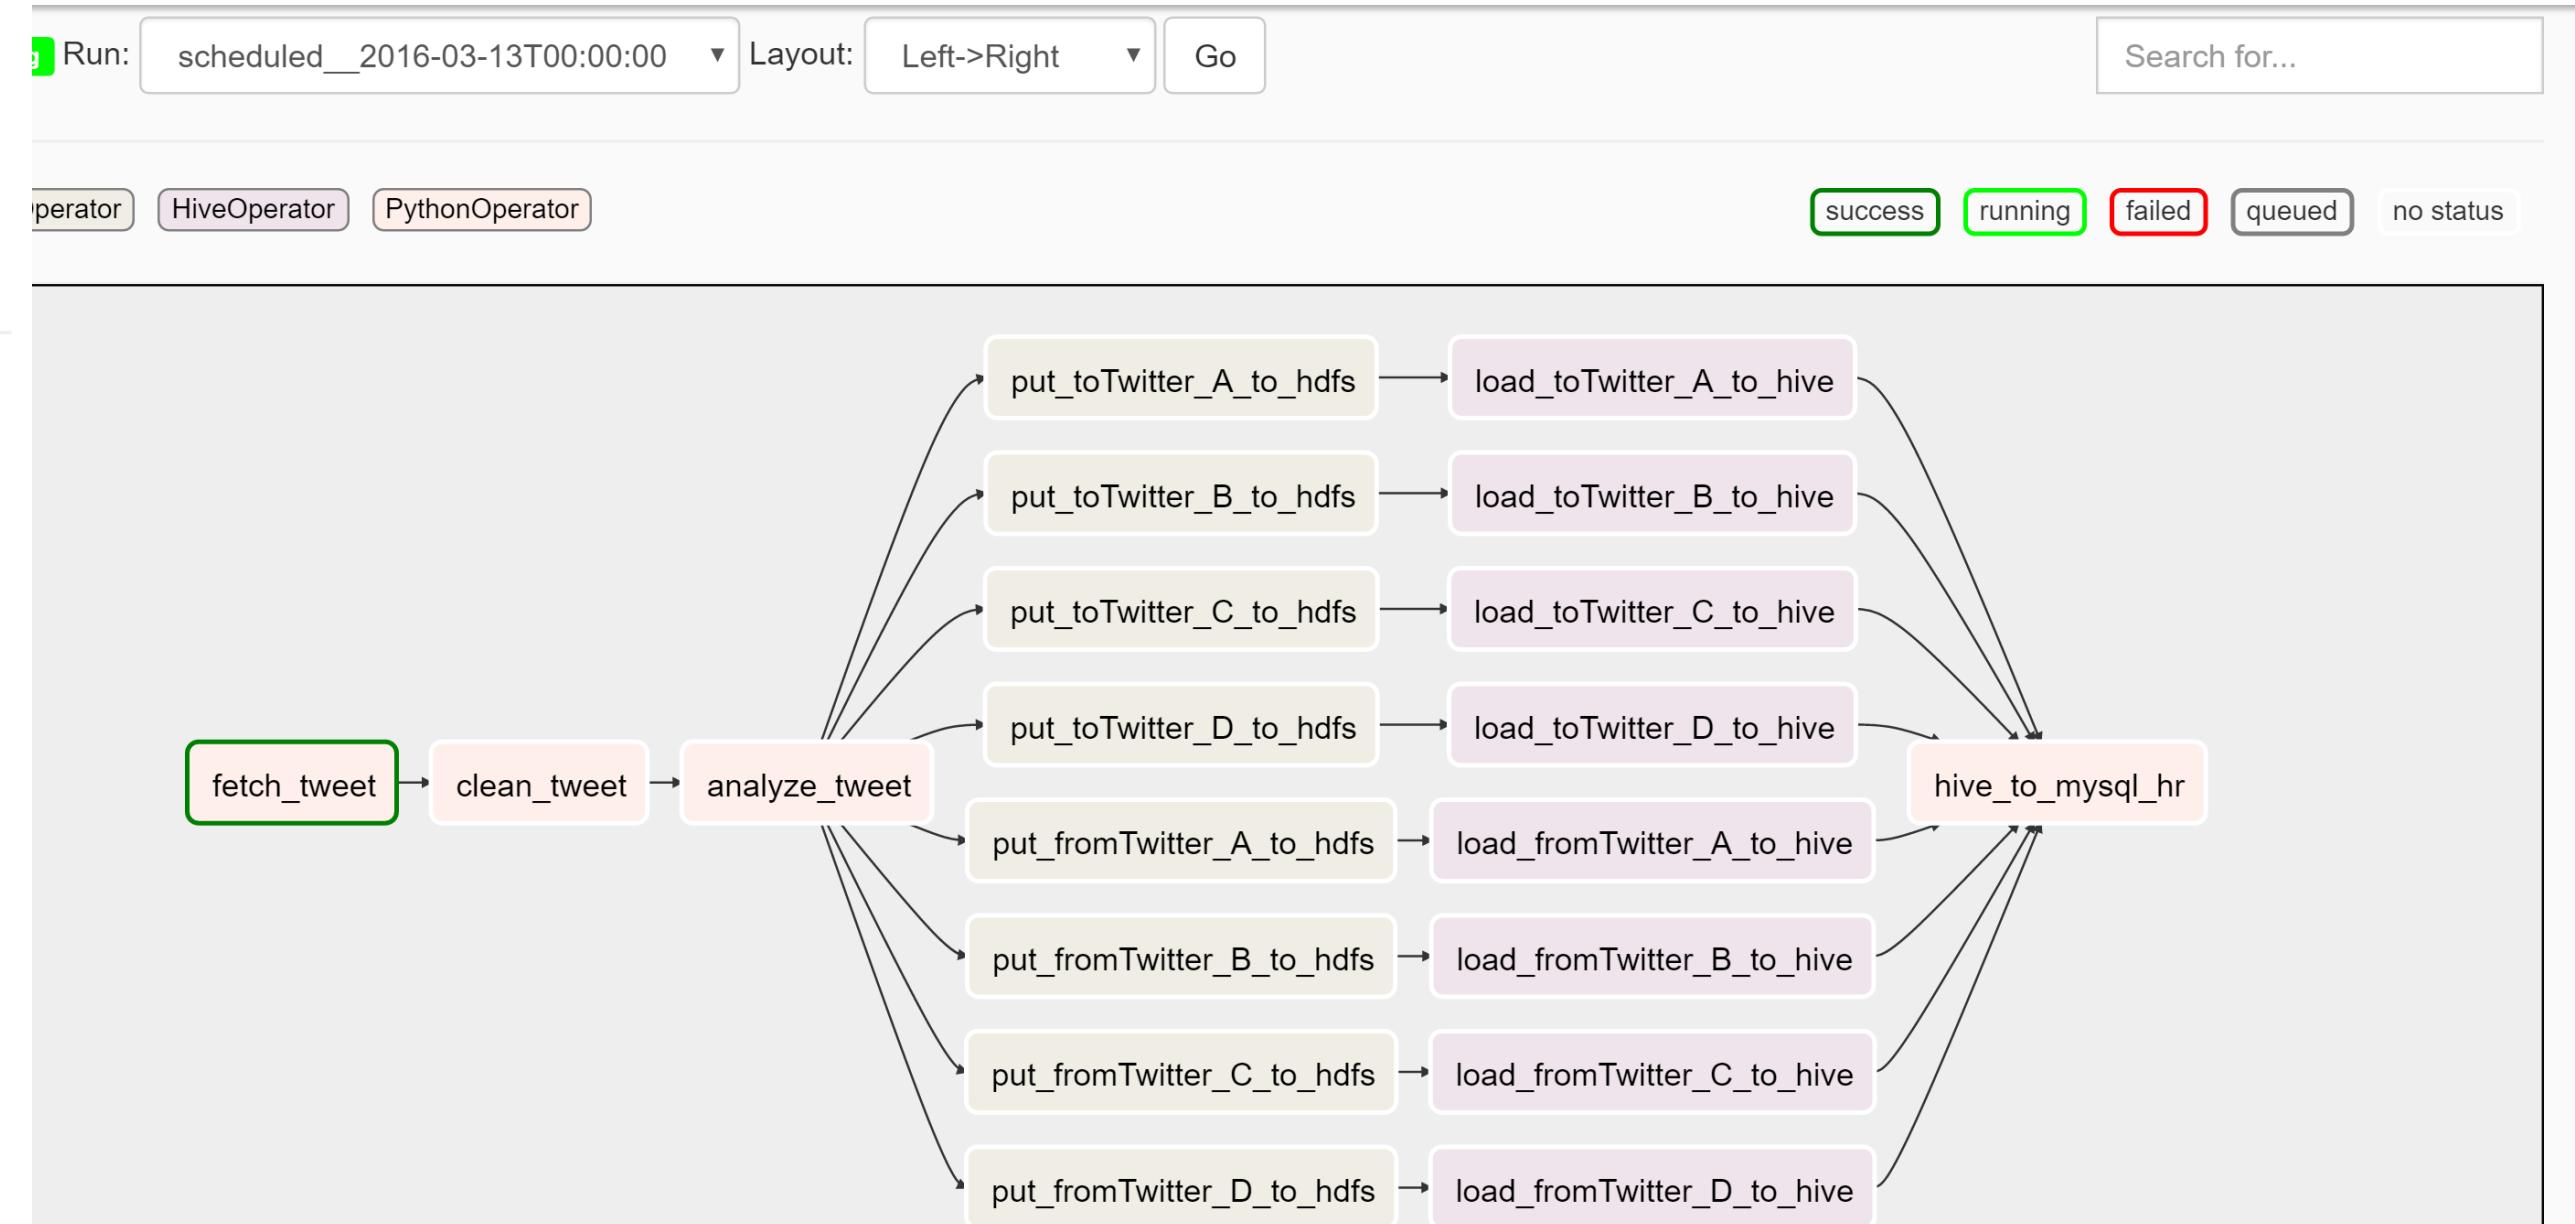
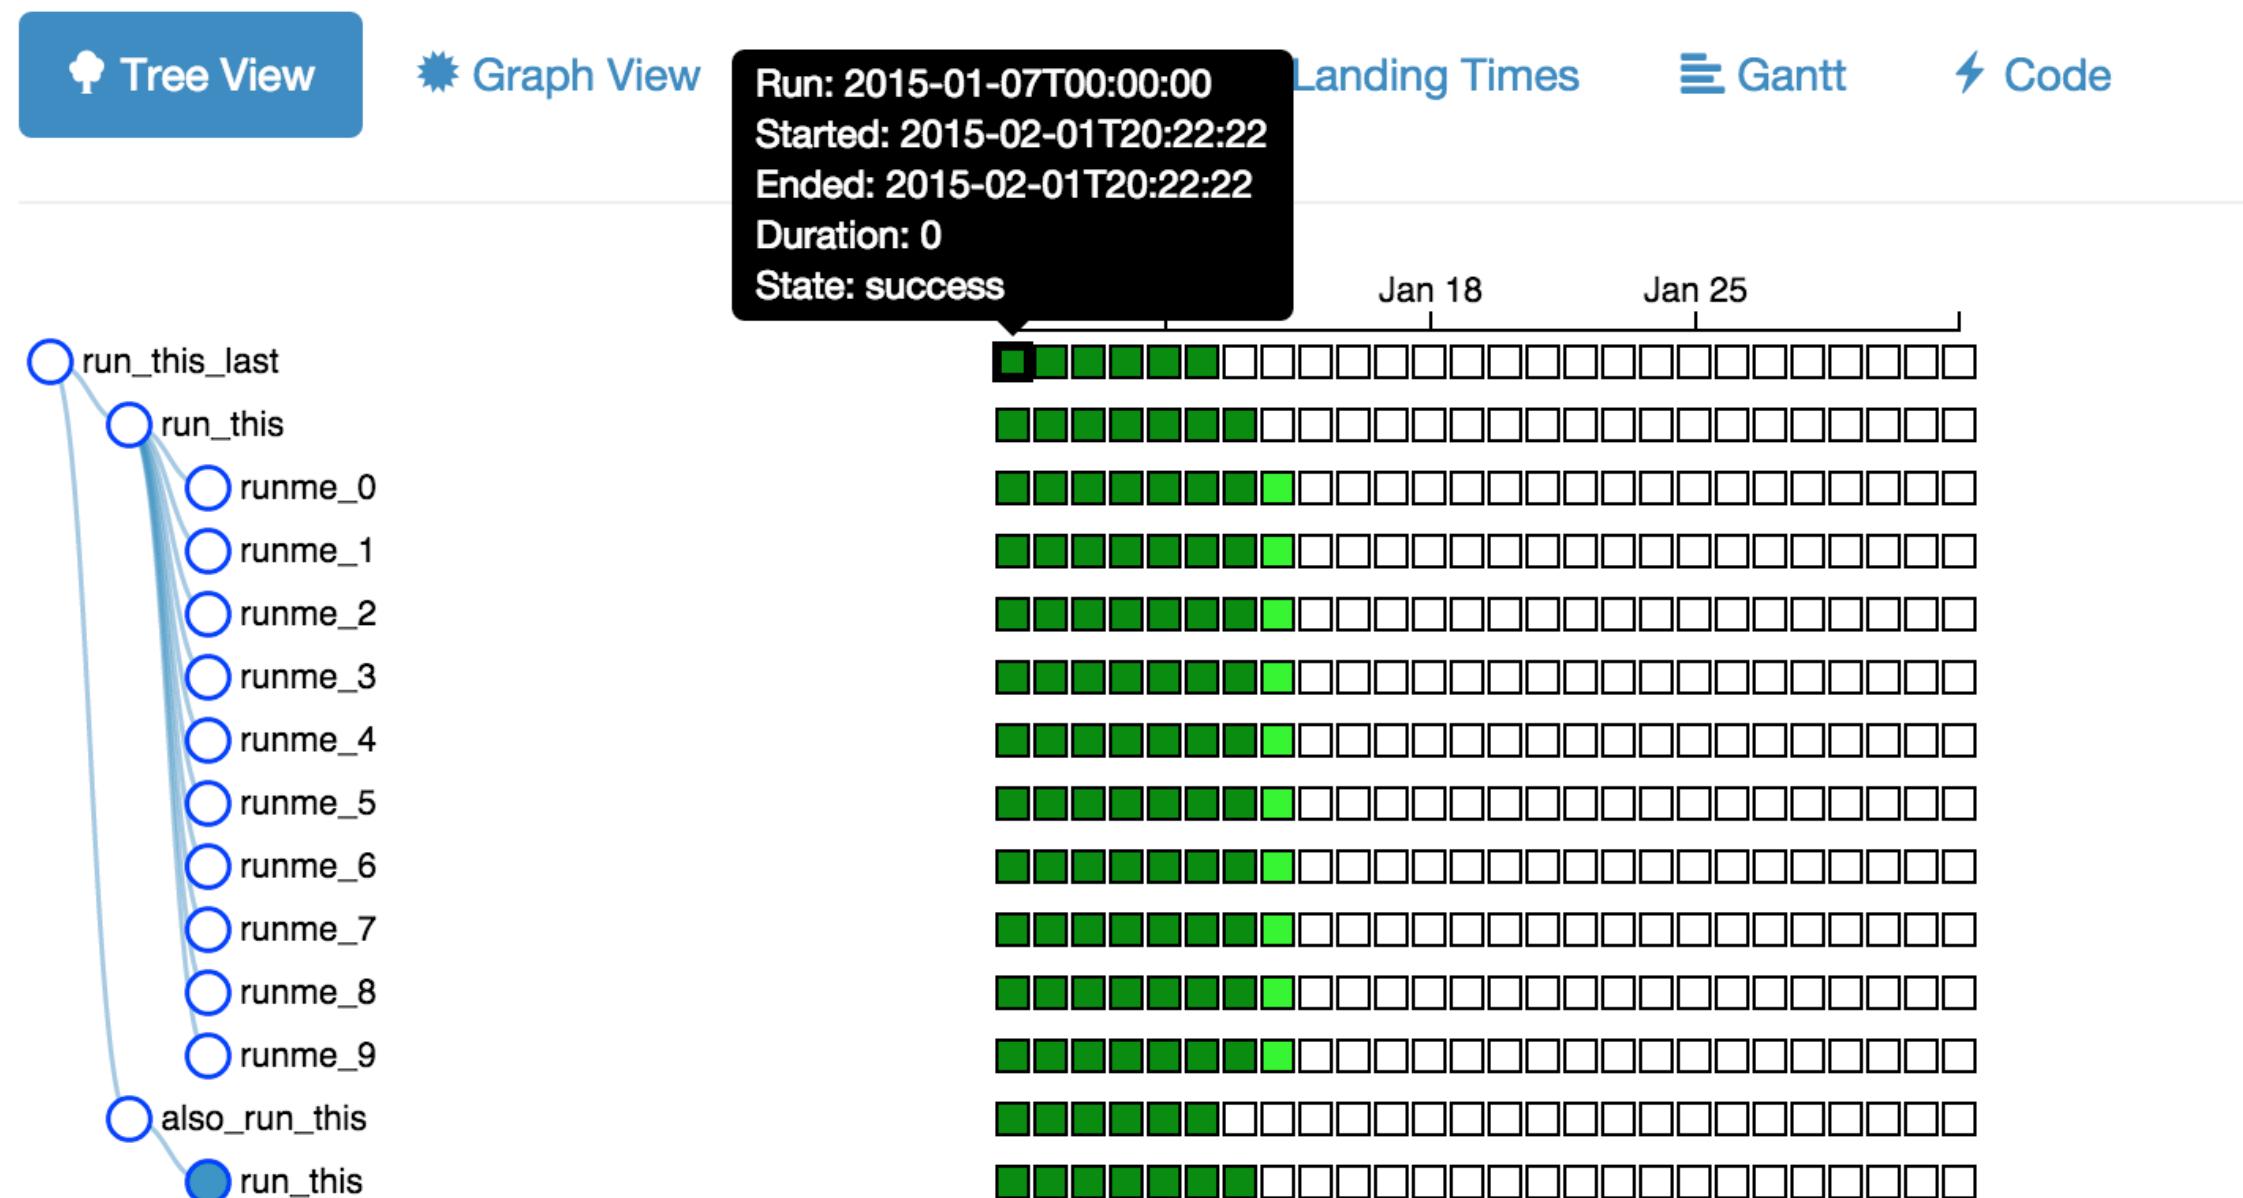
[Hide Paused DAGs](#)



Apache Airflow



DAG: example2



Creating a pipeline with Airflow

```
dag = DAG('tutorial', default_args=default_args)

# t1, t2 and t3 are examples of tasks created by instantiating operators
t1 = BashOperator(
    task_id='print_date',
    bash_command='date',
    dag=dag)

t2 = BashOperator(
    task_id='sleep',
    bash_command='sleep 5',
    retries=3,
    dag=dag)

templated_command = """
    {% for i in range(5) %}
        echo "{{ ds }}"
        echo "{{ macros.ds_add(ds, 7)}}"
        echo "{{ params.my_param }}"
    {% endfor %}
"""

t3 = BashOperator(
    task_id='templated',
    bash_command=templated_command,
    params={'my_param': 'Parameter I passed in'},
    dag=dag)

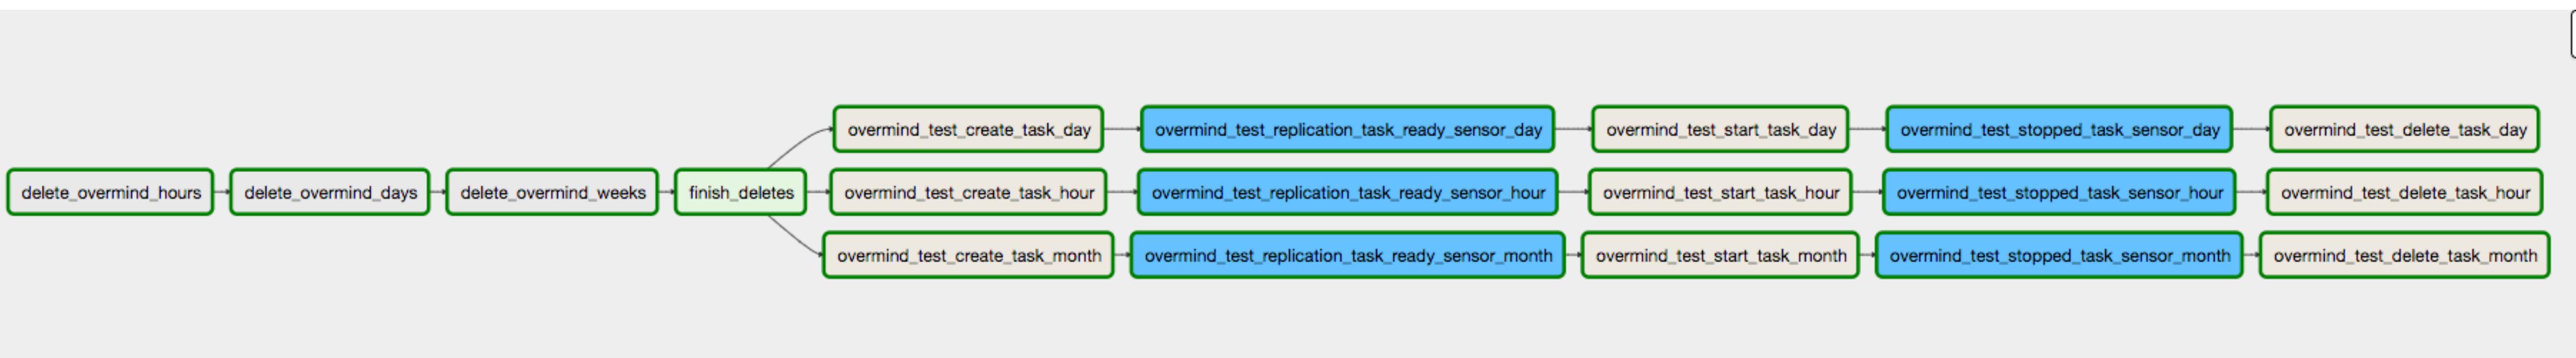
t2.set_upstream(t1)
t3.set_upstream(t1)
```

Define Operators

Set Dependencies



Managing a pipeline with Airflow



Why Airflow on Kubernetes?

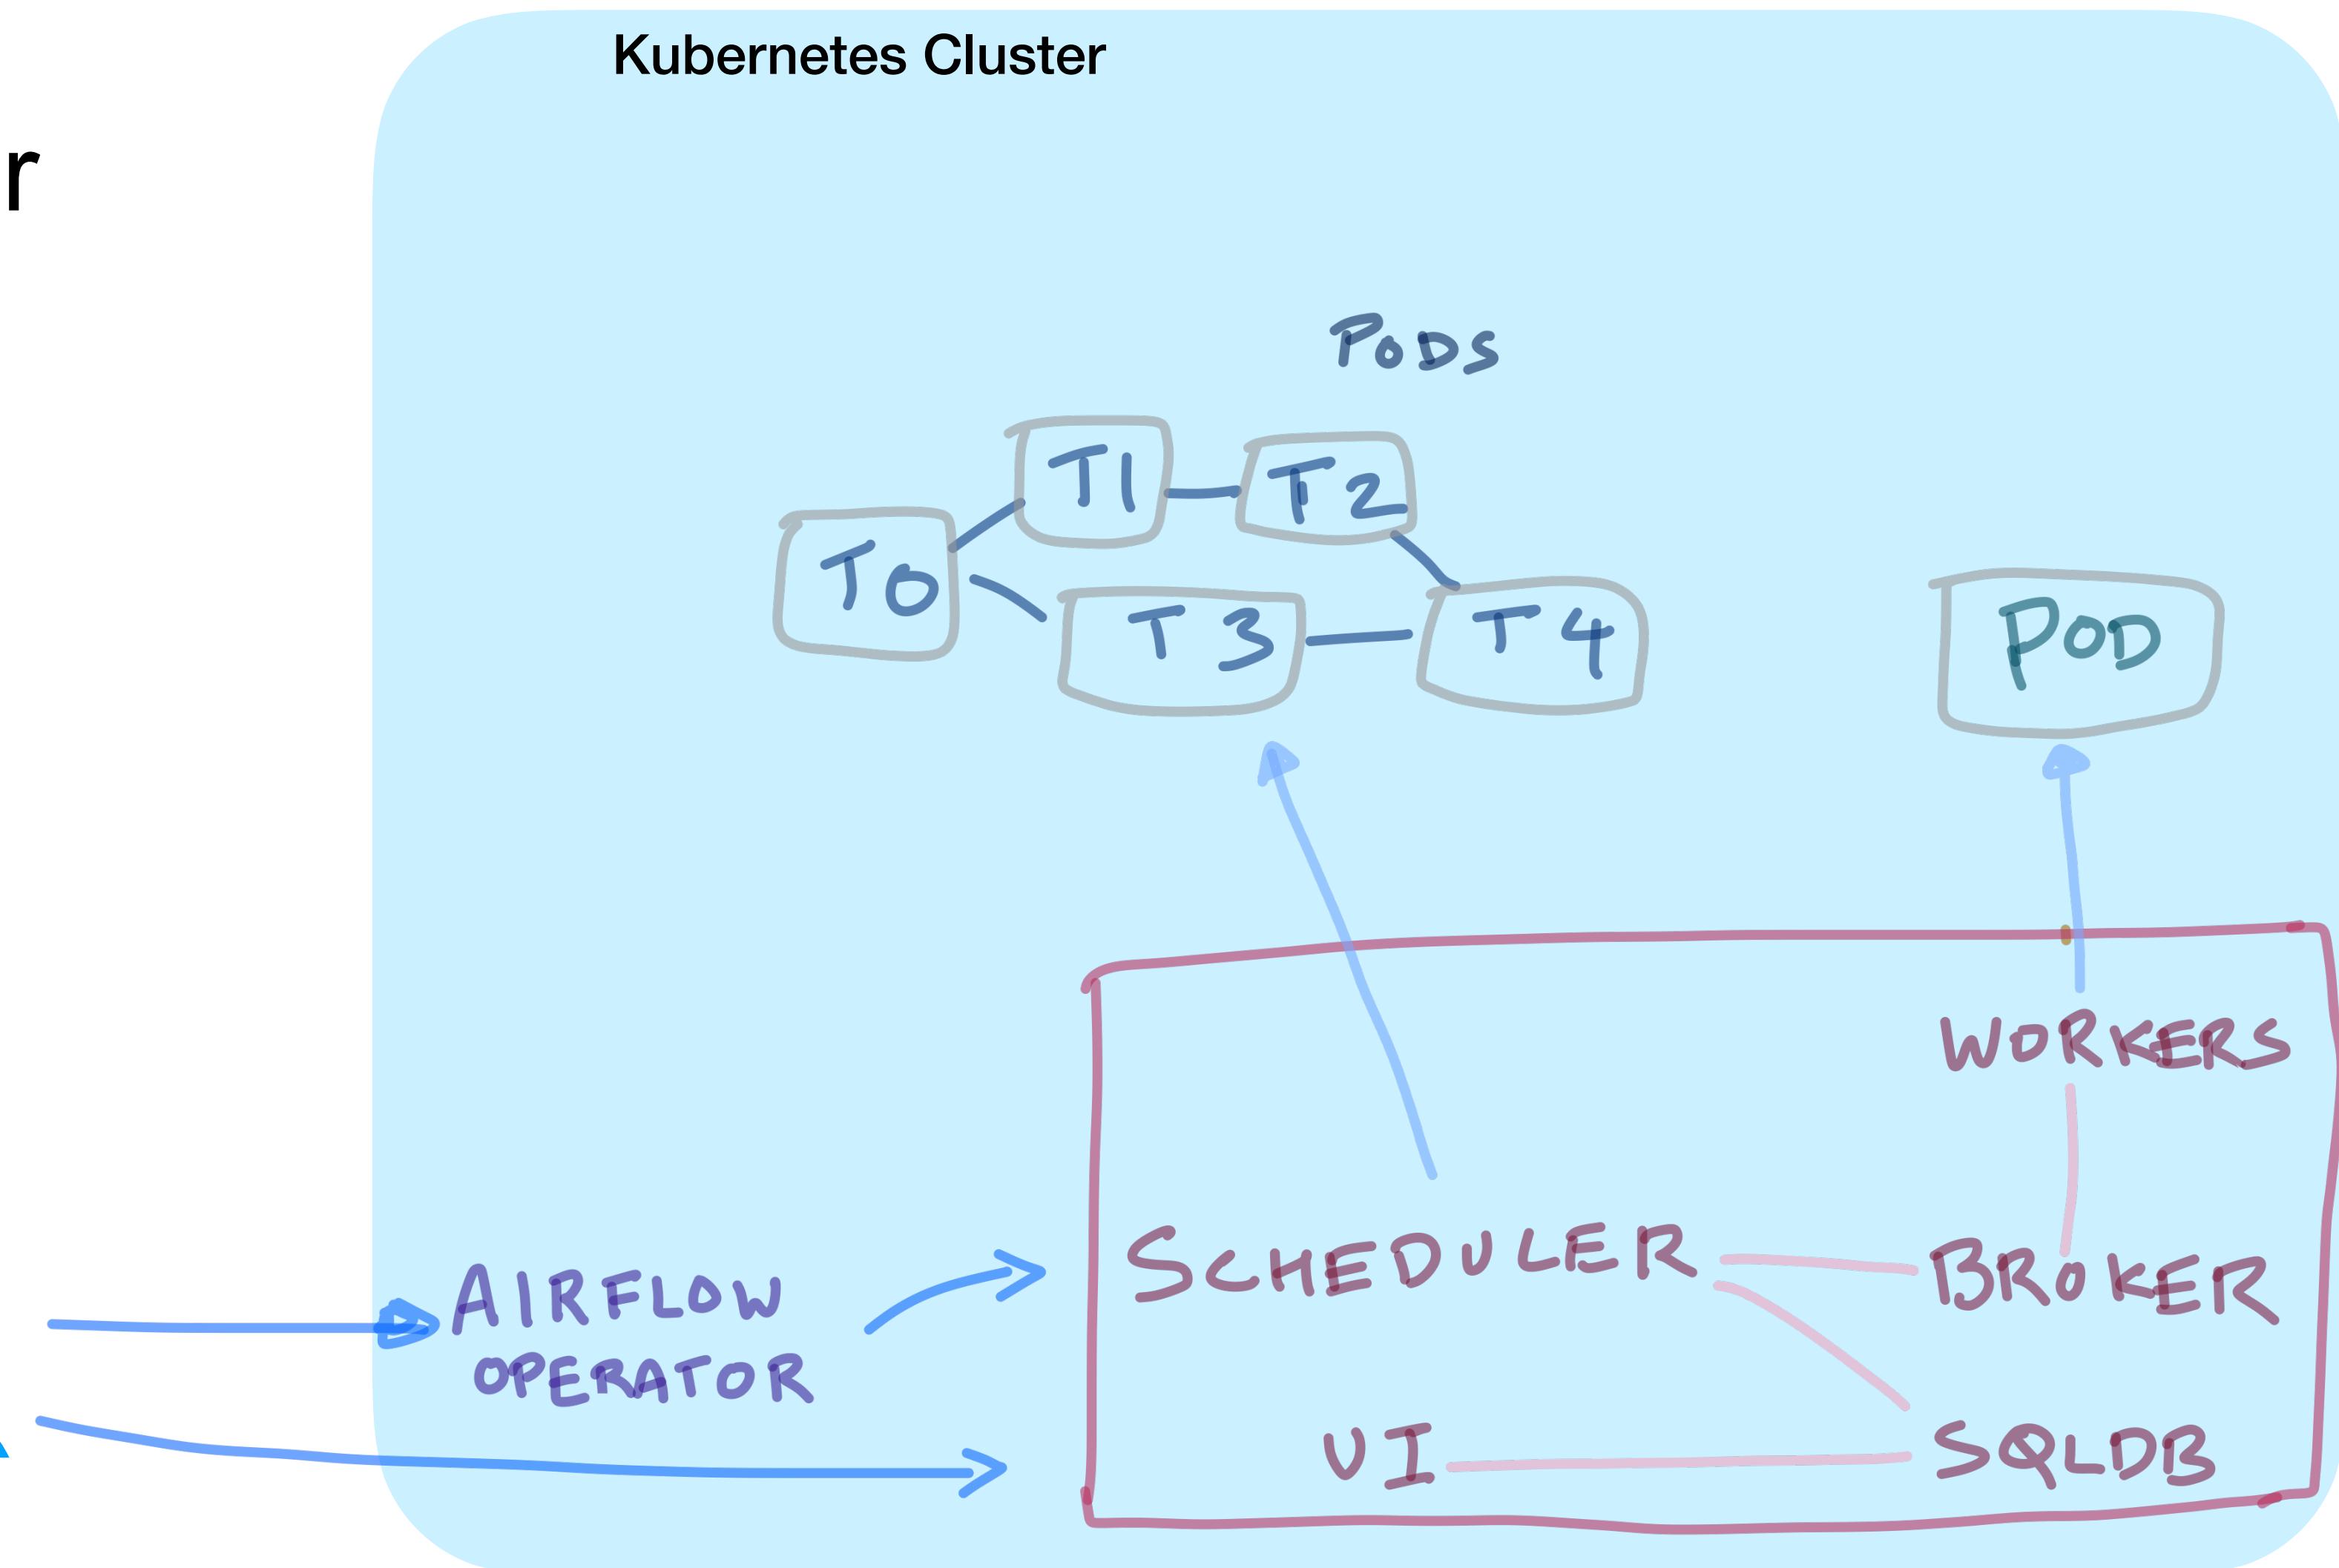
- More flexibility
- Dynamic Resource Allocation
- Lower Barrier to Entry



kubernetes

Airflow + Kubernetes

- Kubernetes Pod Operator
 - Flexibility
- Kubernetes Executor
 - Dynamic Allocation
- AirflowOperator (GCP)
 - One-step deployment

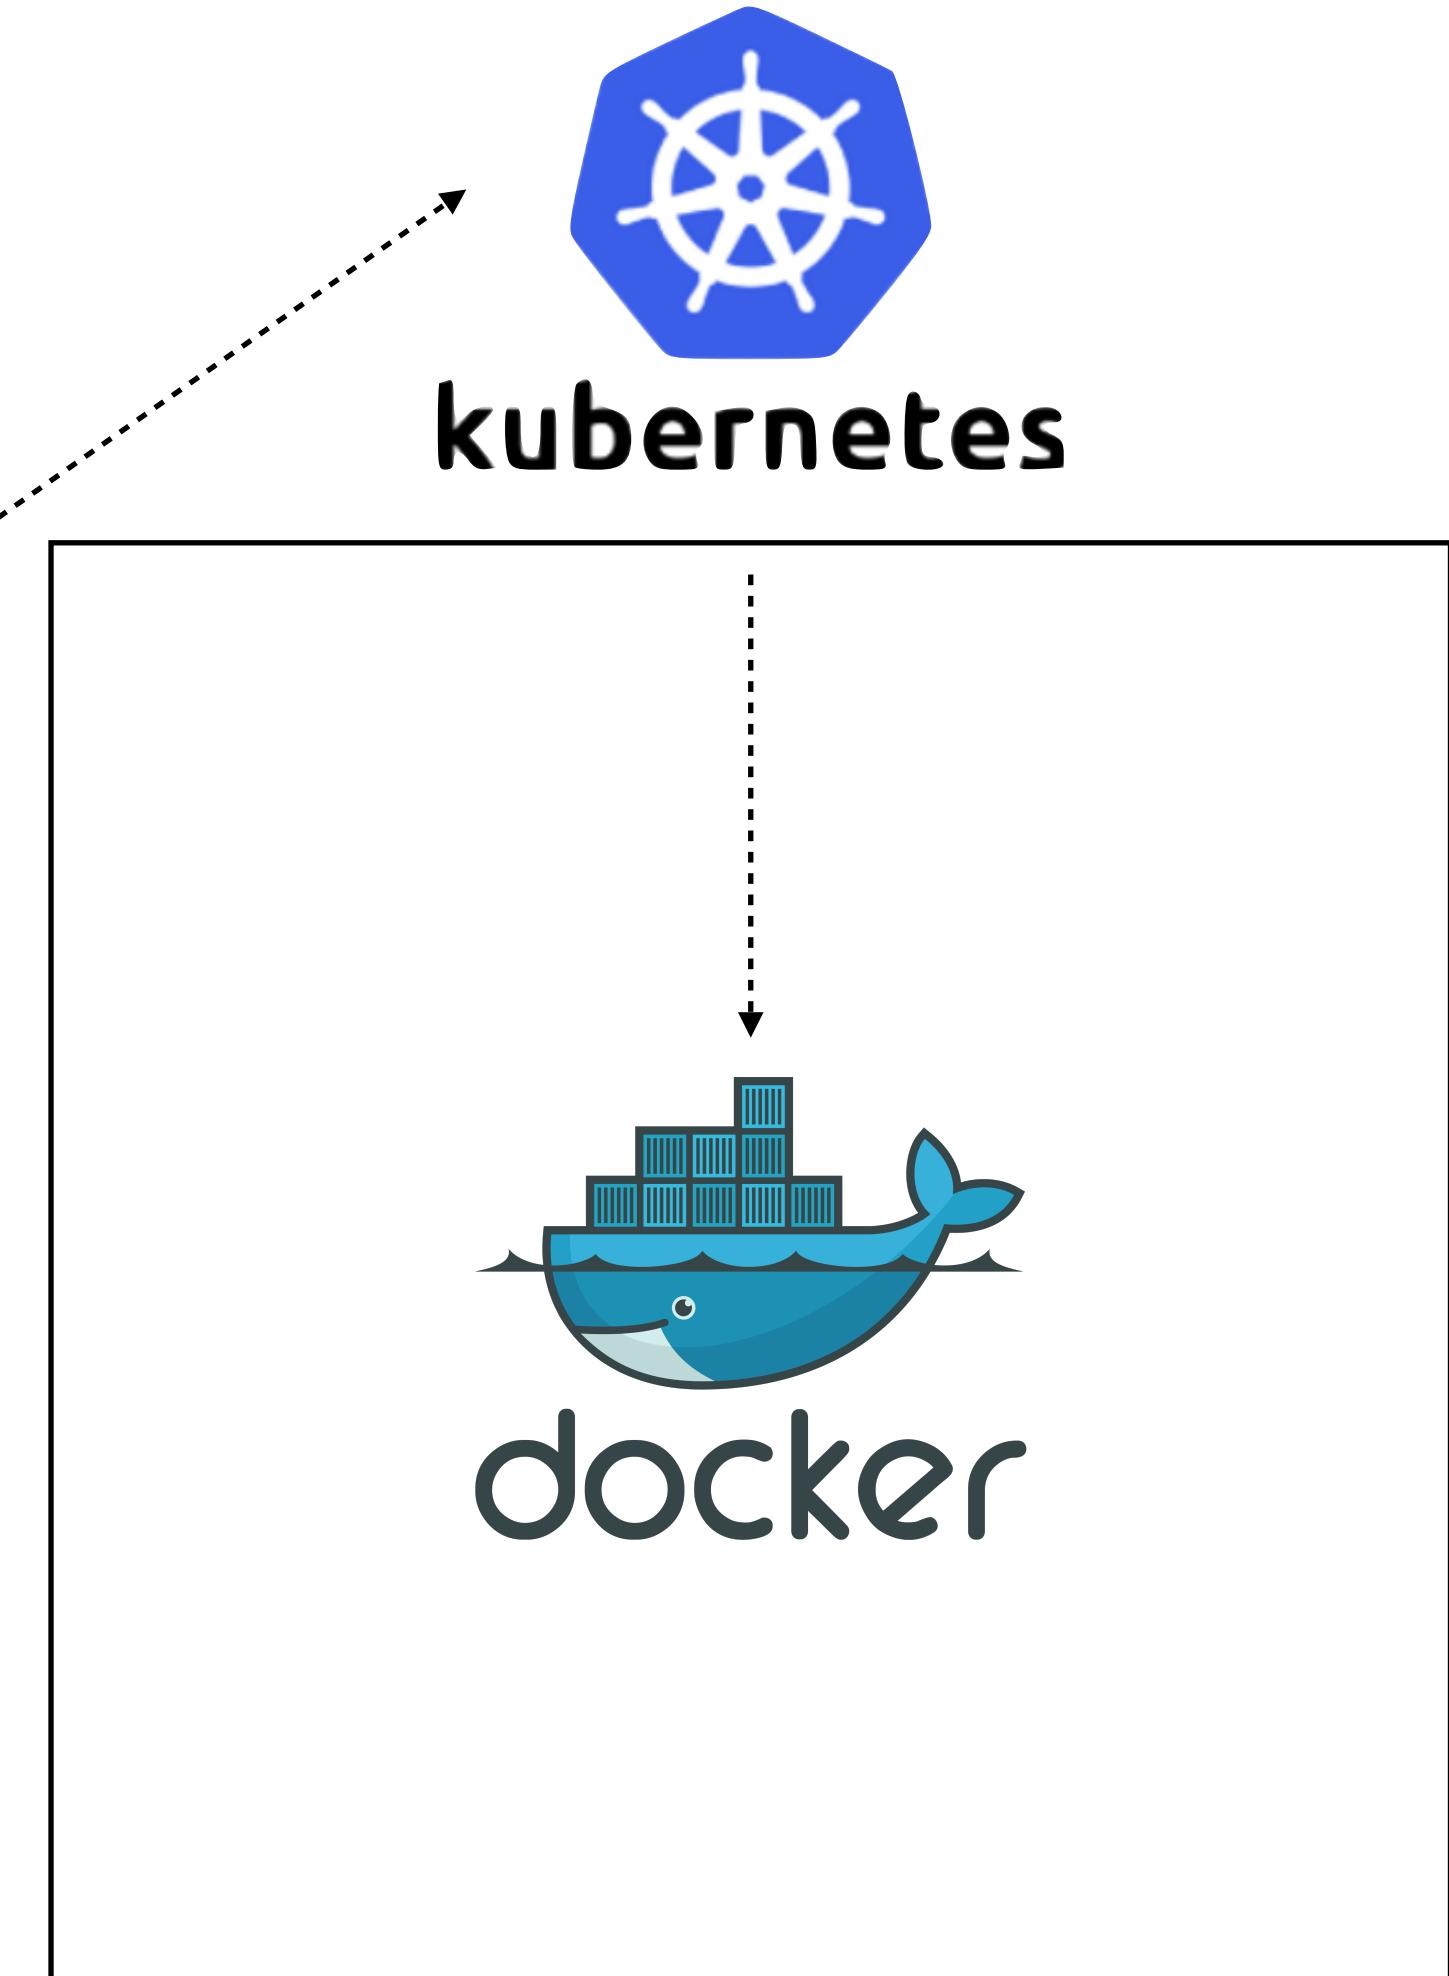


KubernetesPodOperator

- Allow users to deploy arbitrary Docker images
- Users can offload dependencies to containers
- “Lets Airflow focus on scheduling tasks”



Scheduler



KubernetesPodOperator

- Airflow workers are much lighter (don't require extra libraries)
- Easy rollbacks + deployments through tags

```
dag = DAG(
    'kubernetes_sample', default_args=default_args, schedule_interval=timedelta(minutes=10))

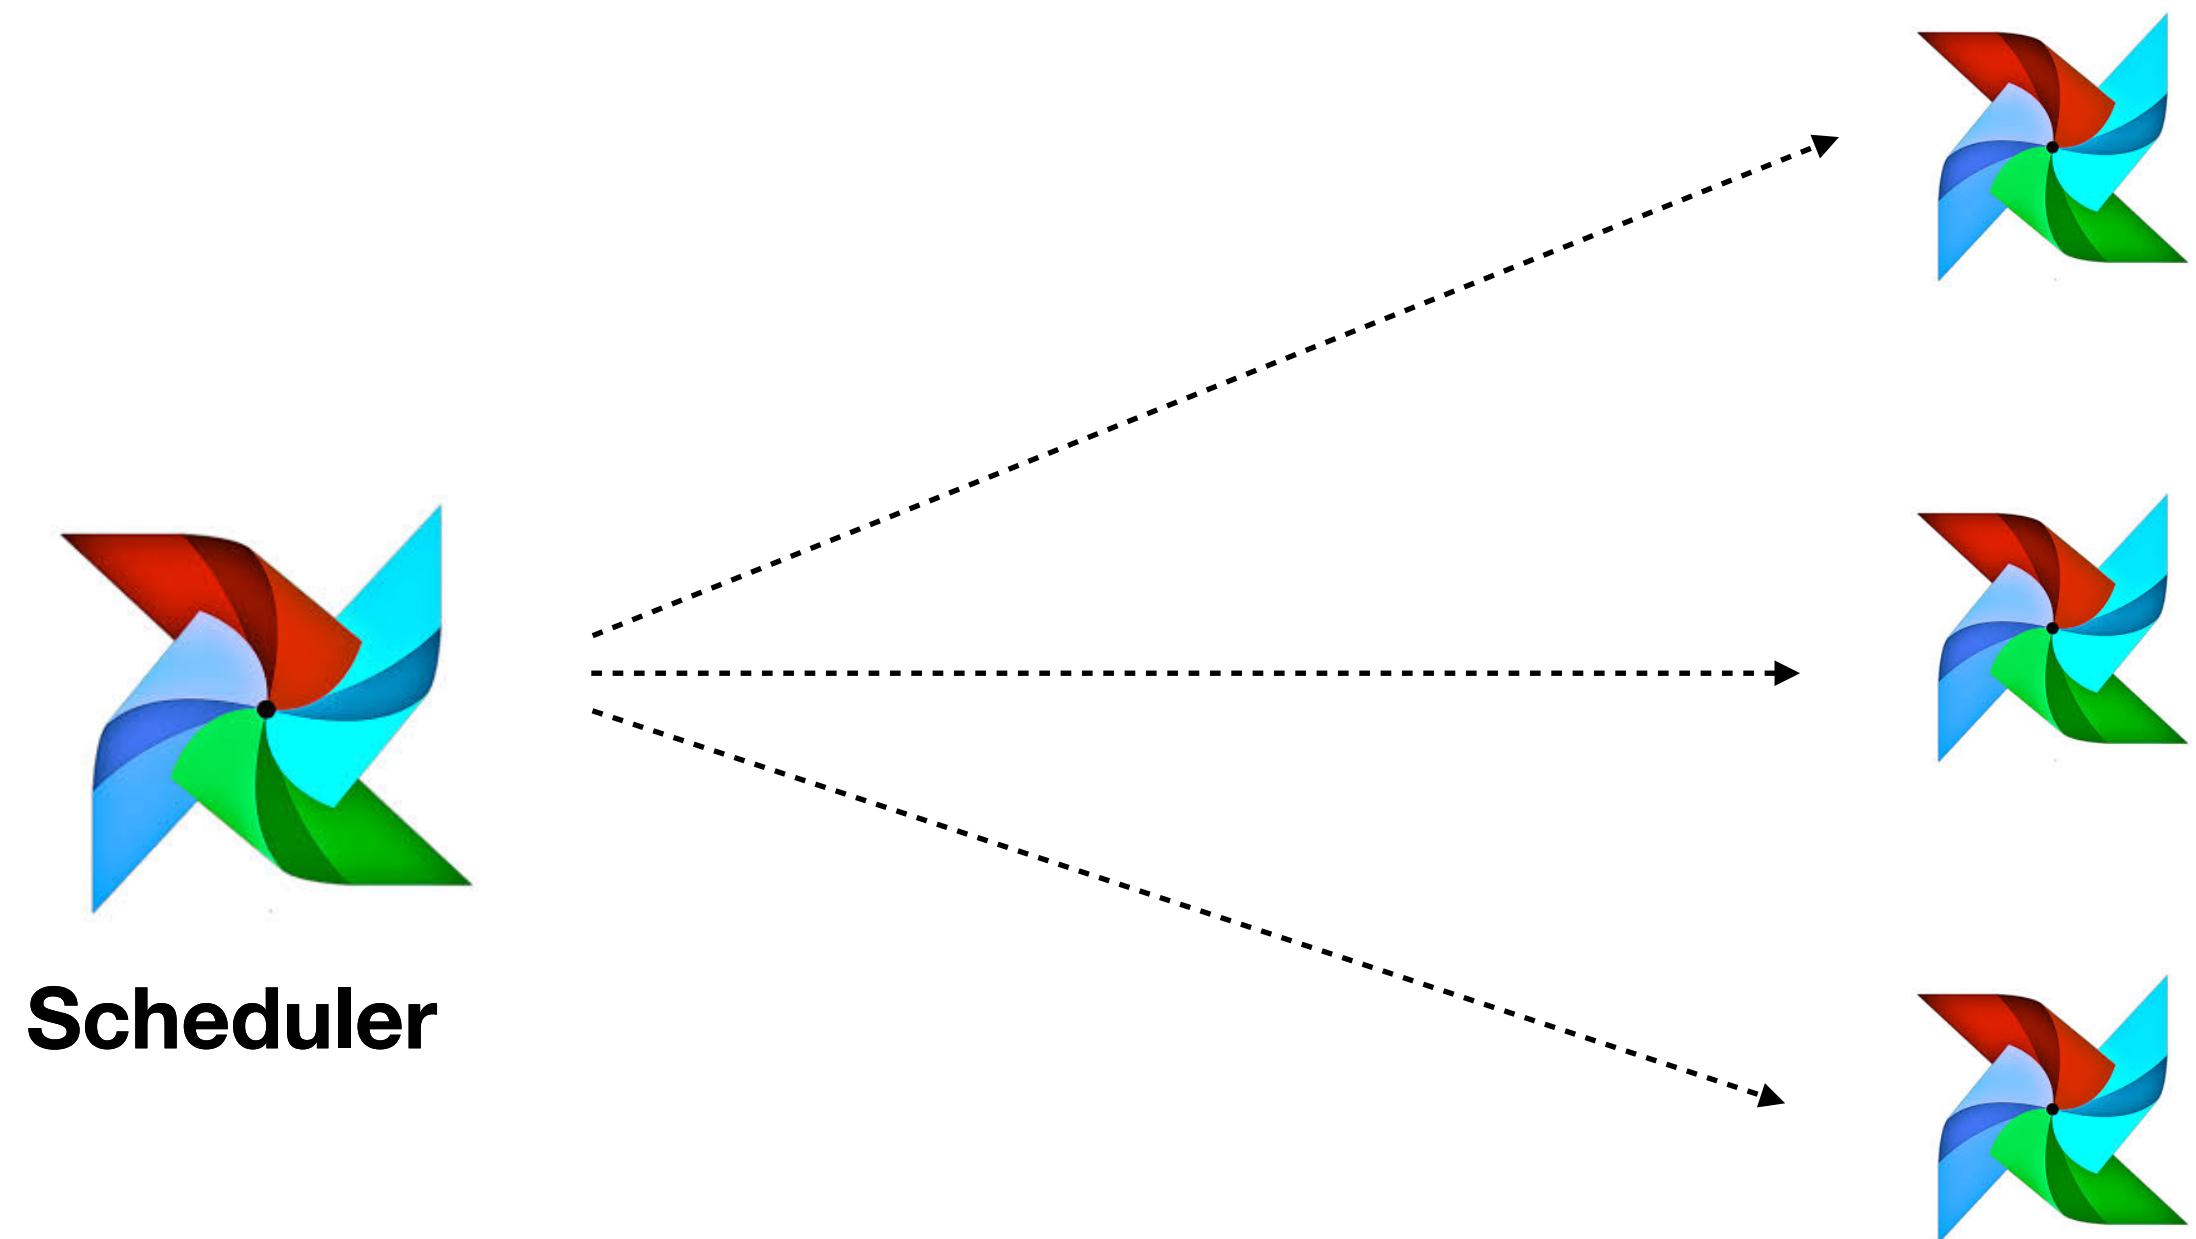
start = DummyOperator(task_id='run_this_first', dag=dag)

passing = KubernetesPodOperator(namespace='default',
                                 image="Python:3.6",
                                 cmd=["Python", "-c"],
                                 arguments=["print('hello world')"],
                                 labels={"foo": "bar"},
                                 name="passing-test",
                                 task_id="passing-task",
                                 get_logs=True,
                                 dag=dag
                                )

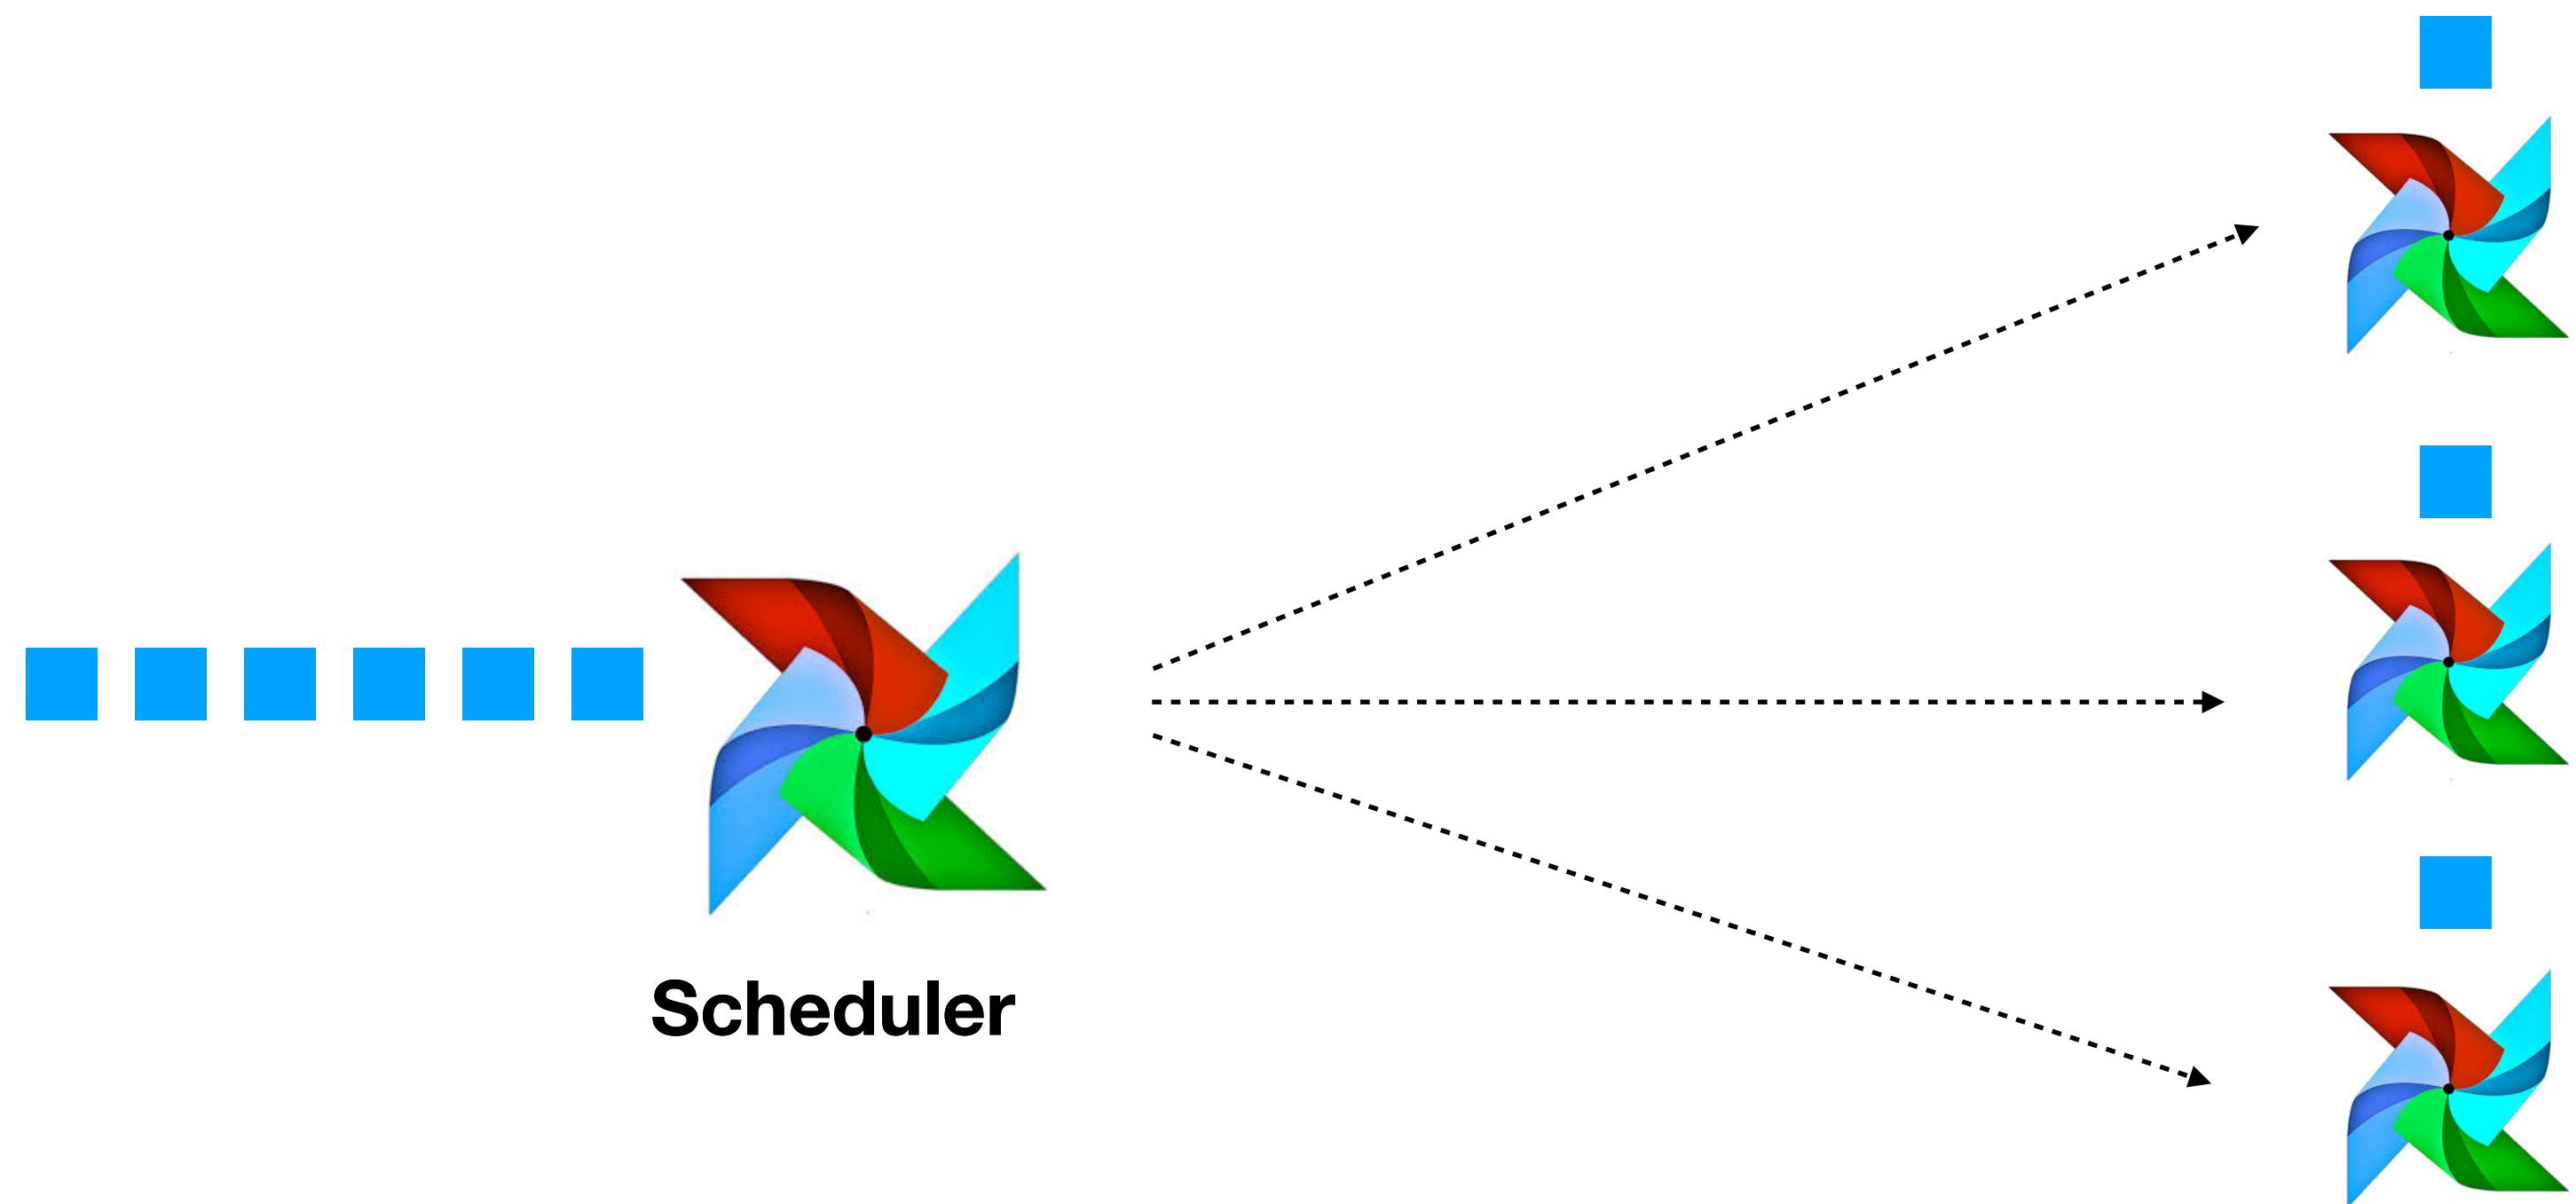
failing = KubernetesPodOperator(namespace='default',
                                 image="ubuntu:1604",
                                 cmd=["Python", "-c"],
                                 arguments=["print('hello world')"],
                                 labels={"foo": "bar"},
                                 name="fail",
                                 task_id="failing-task",
                                 get_logs=True,
                                 dag=dag
                                )

passing.set_upstream(start)
failing.set_upstream(start)
```

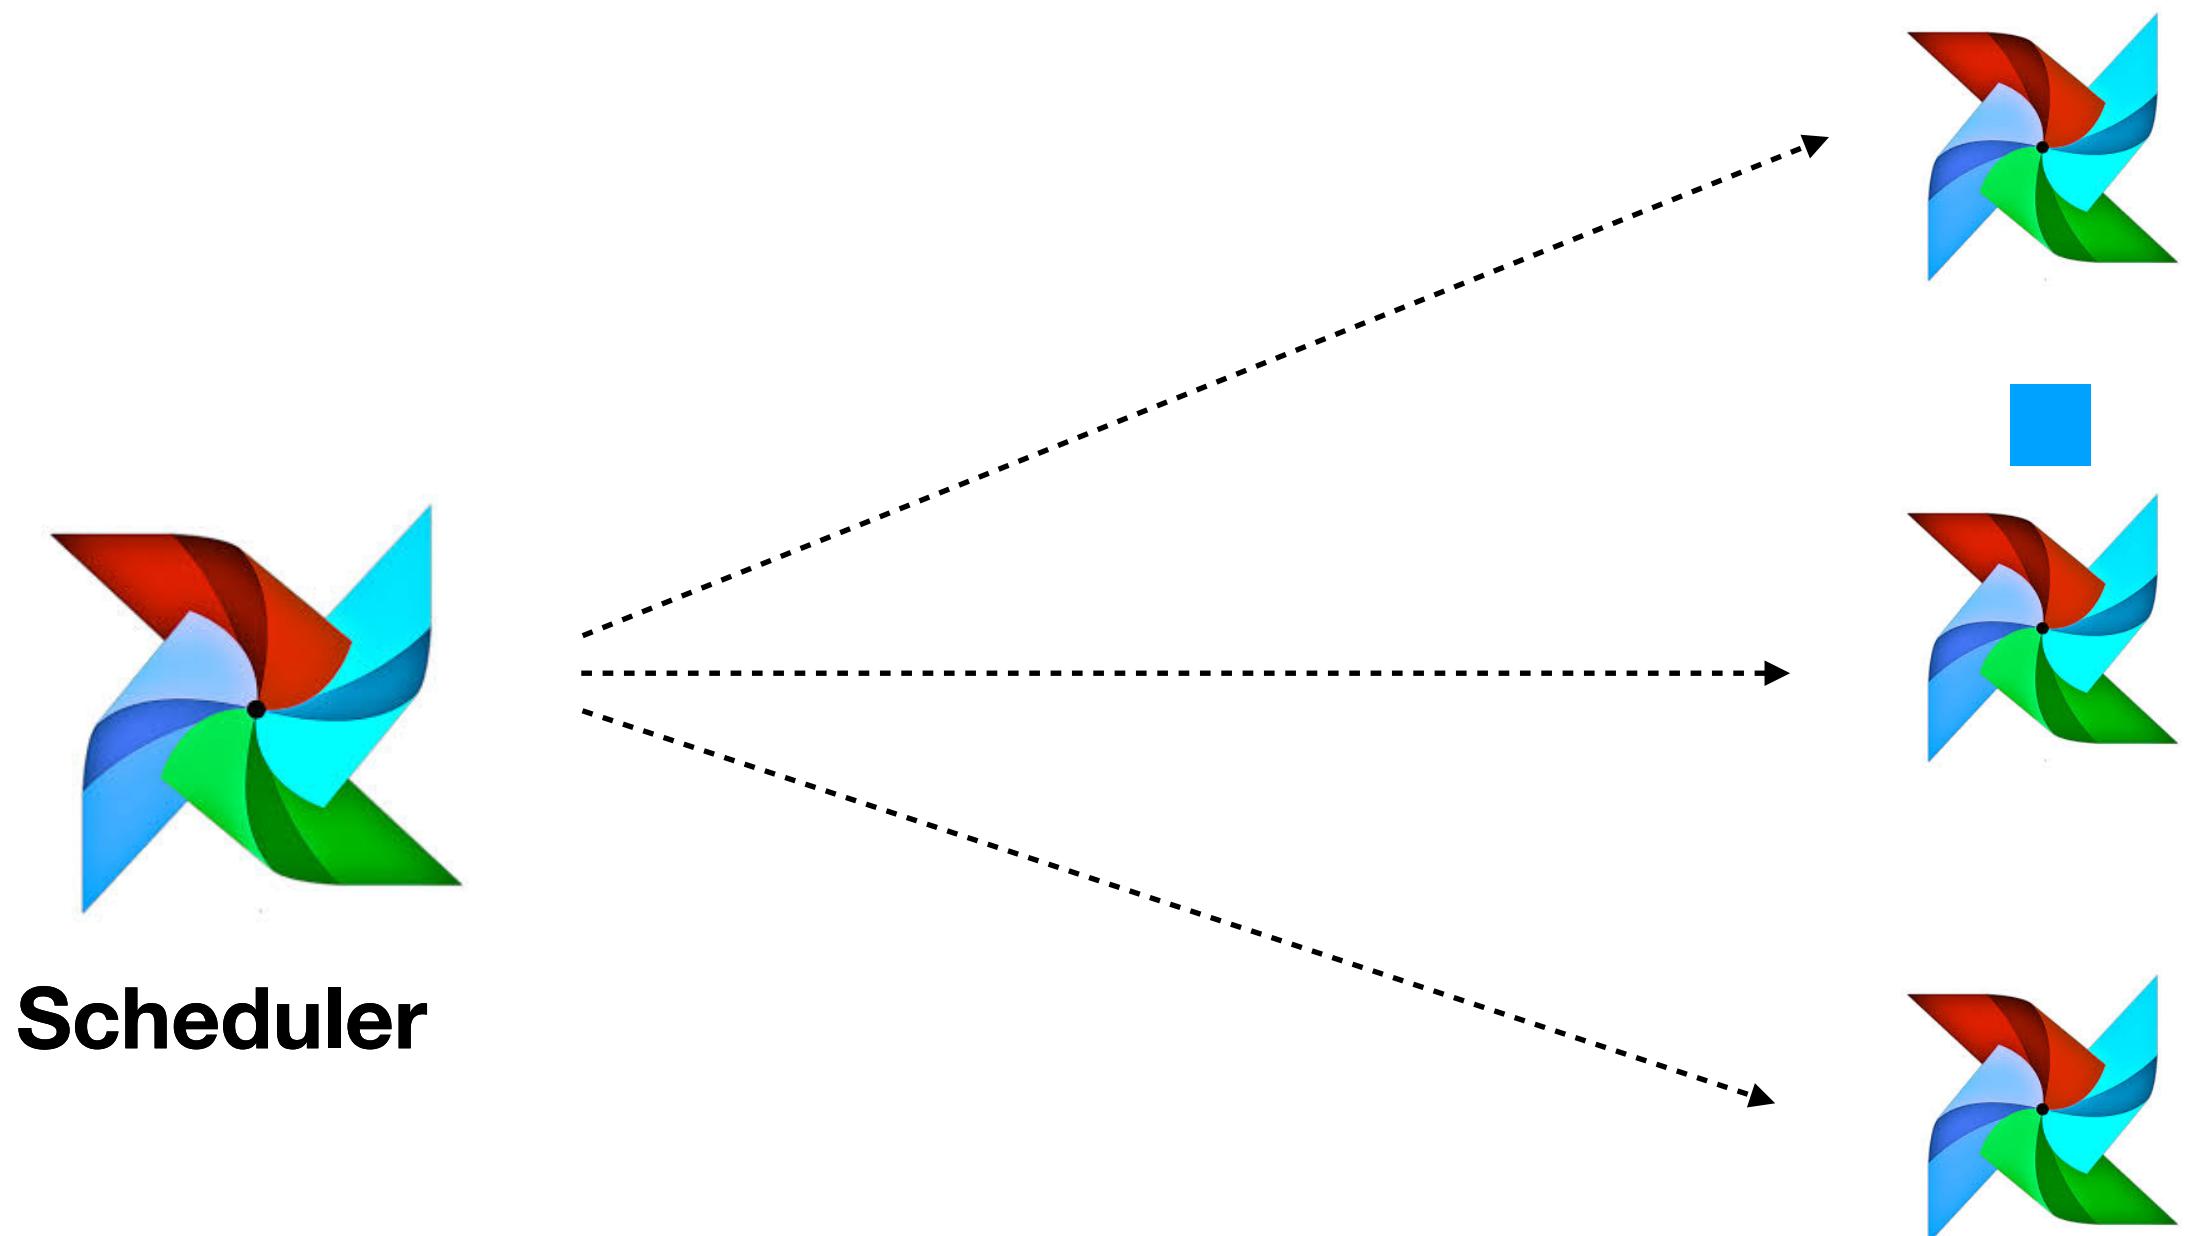
Static Allocation



Static Allocation



Static Allocation



Dynamic Resource Allocation

- All previous Airflow solutions involved static clusters of workers
- Require over-provisioning
- Airflow workers need to handle dependencies for multiple job types

shoiuld we keep this?

Kubernetes Executor

- High levels of parallelism (dynamic allocation)
- Task-level pod configuration
- Fault Tolerance

Dynamic Allocation

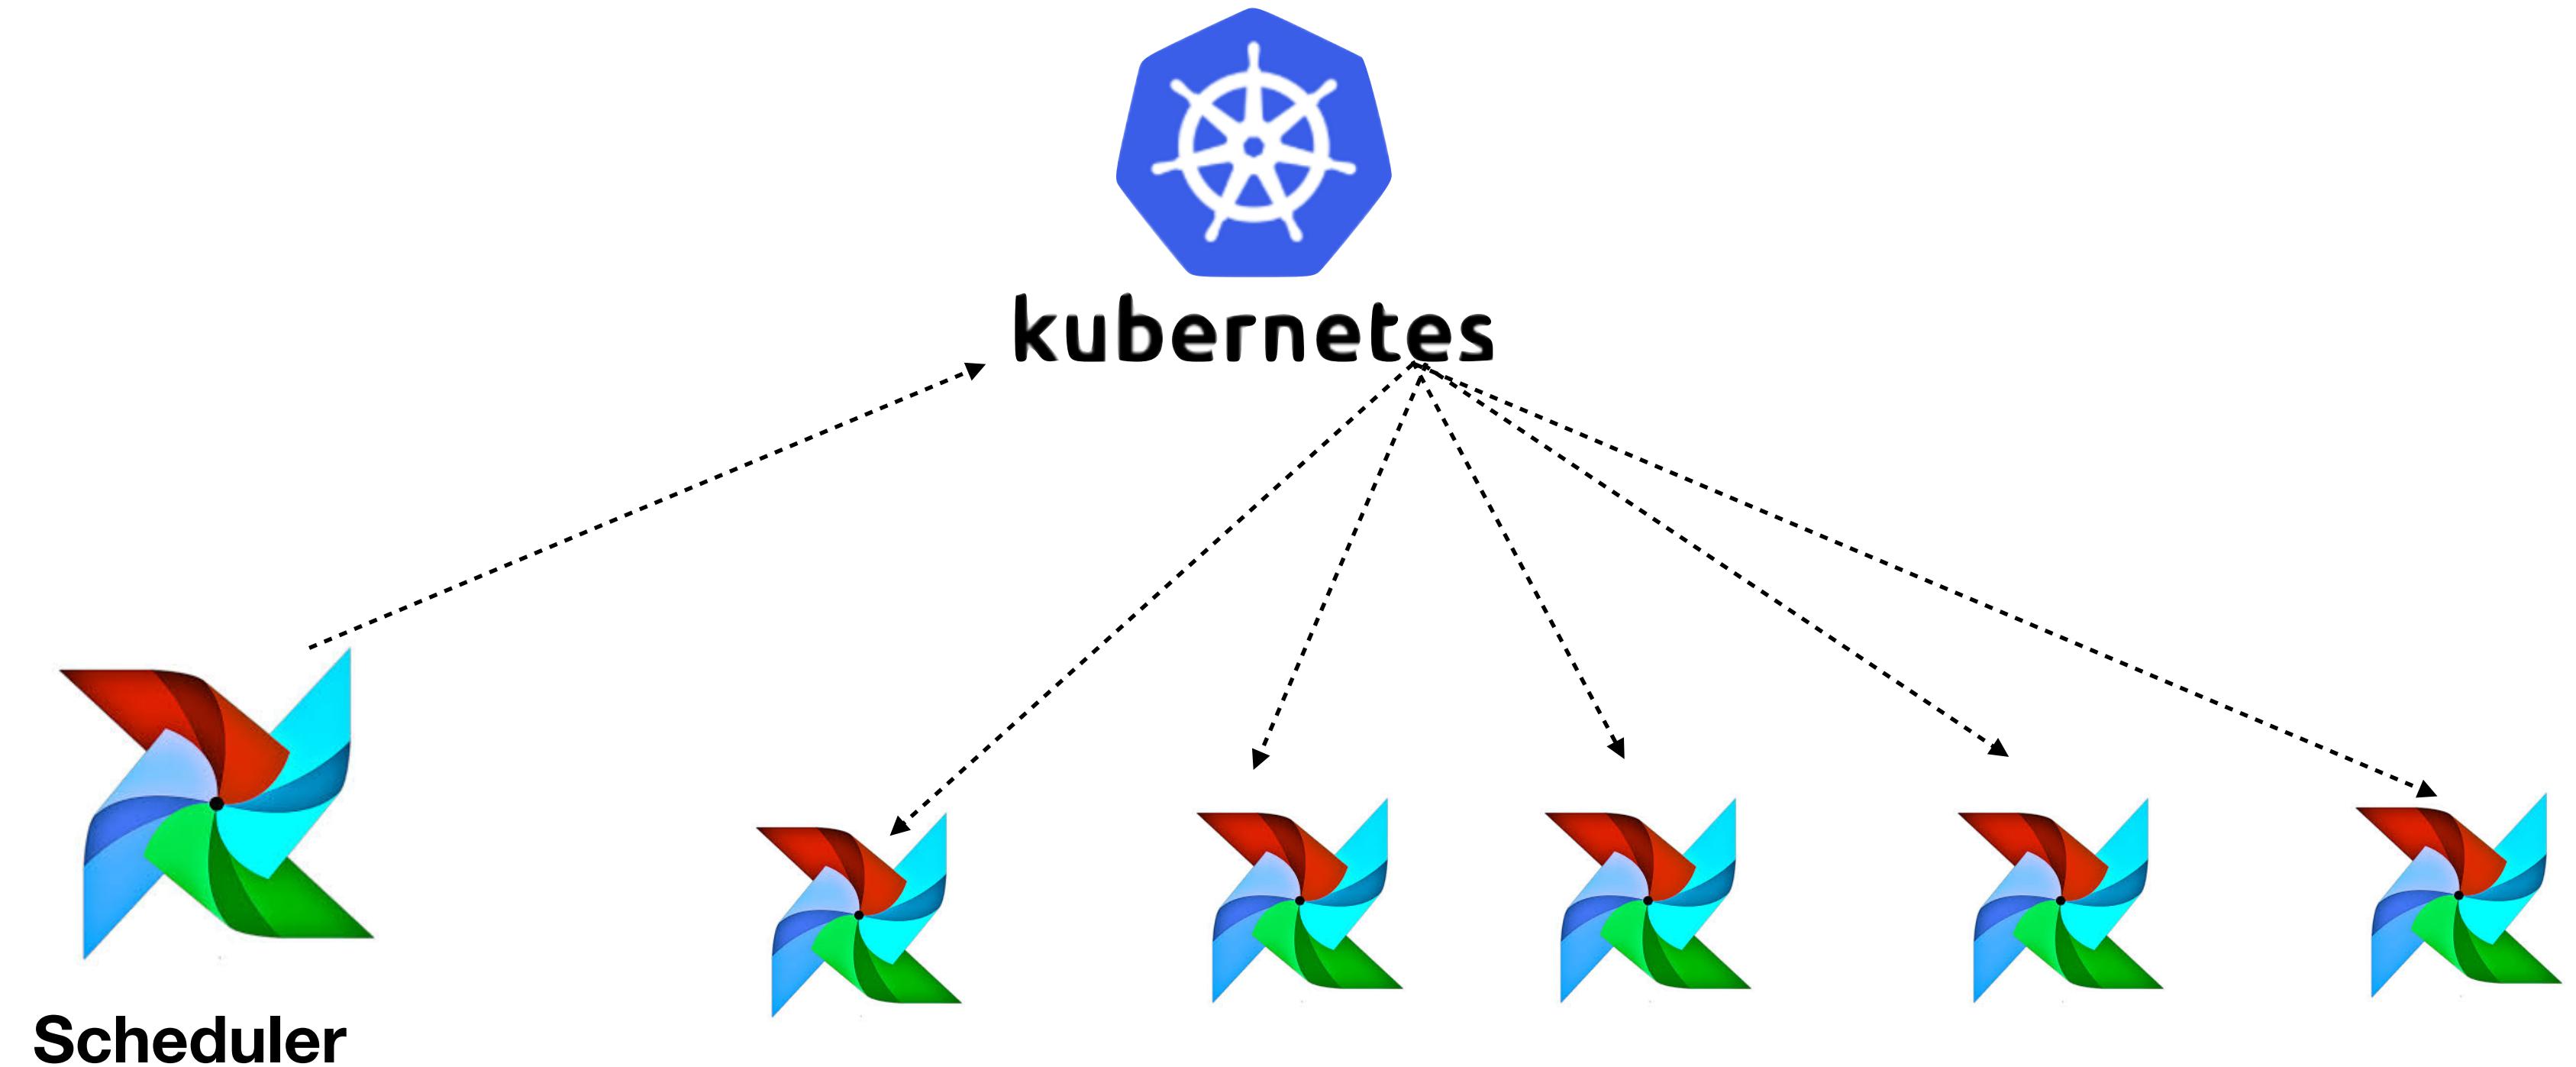


KubeCon



CloudNativeCon

North America 2018



Task Level Configs

```
t = BashOperator(  
    task_id = 'account-test',  
    bash_command = 'gcloud auth application-default login',  
    dag = dag,  
  
    executor_config = {  
        'request_memory': '128Mi',  
        'limit_memory': '128Mi'  
        'image': 'airflow/scipy:1.1.5'  
        'gcp-service-account': 'service-account@xxx.iam.gserviceaccount.com'  
    }  
)
```

Dynamic Allocation

- Don't need to think about capacity planning
- Great for time-sensitive jobs, which would traditionally require over-provisioning

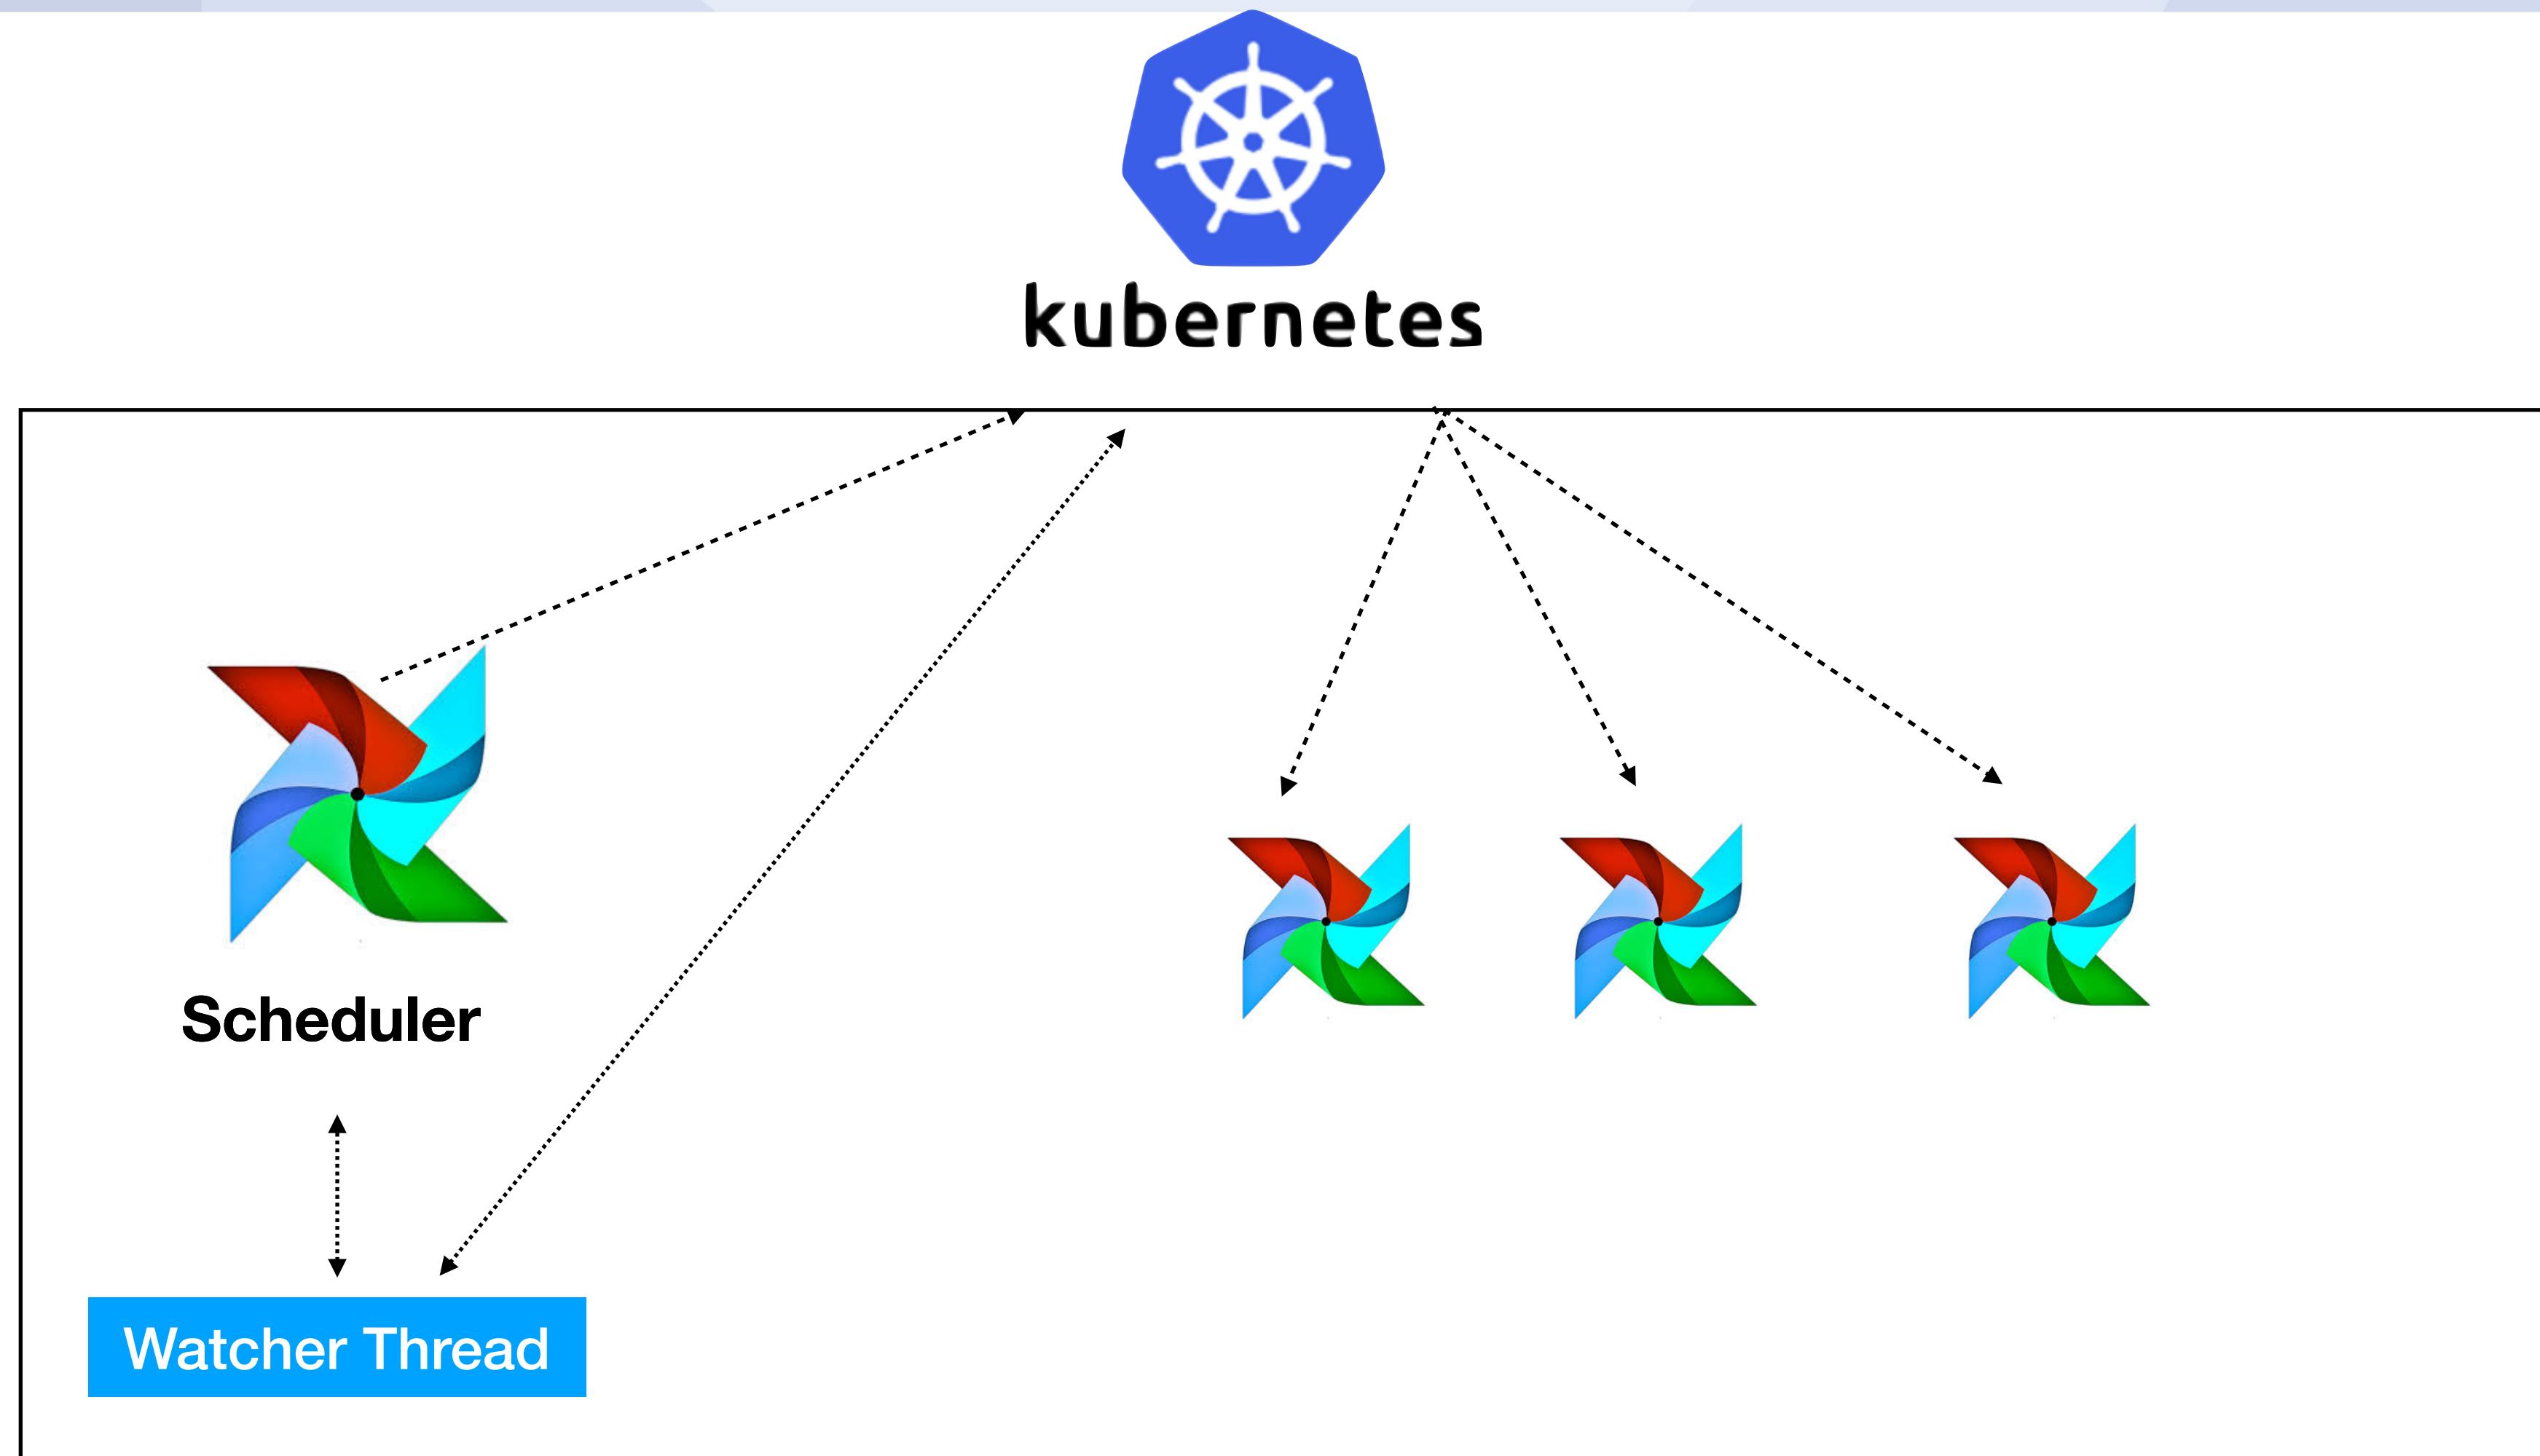
Kubernetes Executor



KubeCon

CloudNativeCon

North America 2018



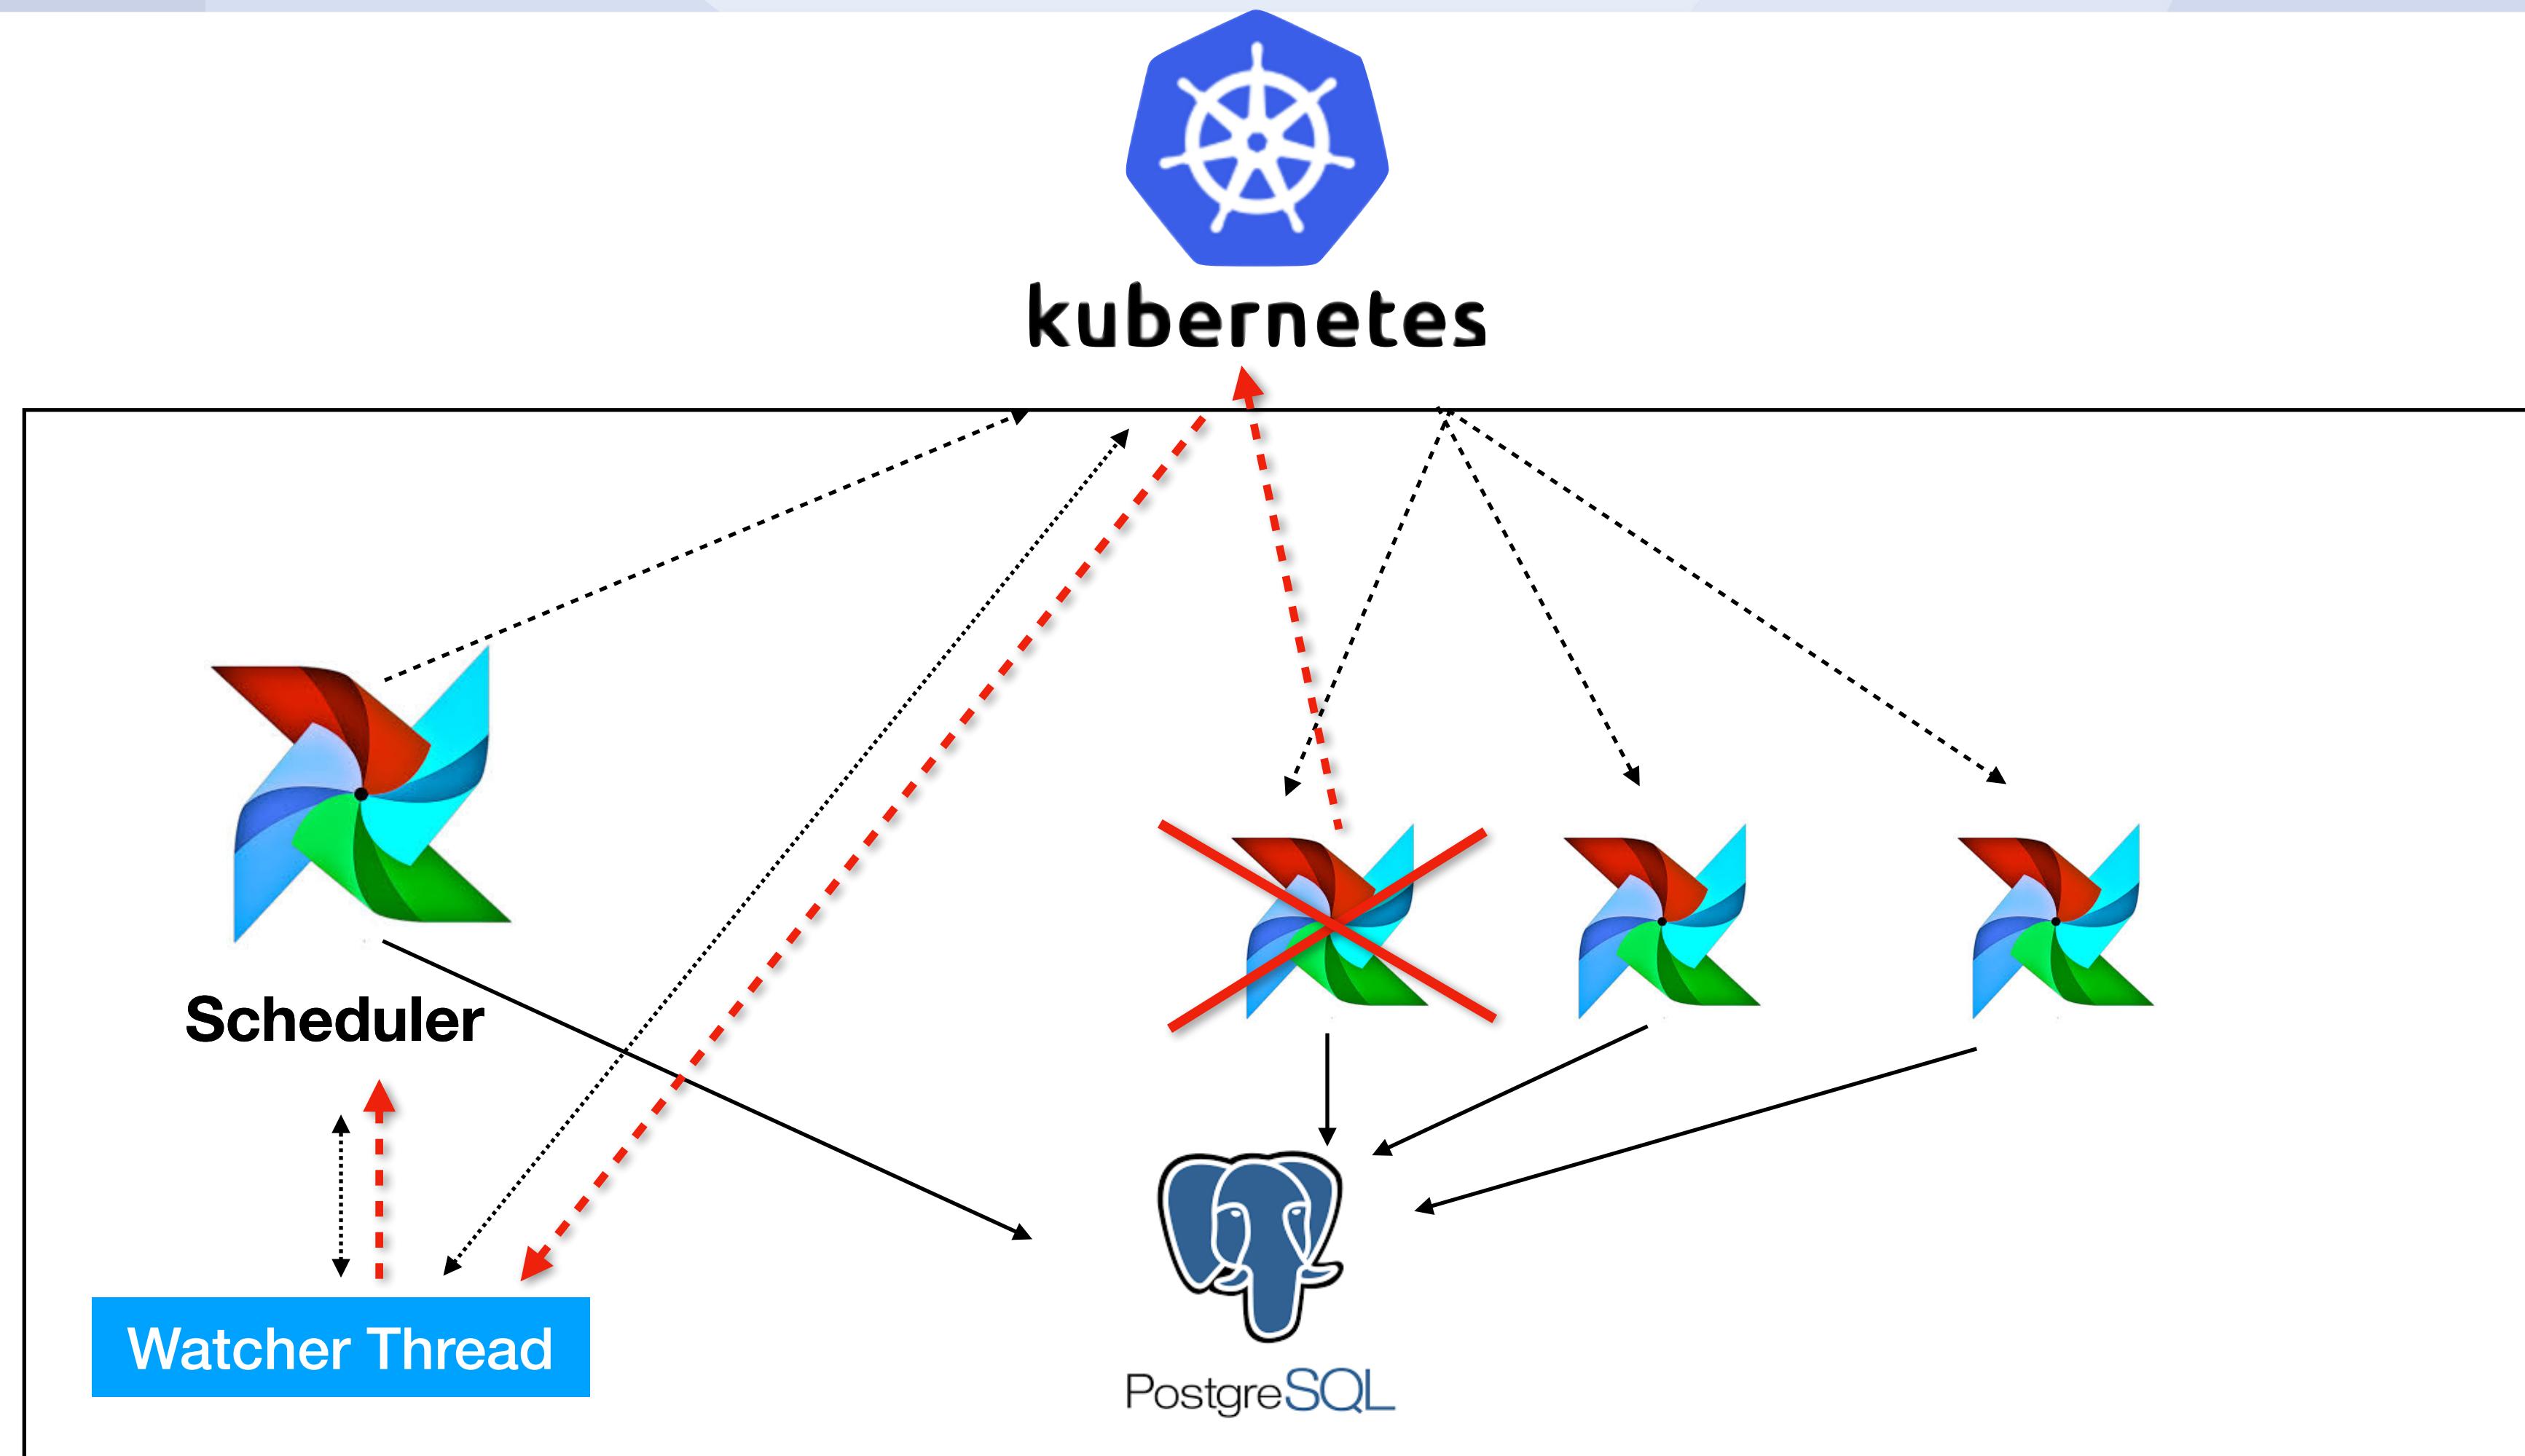
Kubernetes Executor



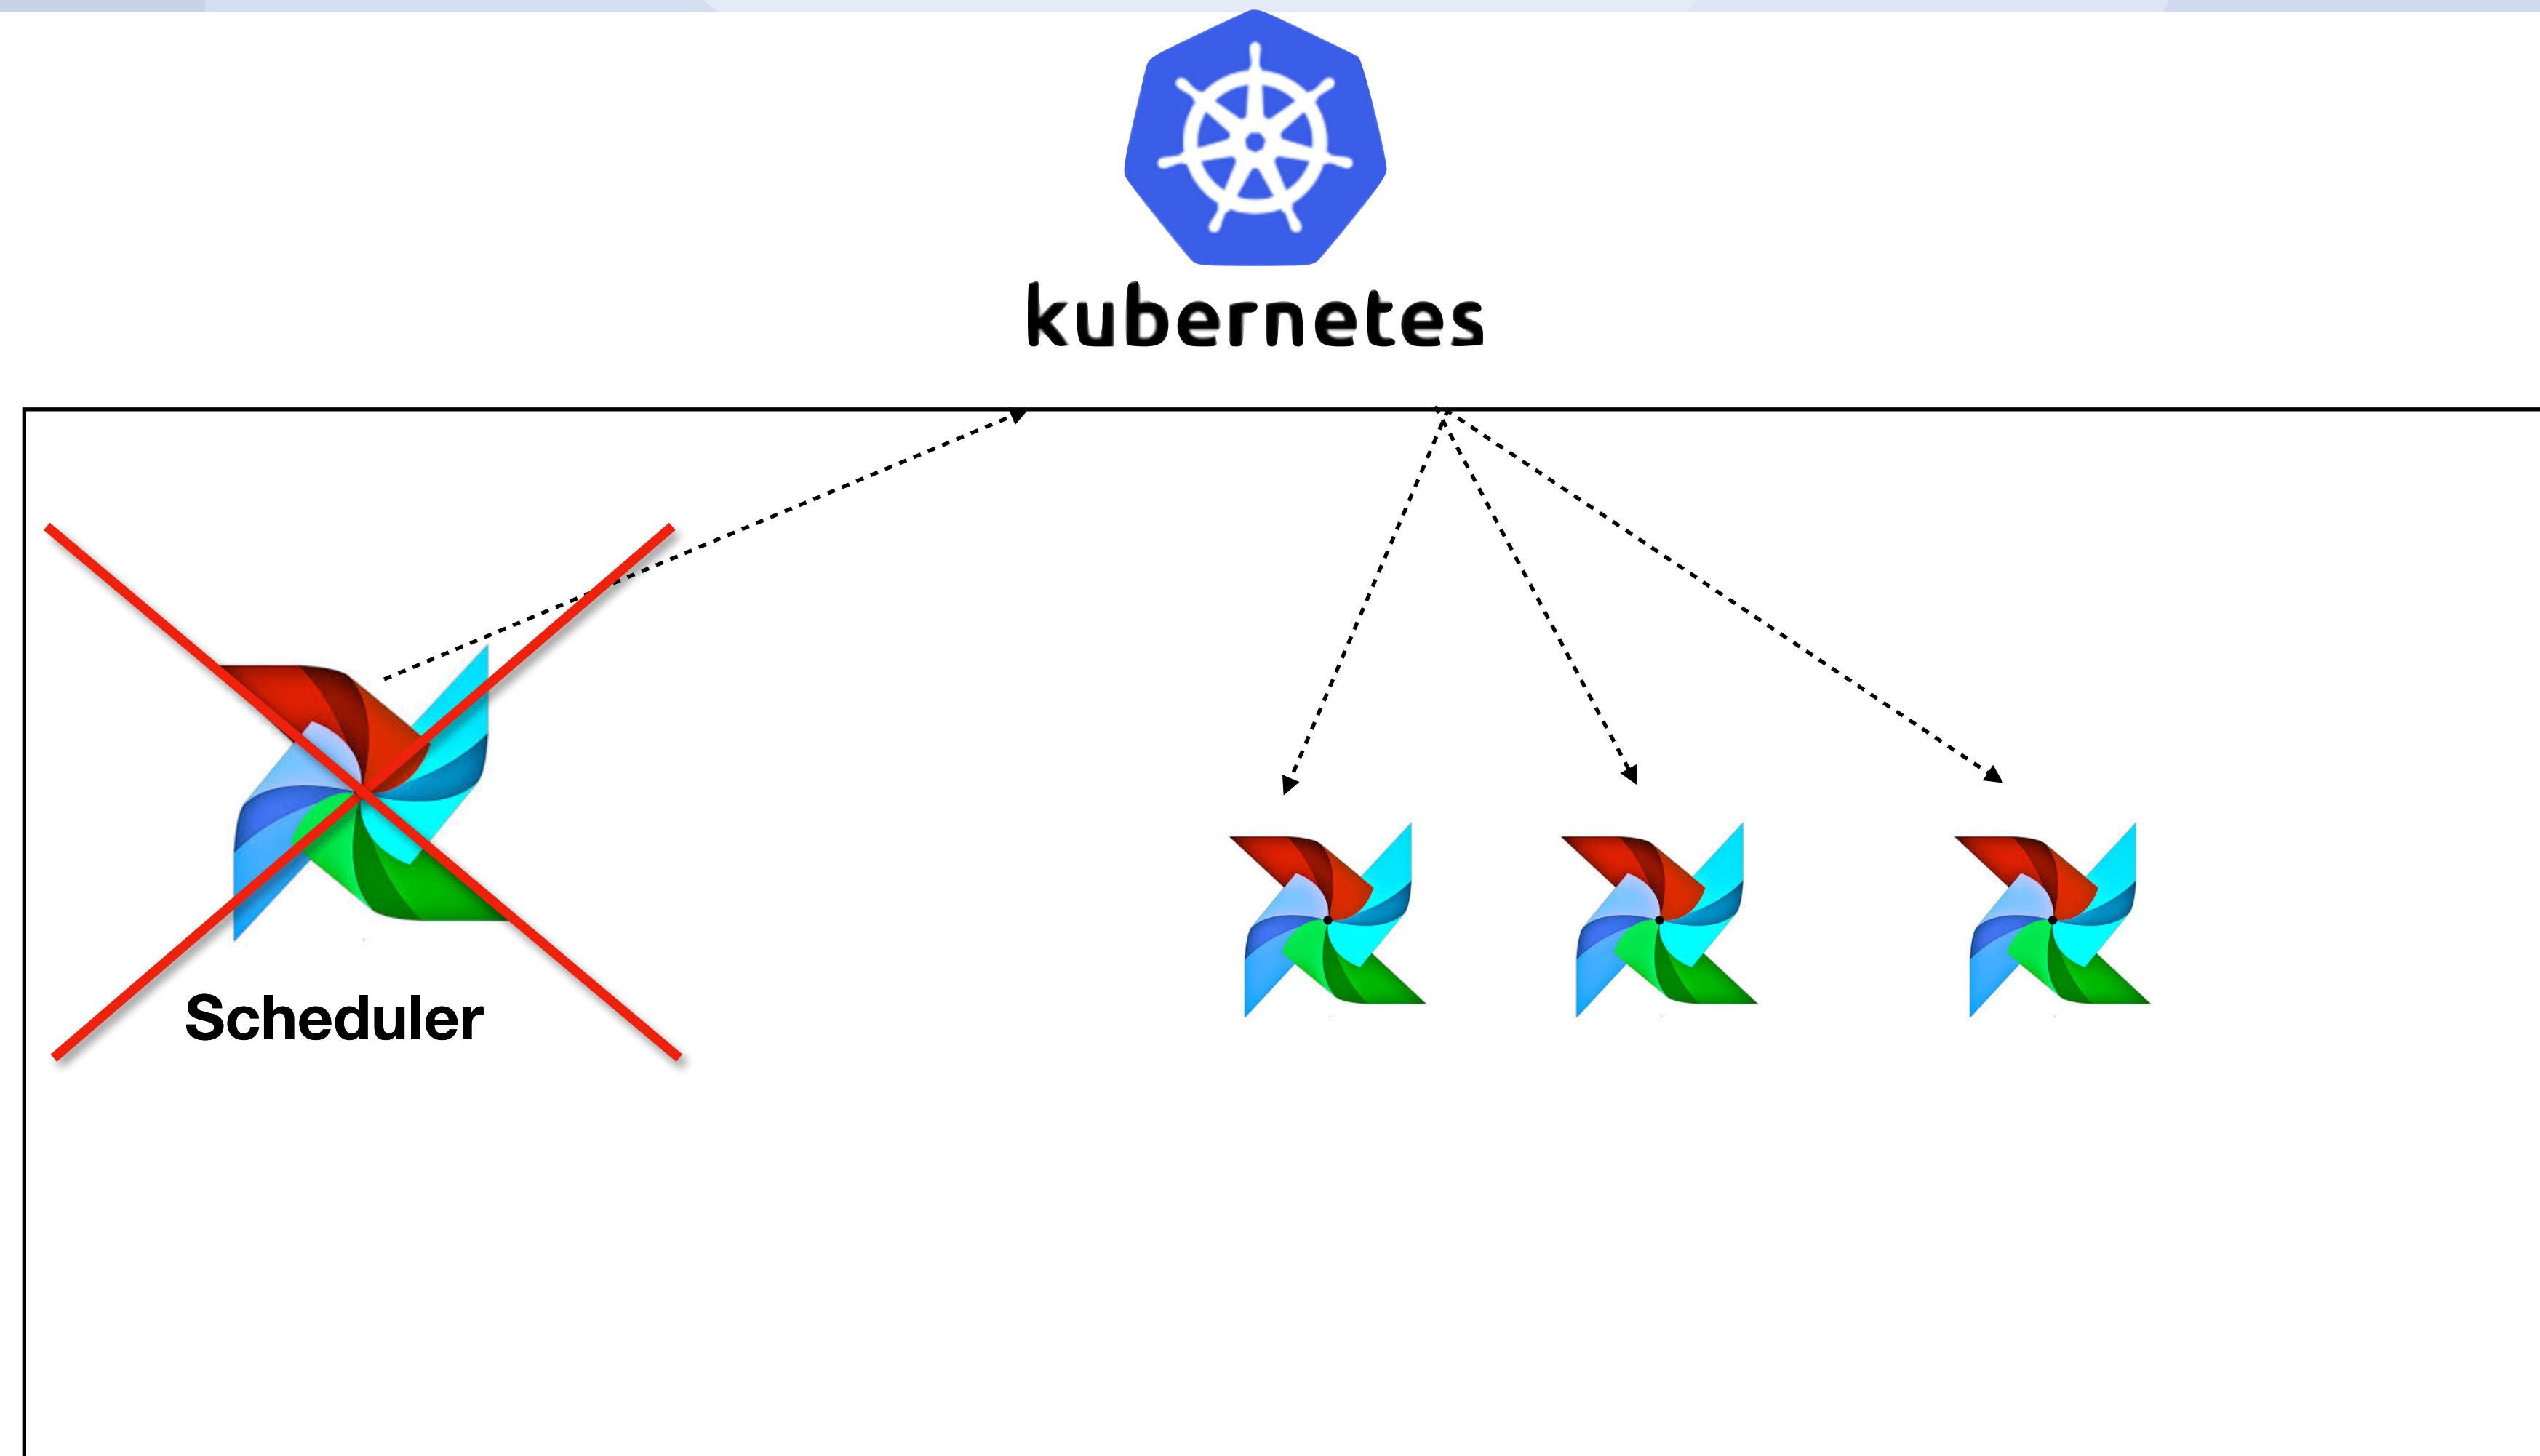
KubeCon

CloudNativeCon

North America 2018



Fault Tolerance



Fault Tolerance

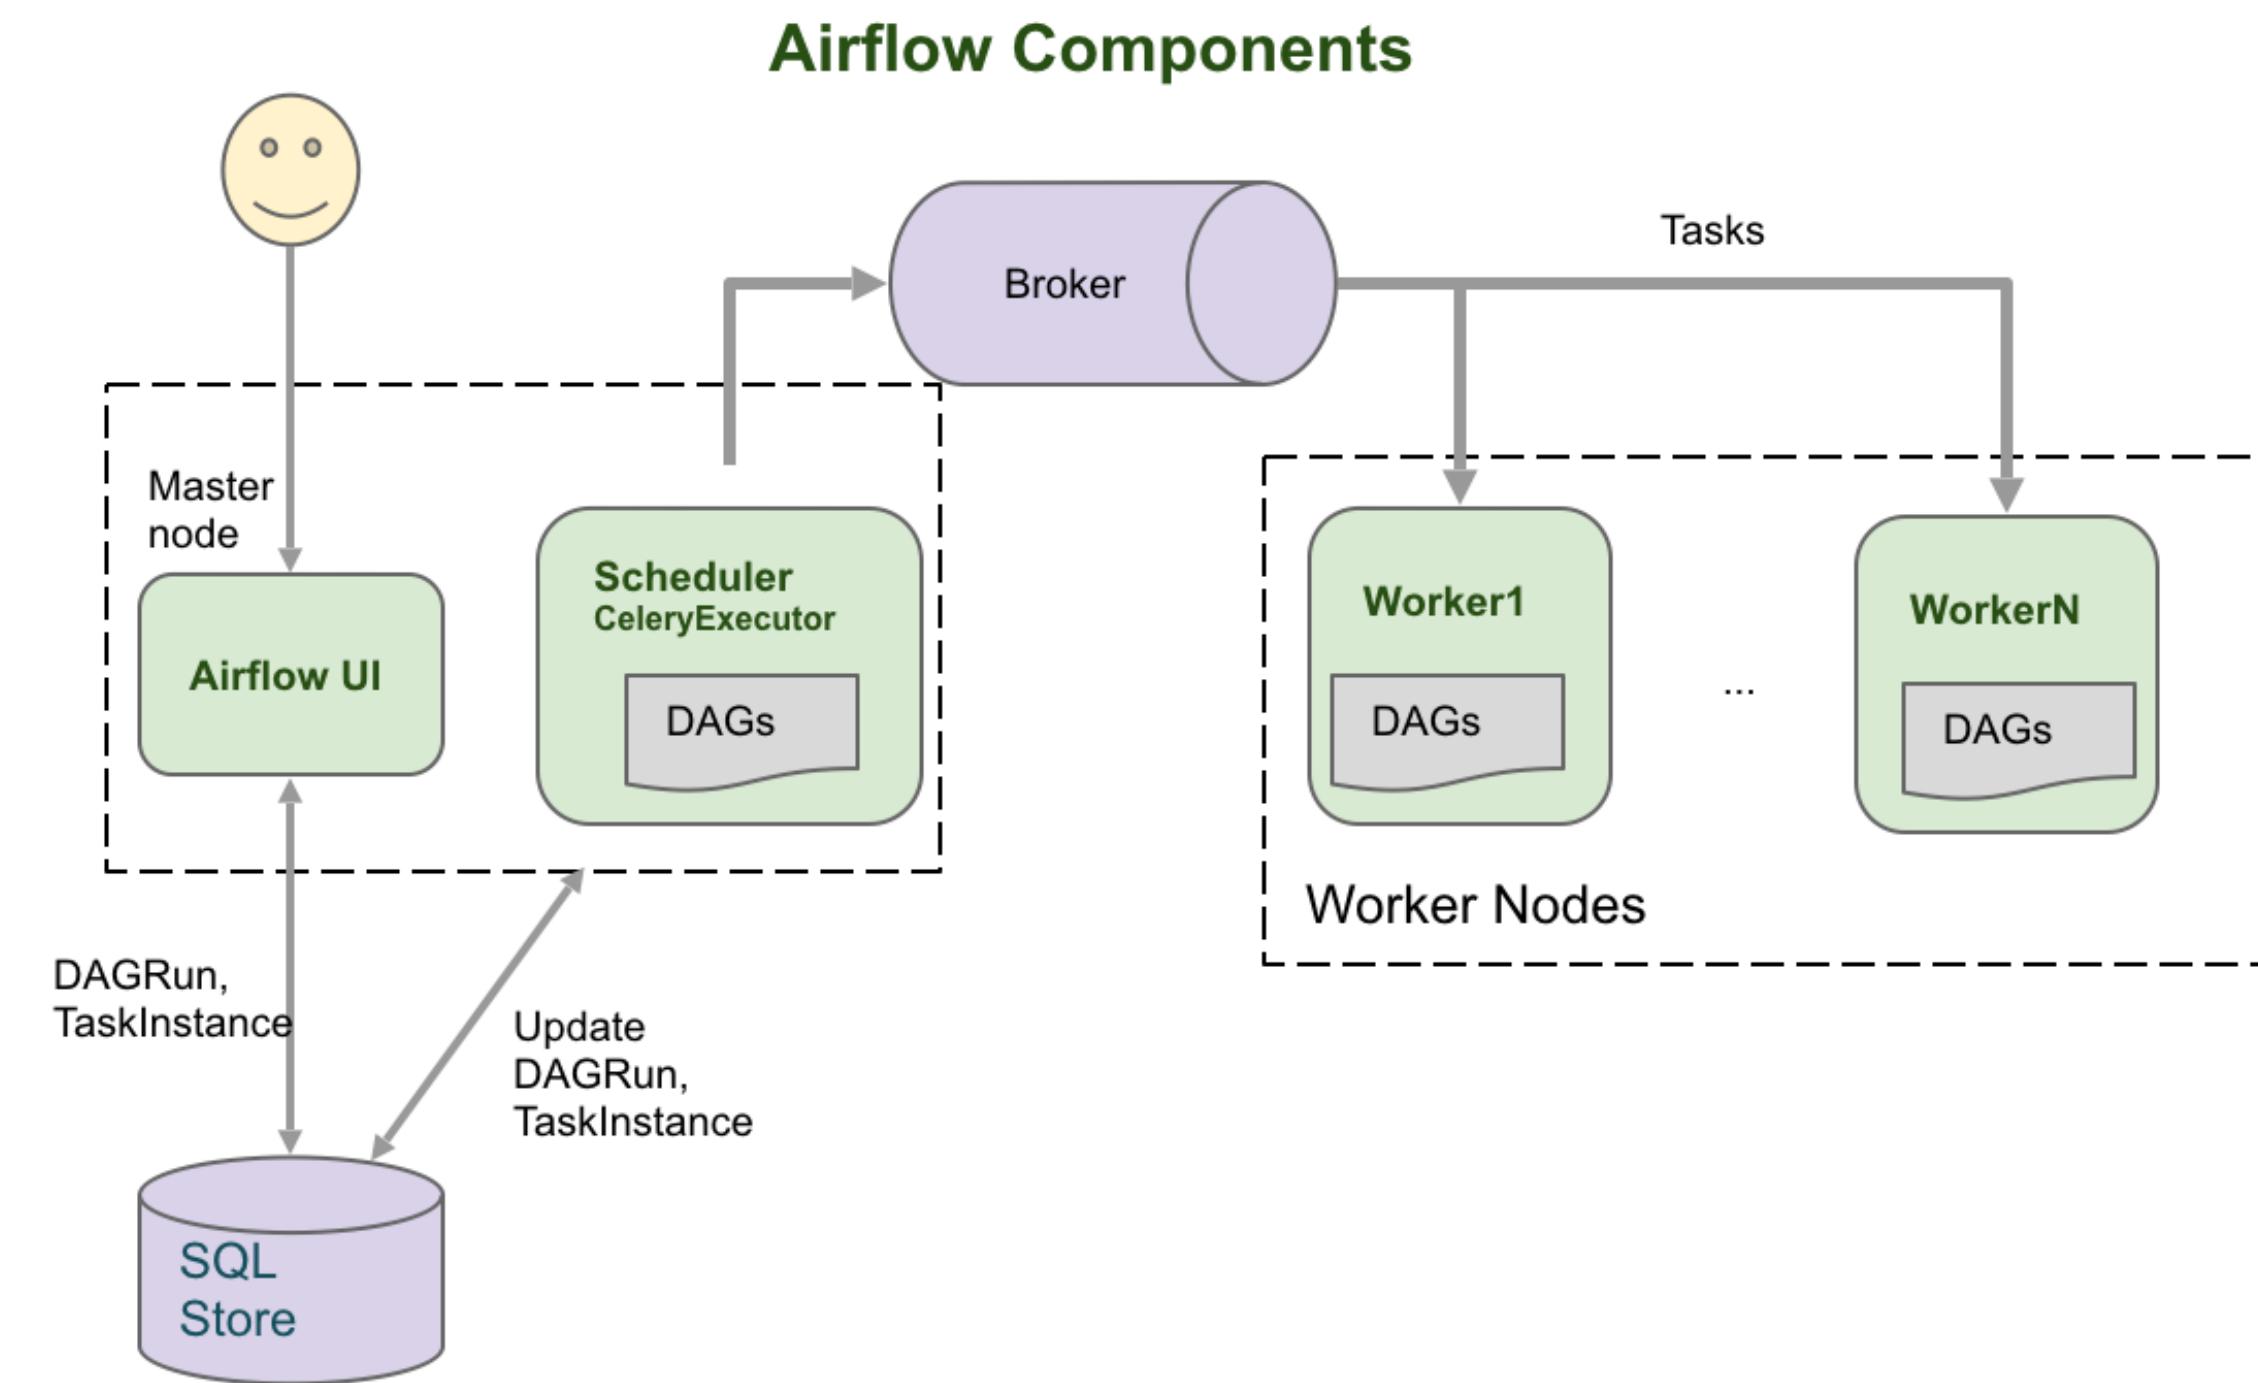
- Uses “resourceVersion” to re-create state
- Maintain a resourceVersion in SQL table for state recovery

AirflowOperator

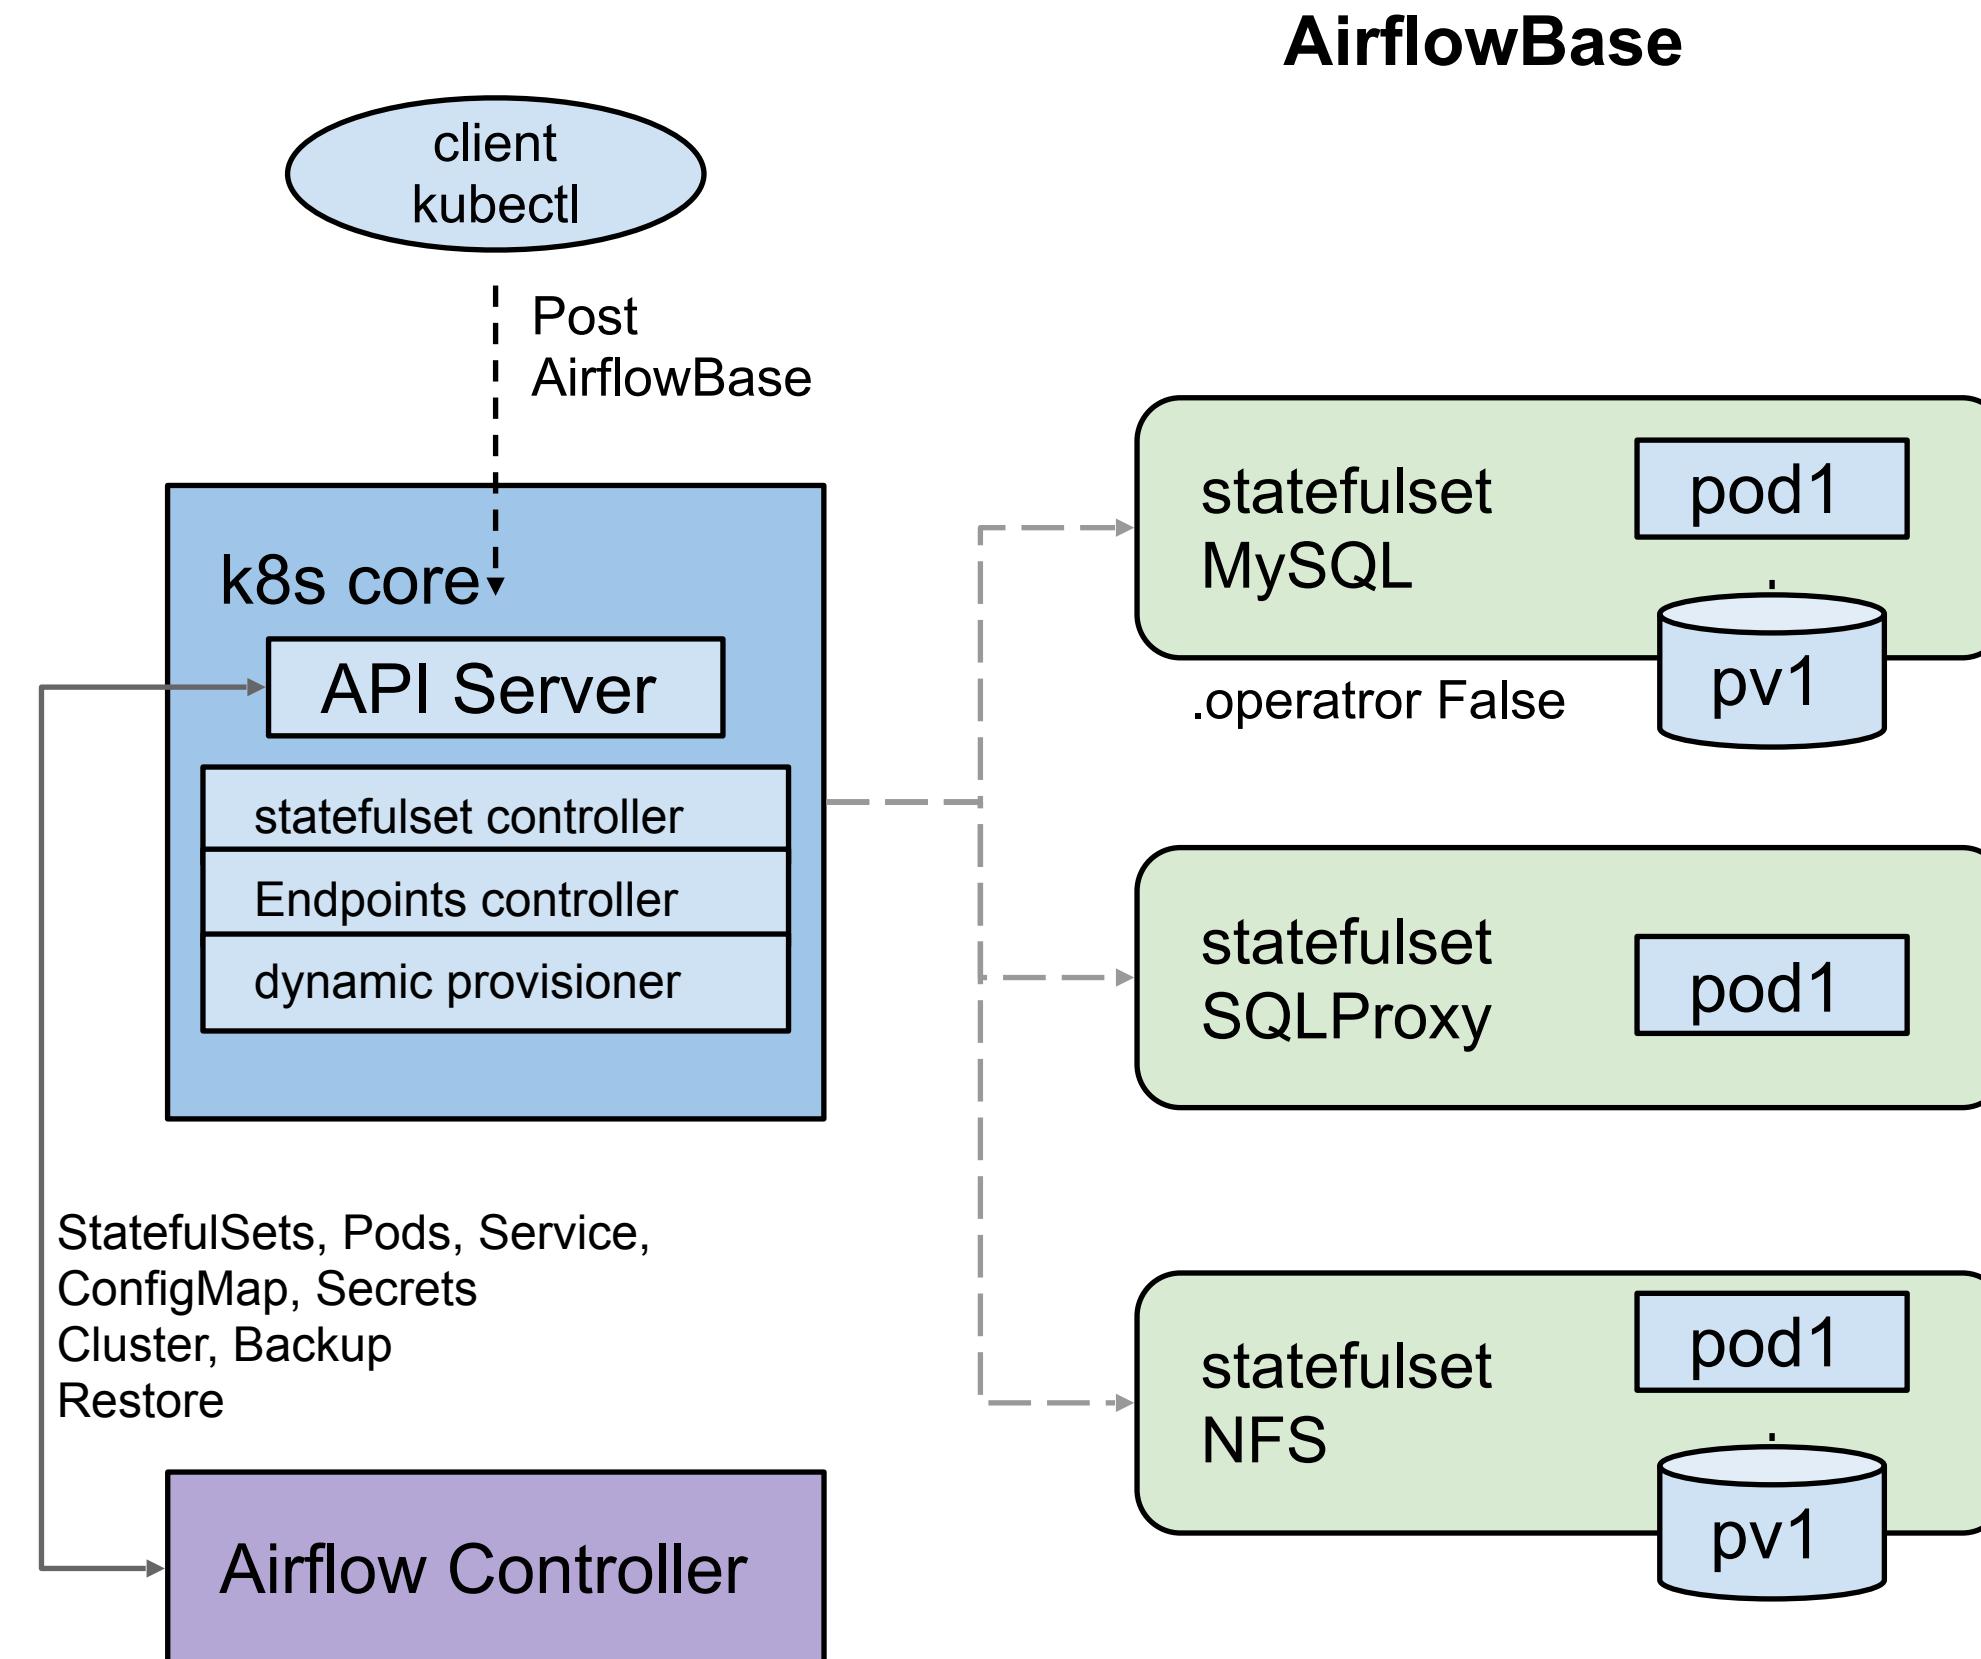
- AirflowOperator custom controller for deploying and managing Airflow
- **CRD(declarative spec) + custom controller**
- AirflowOperator watches for AirflowBase, AirflowCluster CRDs and deploys airflow components

Lower Barrier to Entry

- Airflow can involve a somewhat onerous set-up to run at scale
- Requires multiple technologies (Celery, Redis, SQL, etc)
- Have to set up alerting/monitoring for Airflow system



AirflowBase CRD



- AirflowBase CRD
 - MySQL/Postgres/SQLProxy
 - NFS
- Used by multiple Airflow Clusters

```
apiVersion: airflow.k8s.io/v1alpha1
kind: AirflowBase
metadata:
  name: mc-base
spec:
  mysql:
    operator: False
  storage:
    version: ""
```

```
apiVersion: airflow.k8s.io/v1alpha1
kind: AirflowBase
metadata:
  name: ck-base
spec:
  sqlproxy:
    project: someproject
    region: us-central1
    instance: testsql-cluster
  storage:
    version: ""
```

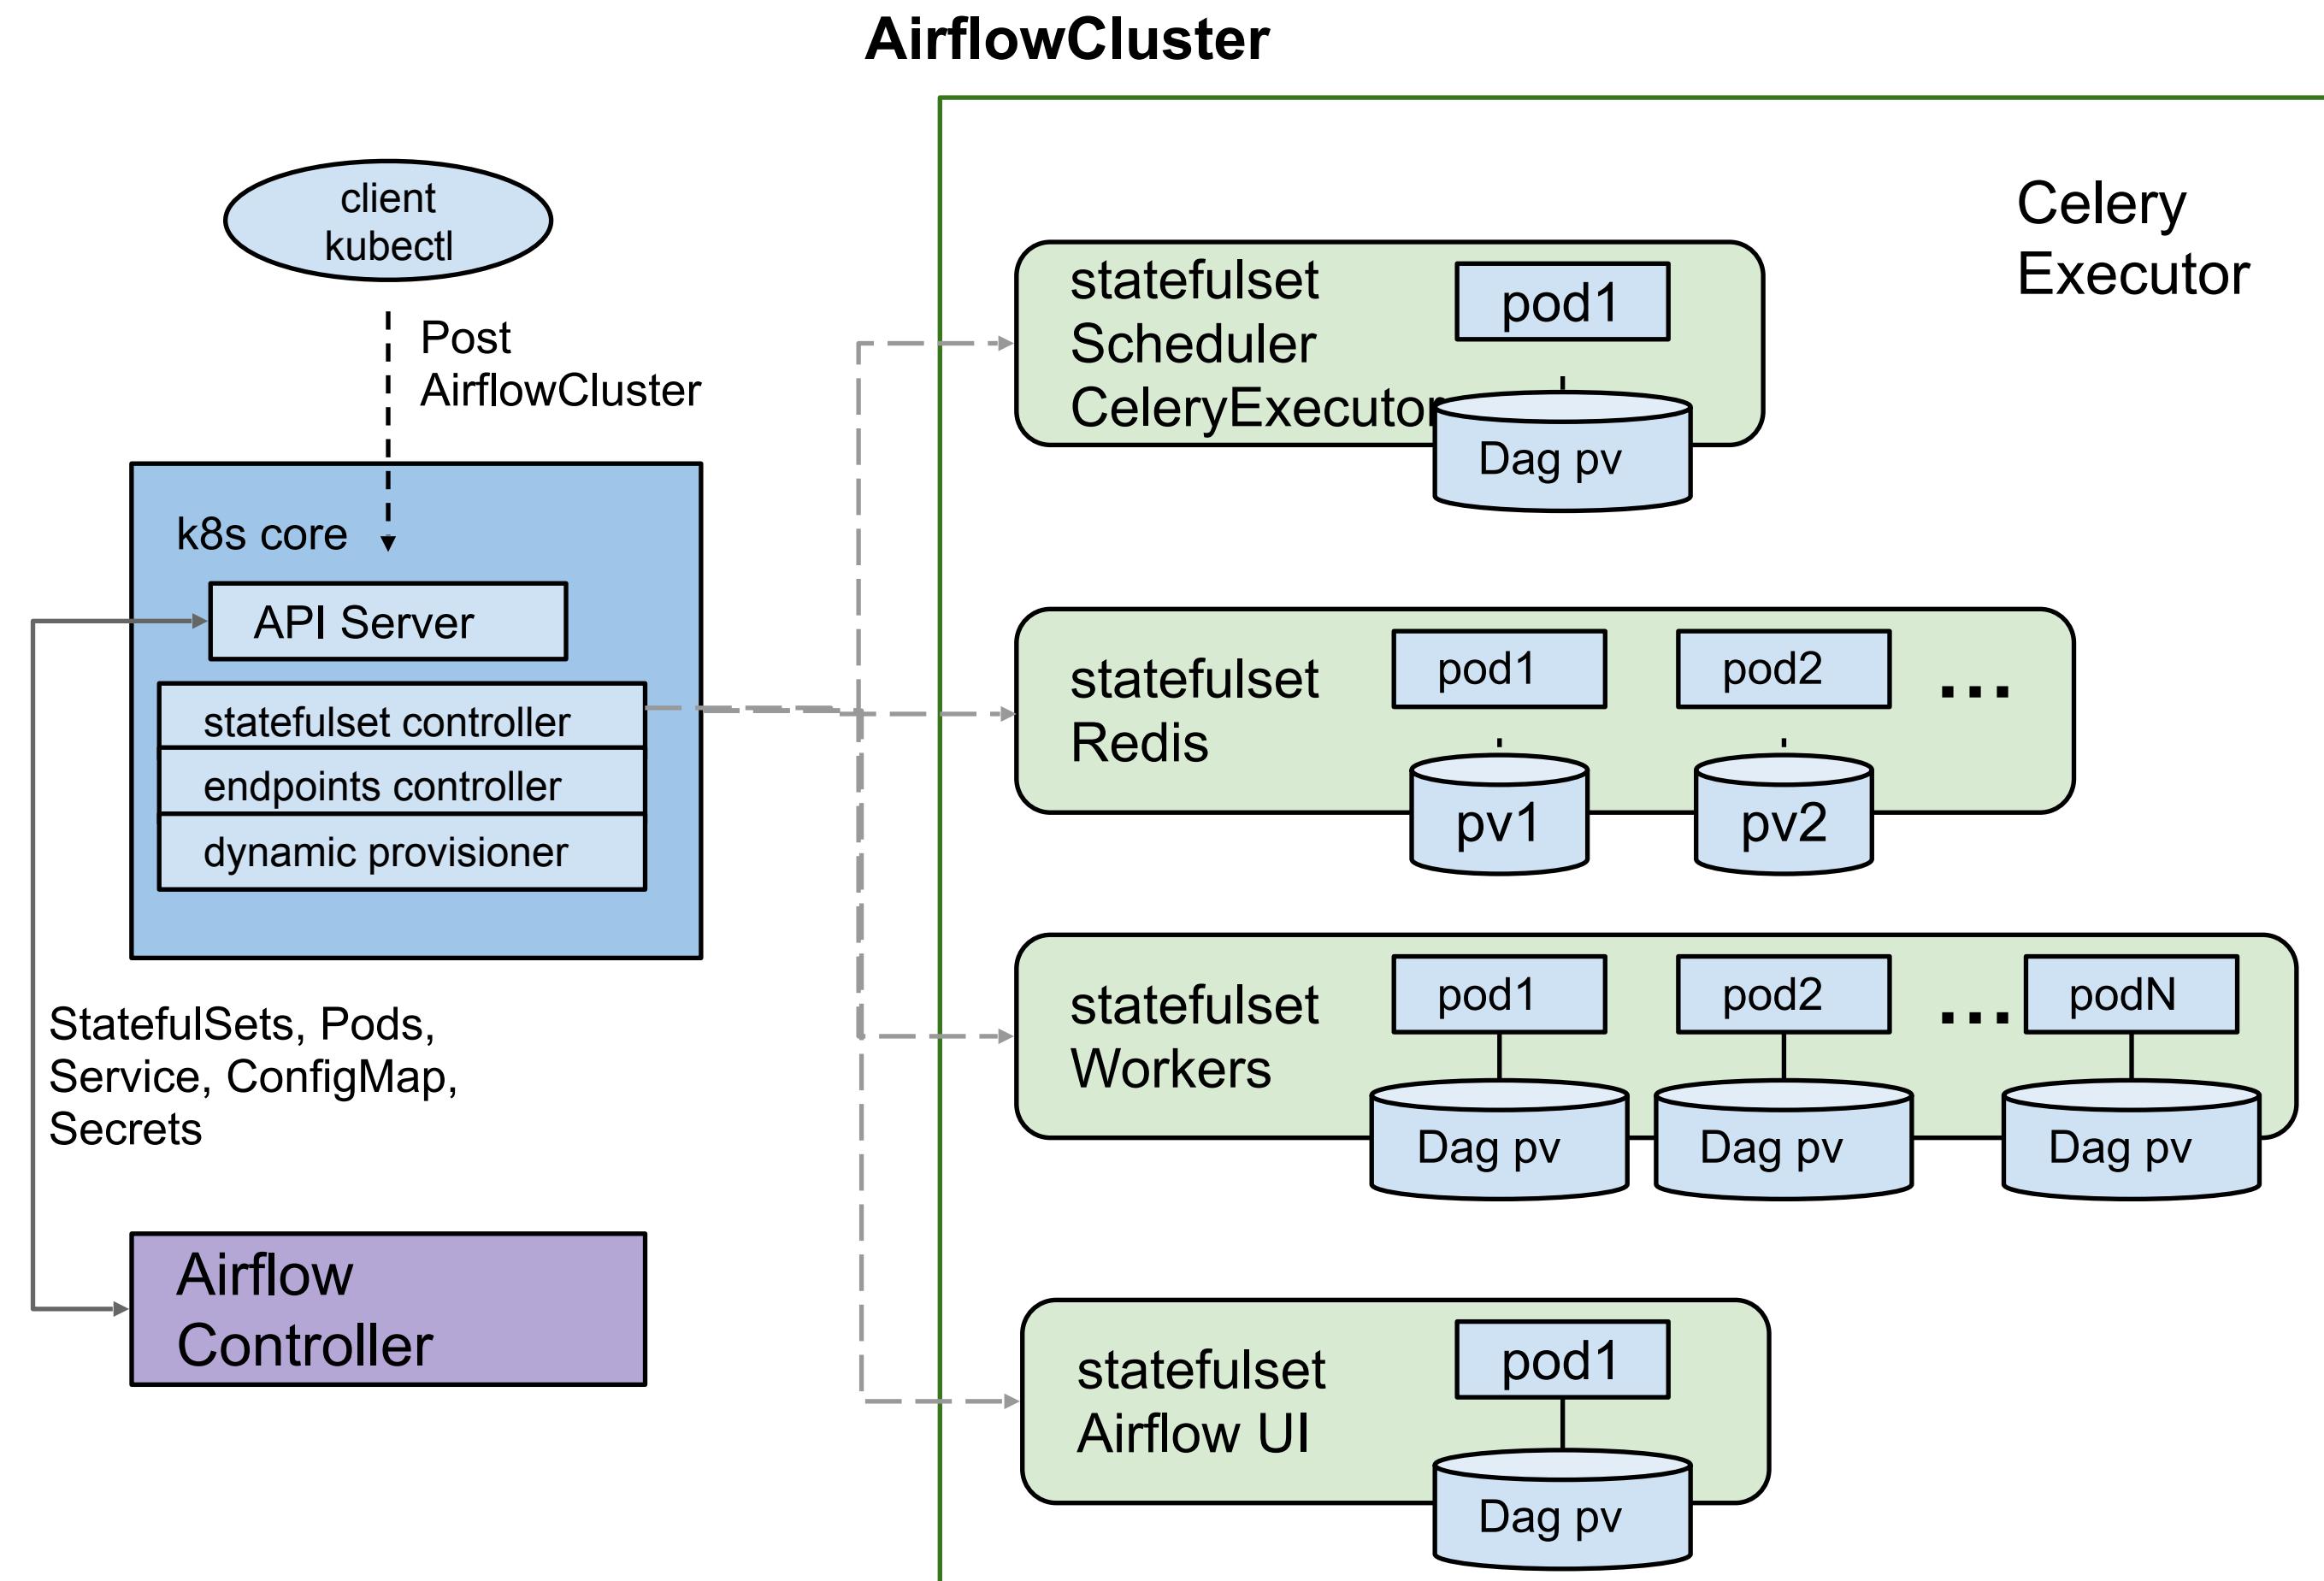
AirflowCluster CRD



KubeCon

CloudNativeCon

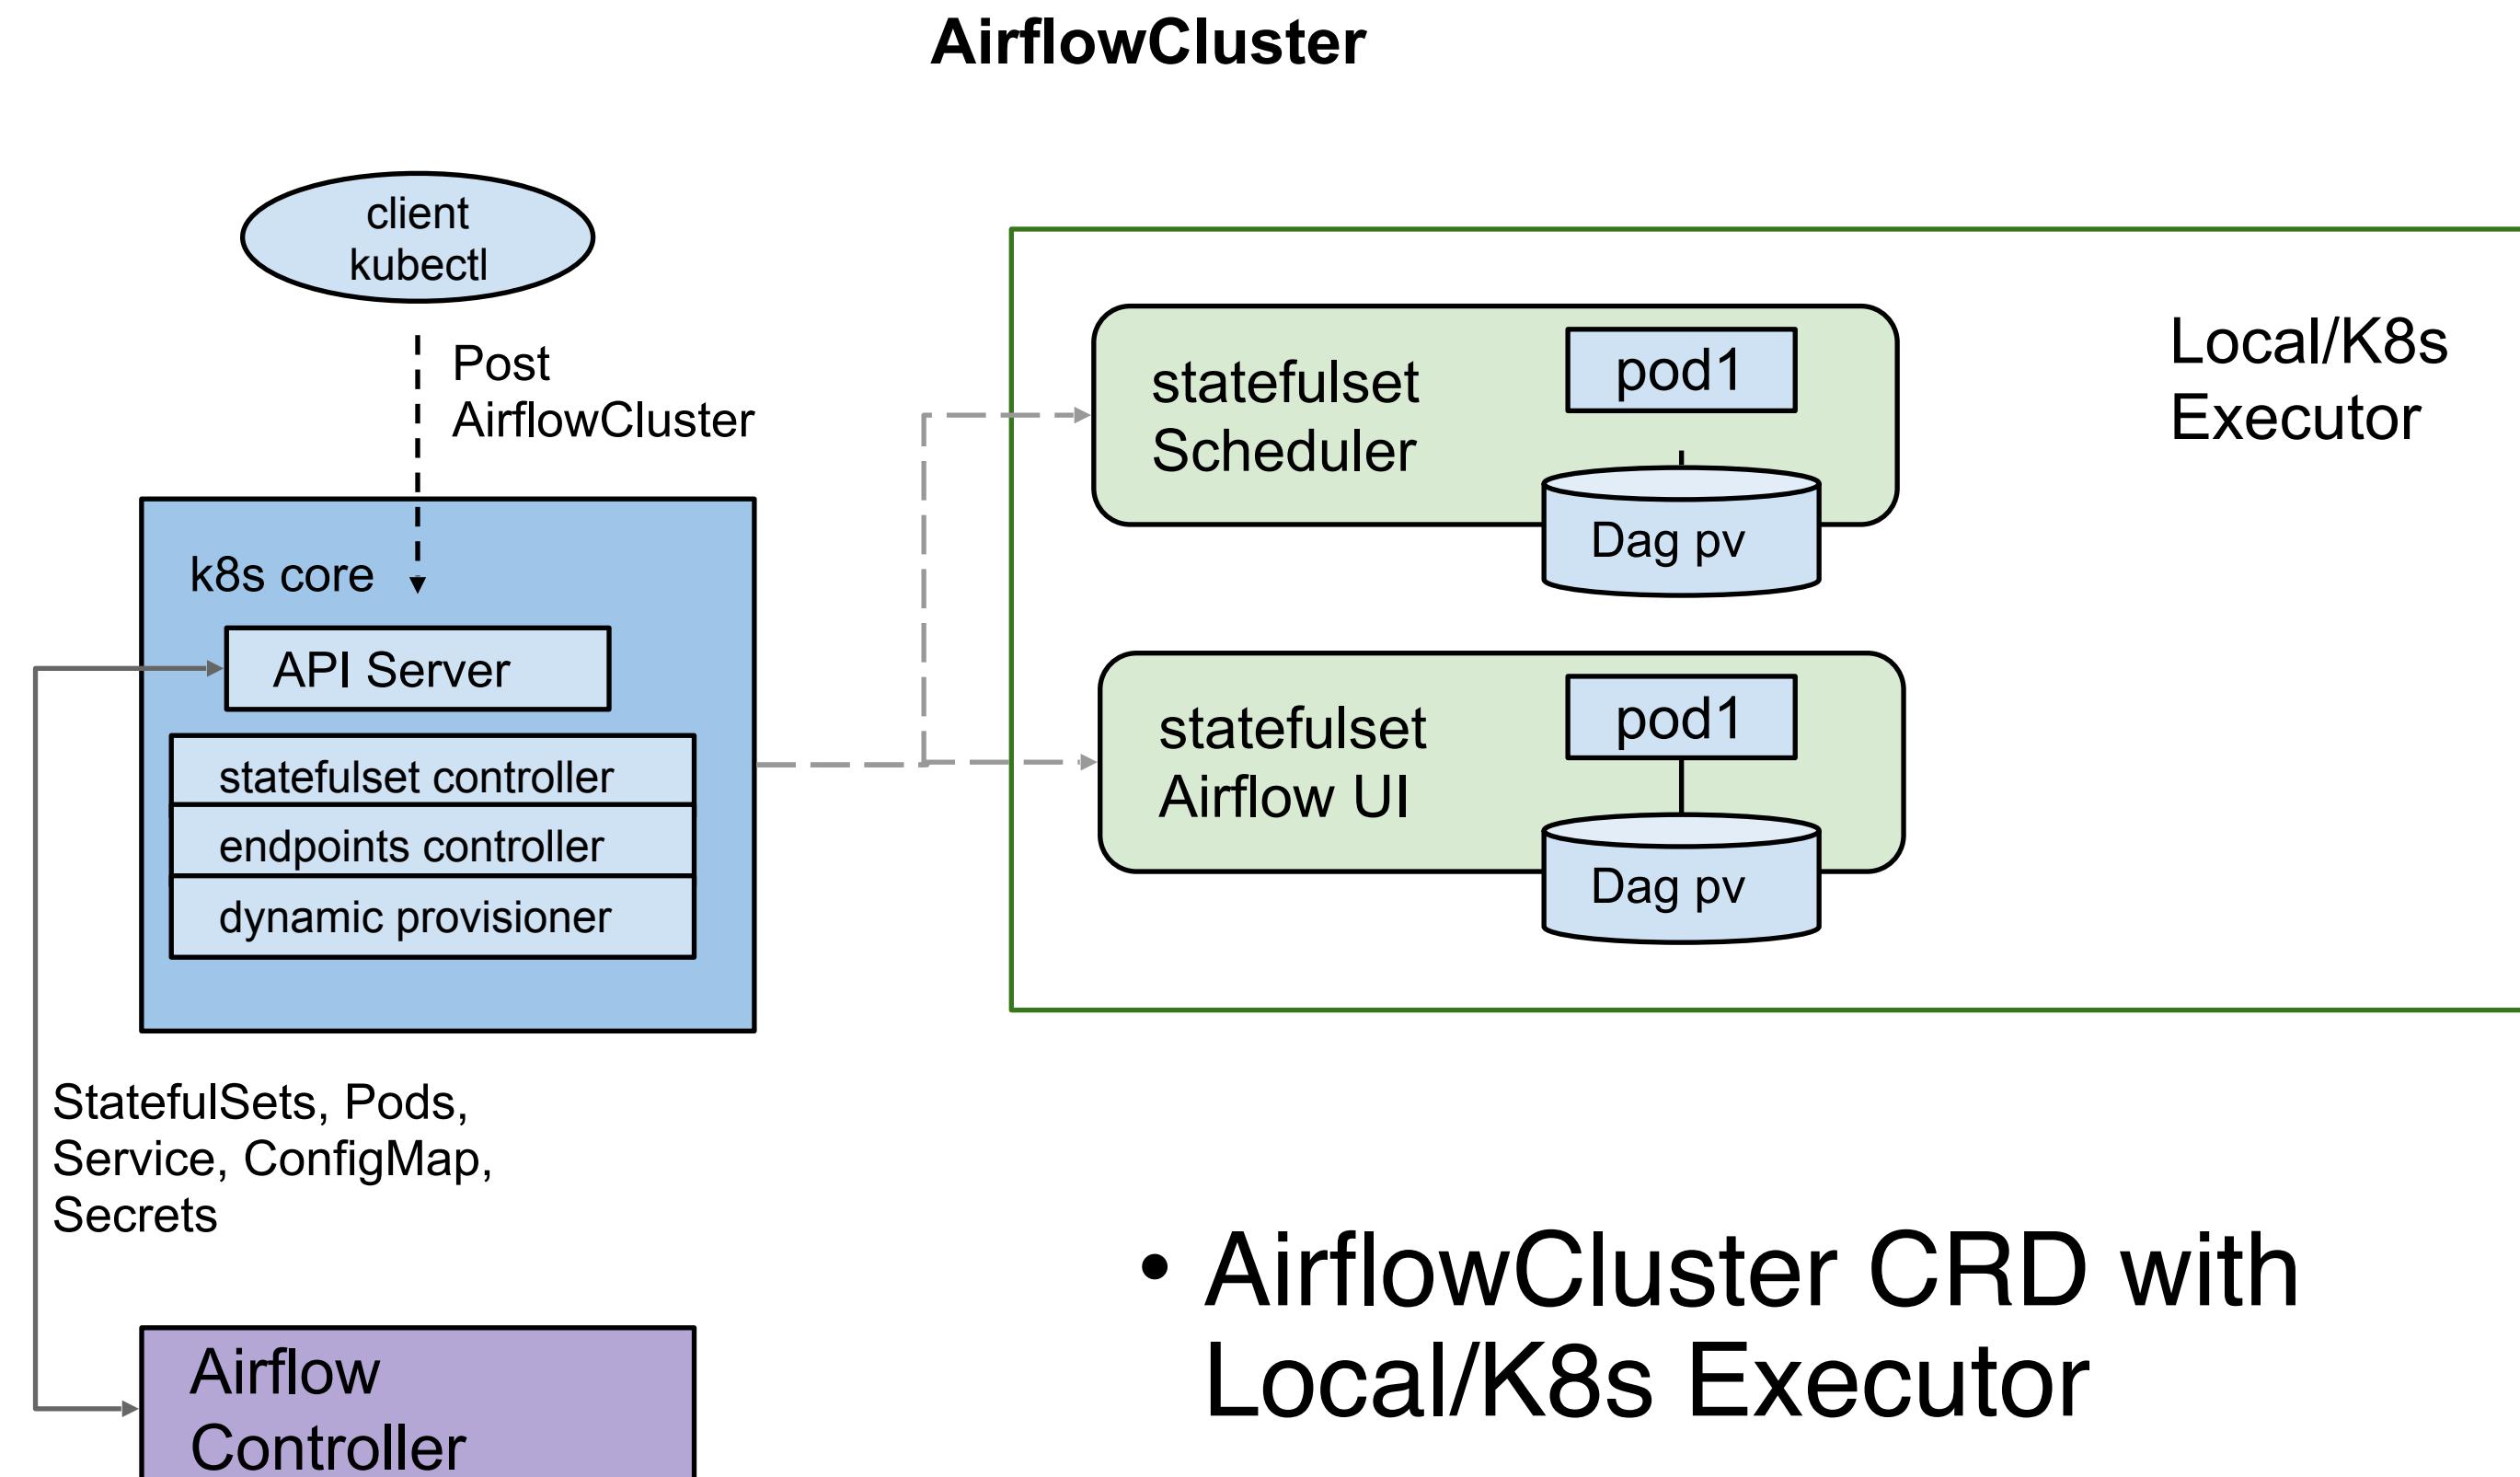
North America 2018



- Celery Executor
 - Redis
 - Airflow UI
 - Airflow Scheduler
 - Airflow Workers
- Each cluster gets its own unique SQL connection string (user:password/dB).

```
apiVersion: airflow.k8s.io/v1alpha1
kind: AirflowCluster
metadata:
  name: mc-cluster
spec:
  executor: Celery
  config:
    airflow:
      AIRFLOW_SOME_CONFIG: SomeValue
    redis:
      operator: False
    scheduler:
      version: "1.10.1"
    ui:
      replicas: 1
      version: "1.10.1"
    worker:
      replicas: 2
      version: "1.10.1"
    flower:
      replicas: 1
      version: "1.10.1"
    dags:
      subdir: "airflow/example_dags/"
      git:
        repo: "https://github.com/apache/incubator-airflow/"
        once: true
    airflowbase:
      name: mc-base
```

AirflowCluster CRD

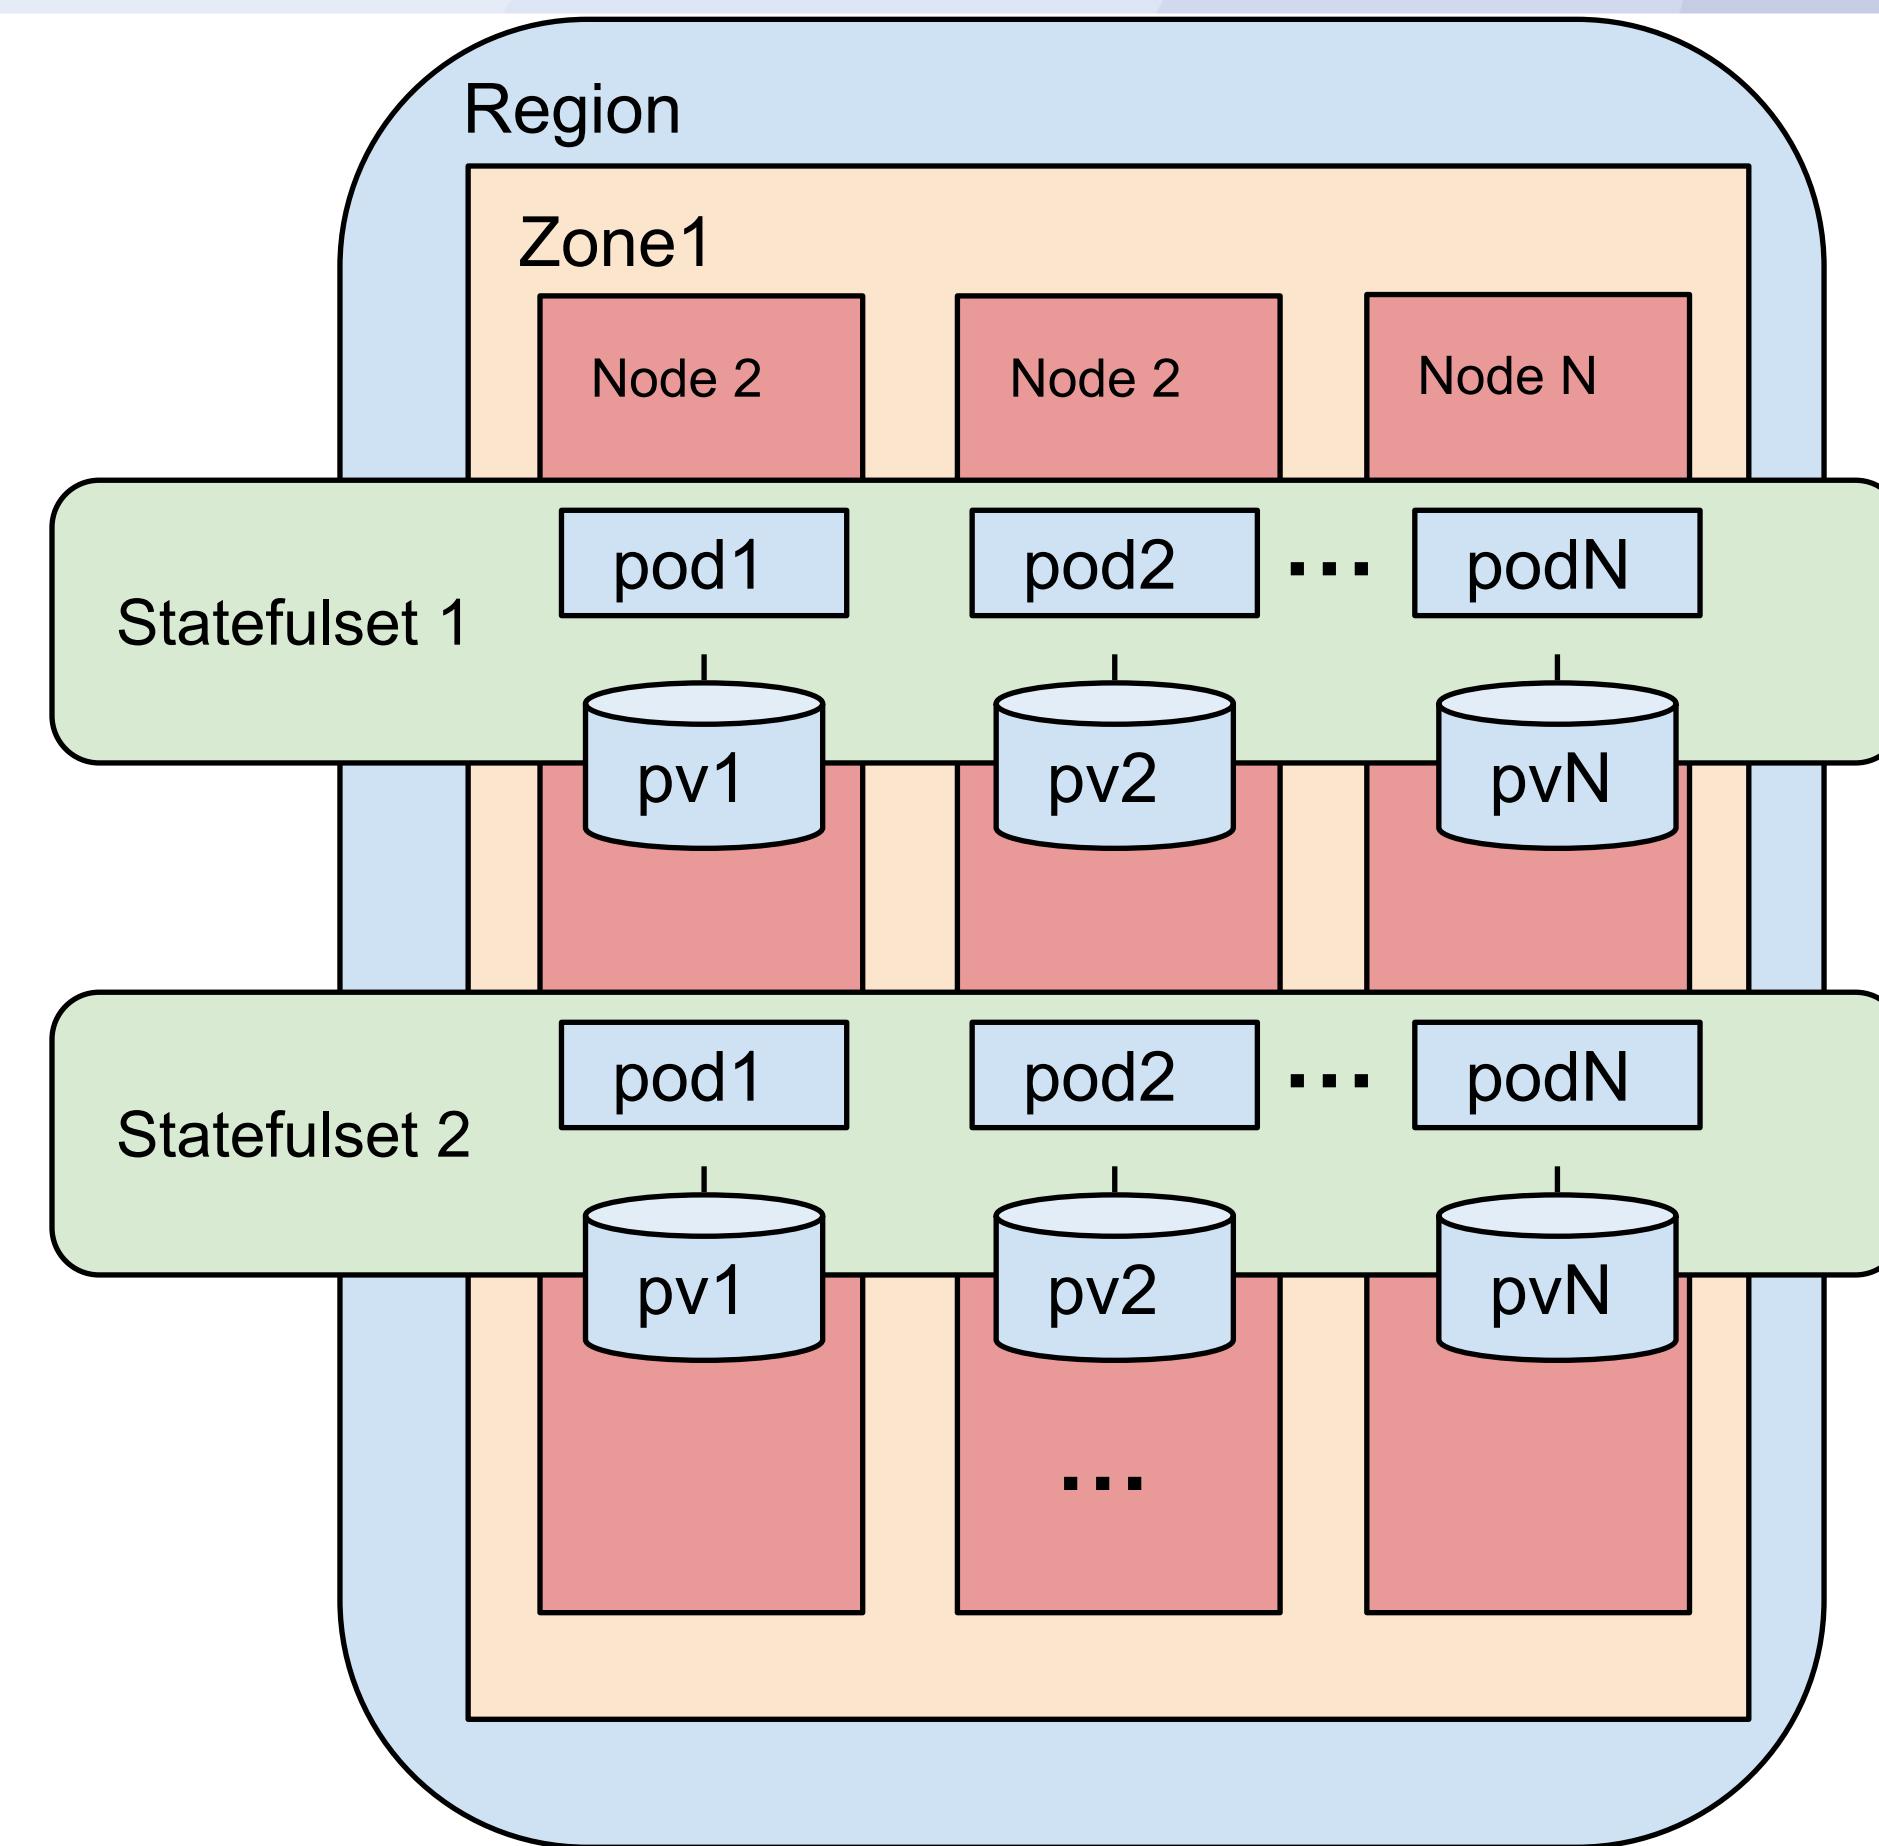


- AirflowCluster CRD with Local/K8s Executor
 - Airflow UI
 - Airflow Scheduler

```
apiVersion: airflow.k8s.io/v1alpha1
kind: AirflowCluster
metadata:
  name: mk-cluster
spec:
  executor: Kubernetes
  ui:
    replicas: 1
    version: "1.10.1"
  scheduler:
    version: "1.10.1"
  worker:
    version: "1.10.1"
  dags:
    subdir: "airflow/example_dags/"
    git:
      repo: "https://github.com/apache/incubator-airflow/"
      once: true
      branch: master
  airflowbase:
    name: mc-base
```

AirflowOperator

- `cluster.Spec.Affinity.*.topology` can be set to “`kubernetes.io/hostname`” to spread Pods across Nodes within a Zone.
- Limit the effect of node failures within a zone



Pods spread across Nodes in Zone

Monitoring

- Can use existing Kubernetes infrastructure
- Only needed to think about Airflow, not machines

Prometheus

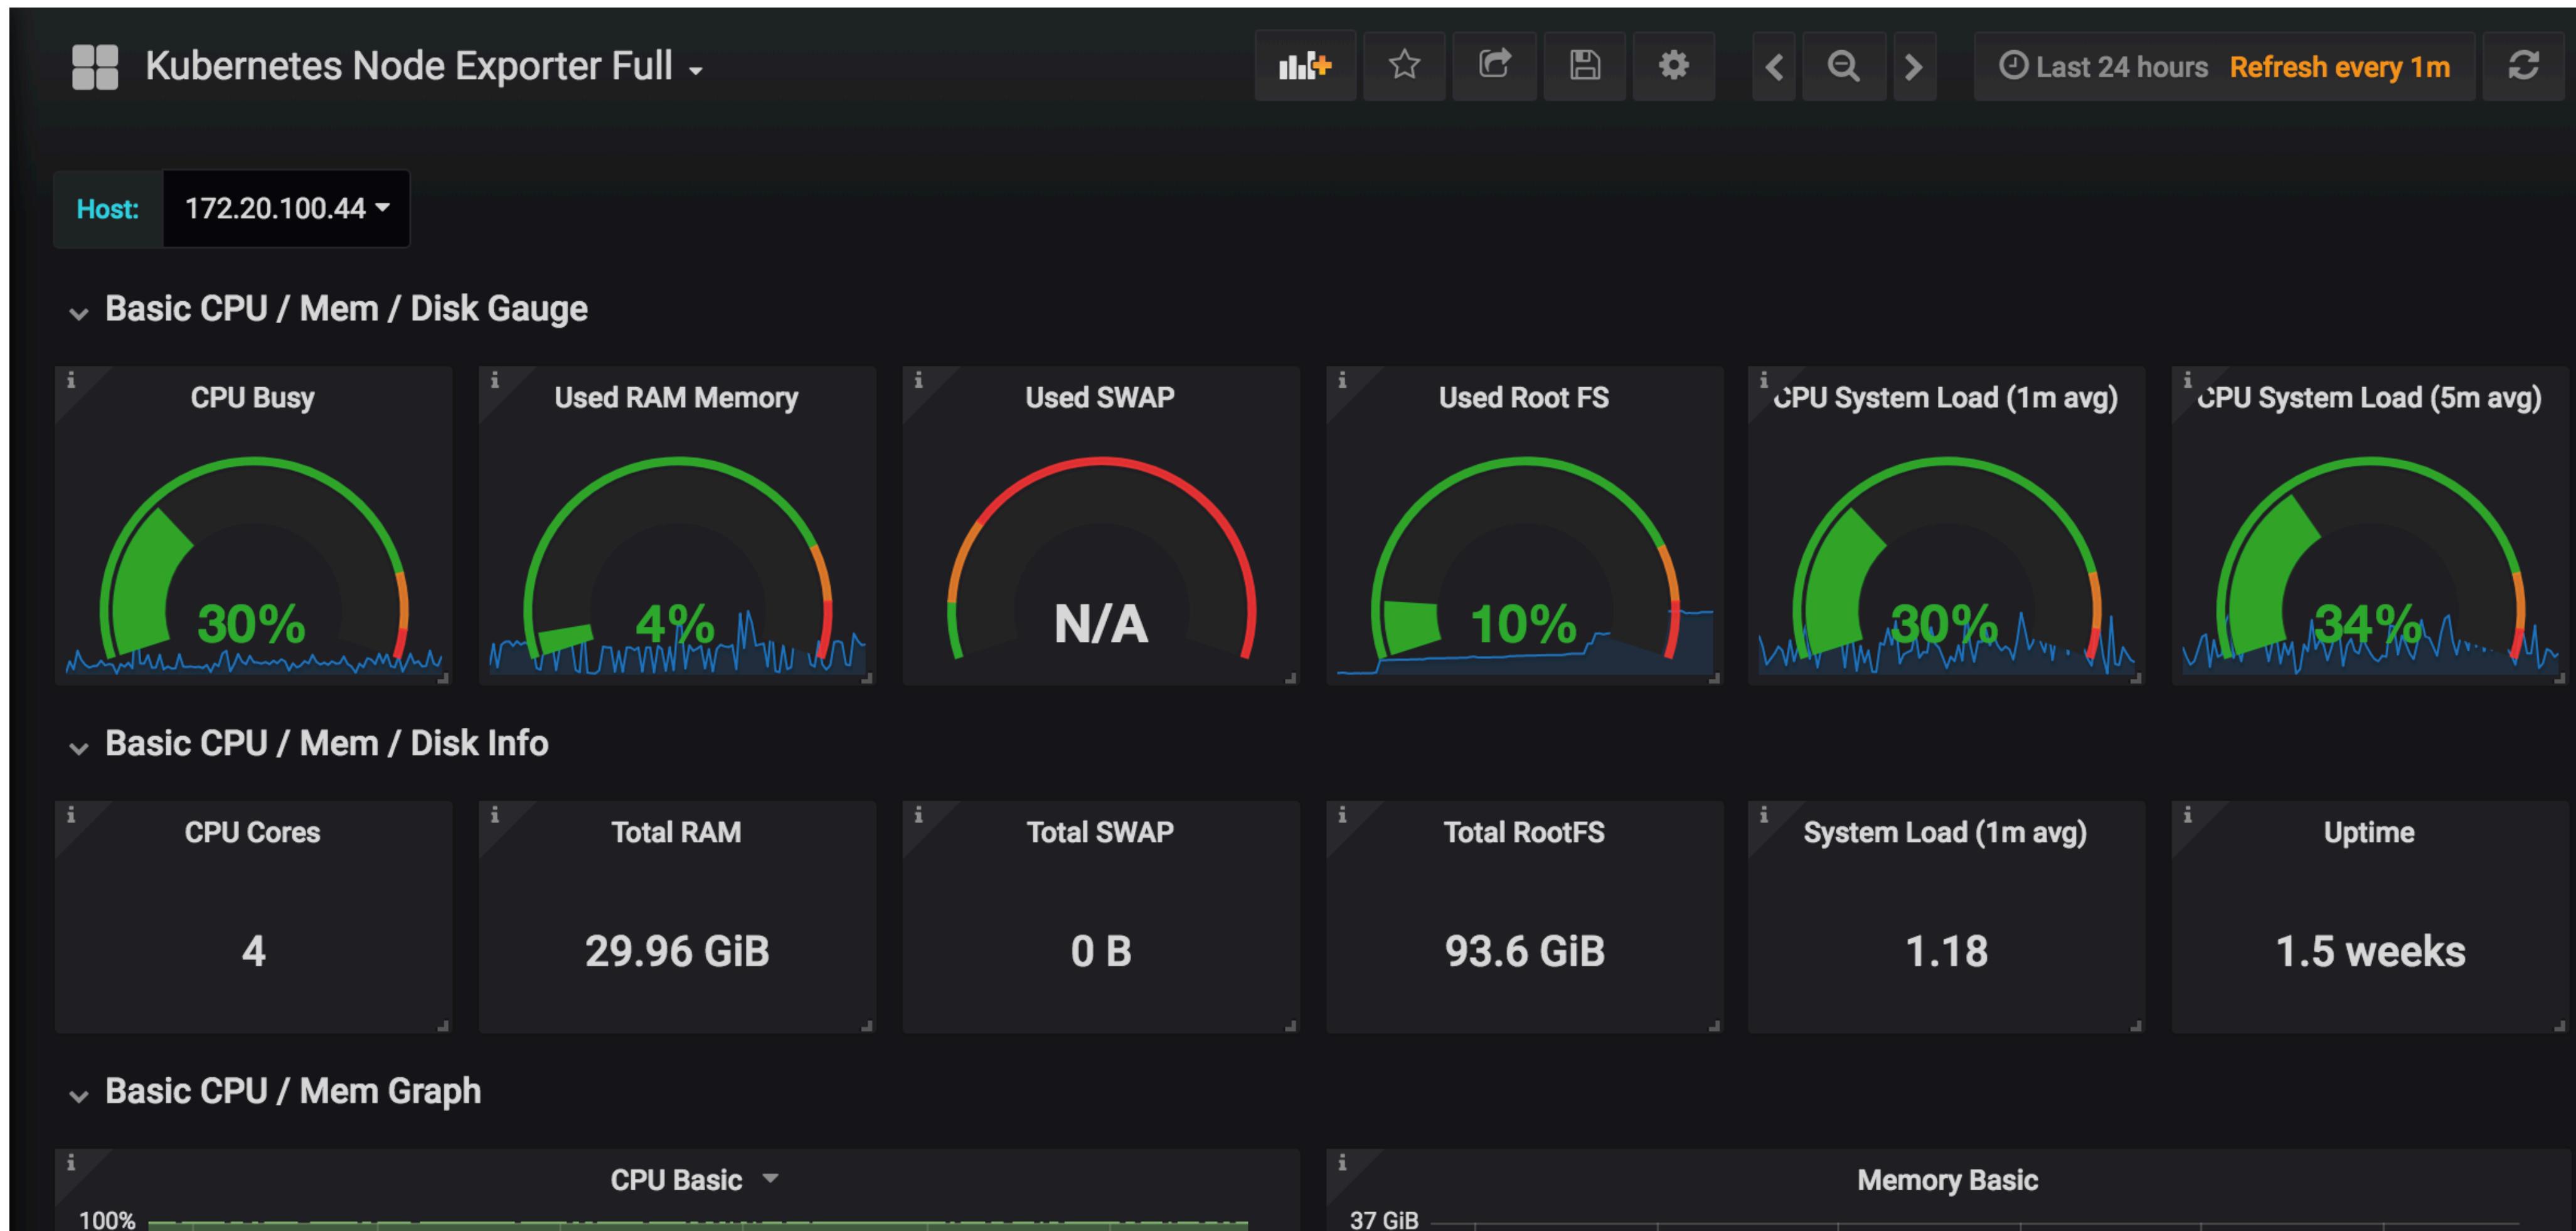


KubeCon



CloudNativeCon

North America 2018



Elasticsearch

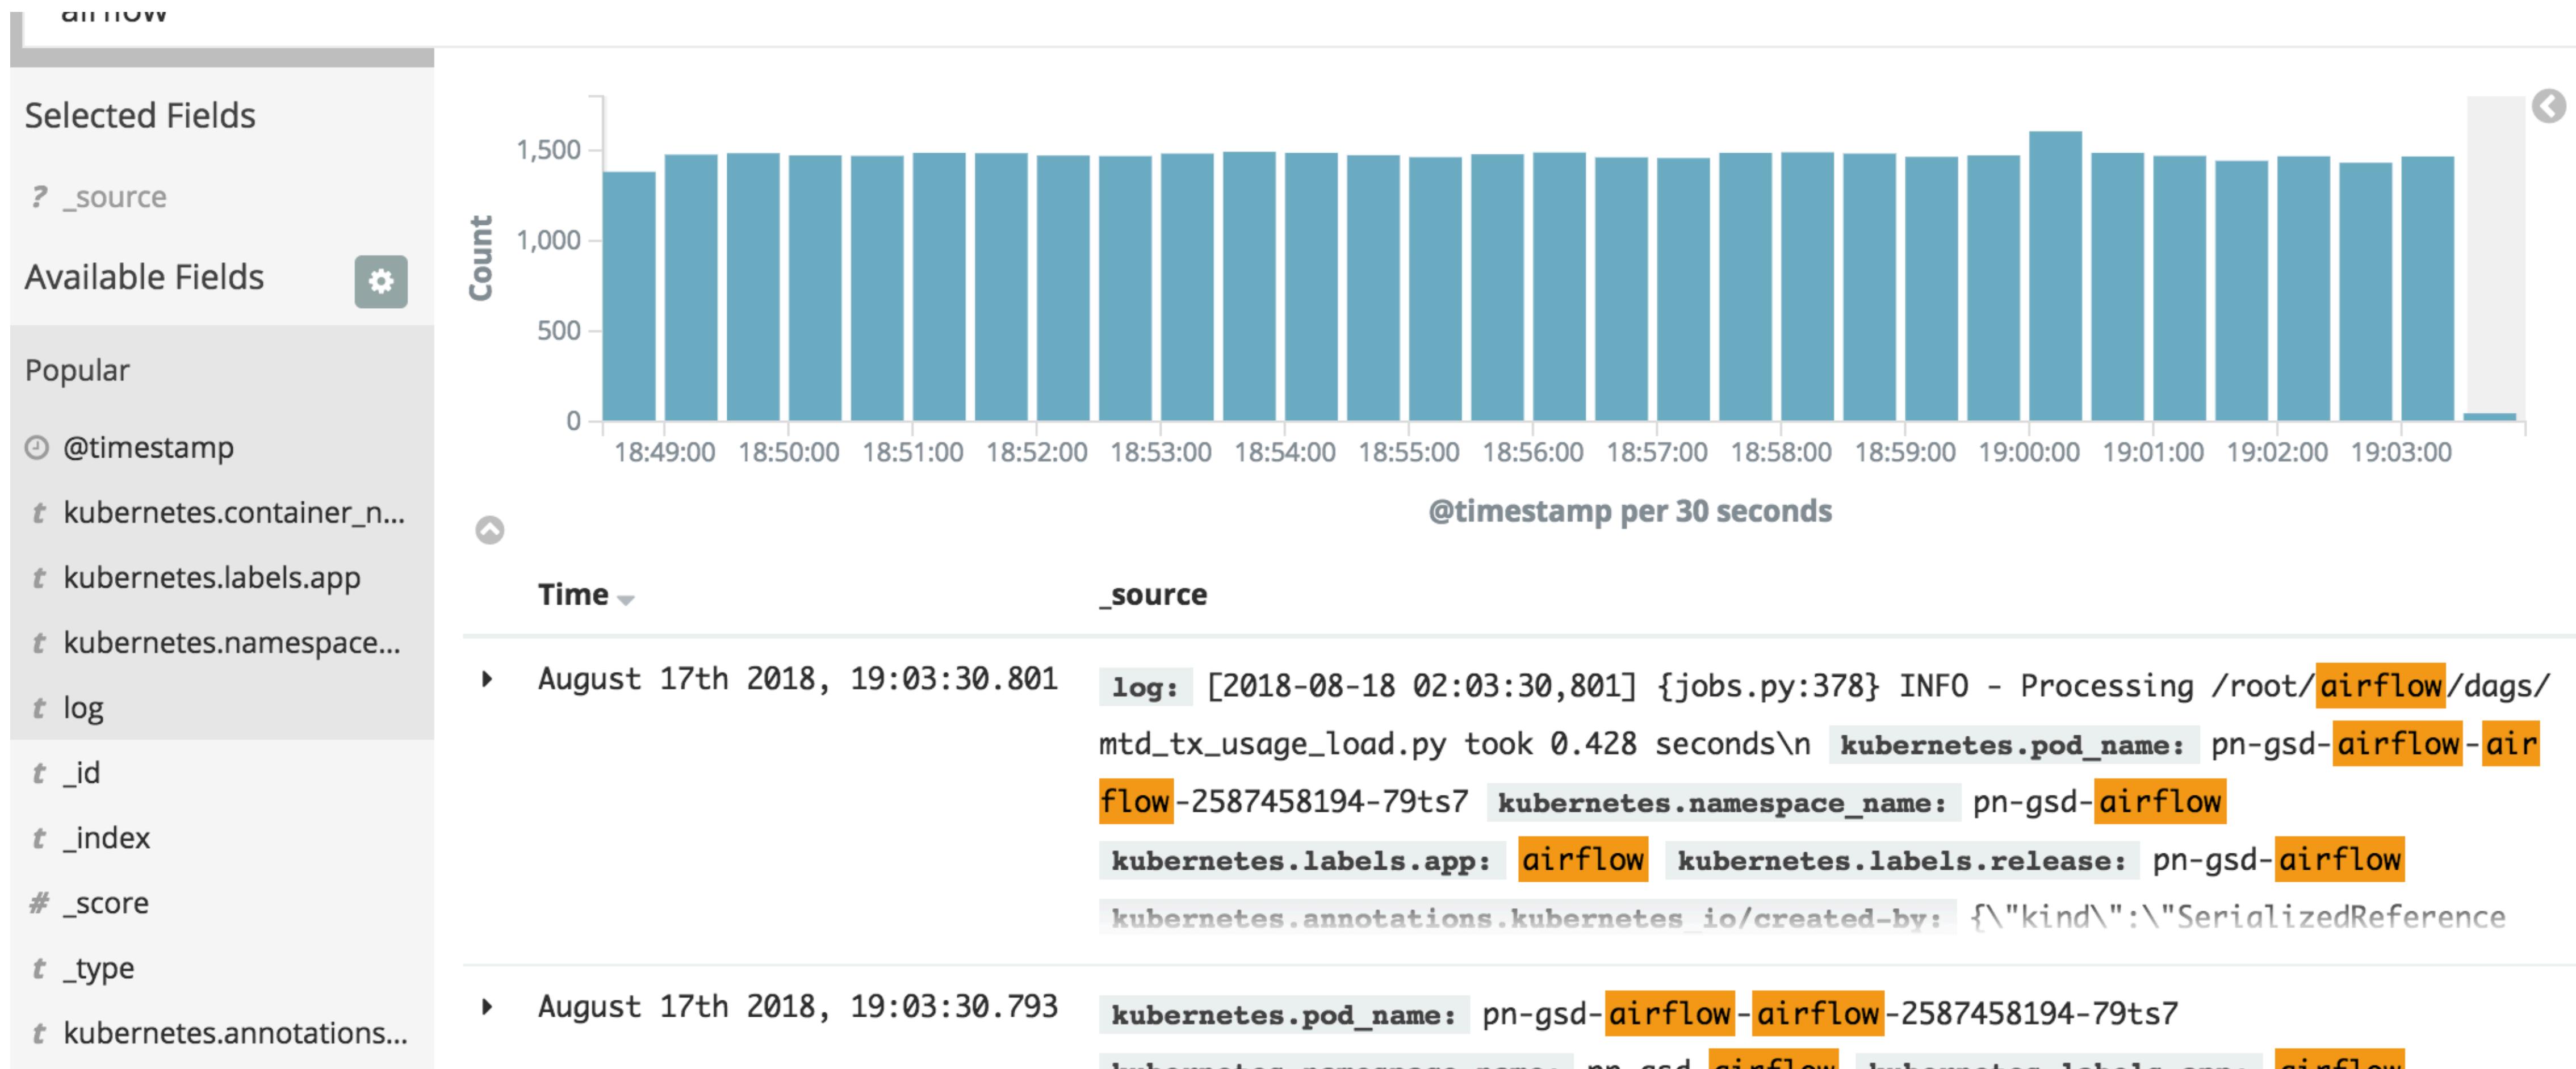


KubeCon



CloudNativeCon

North America 2018



Airflow Dashboard

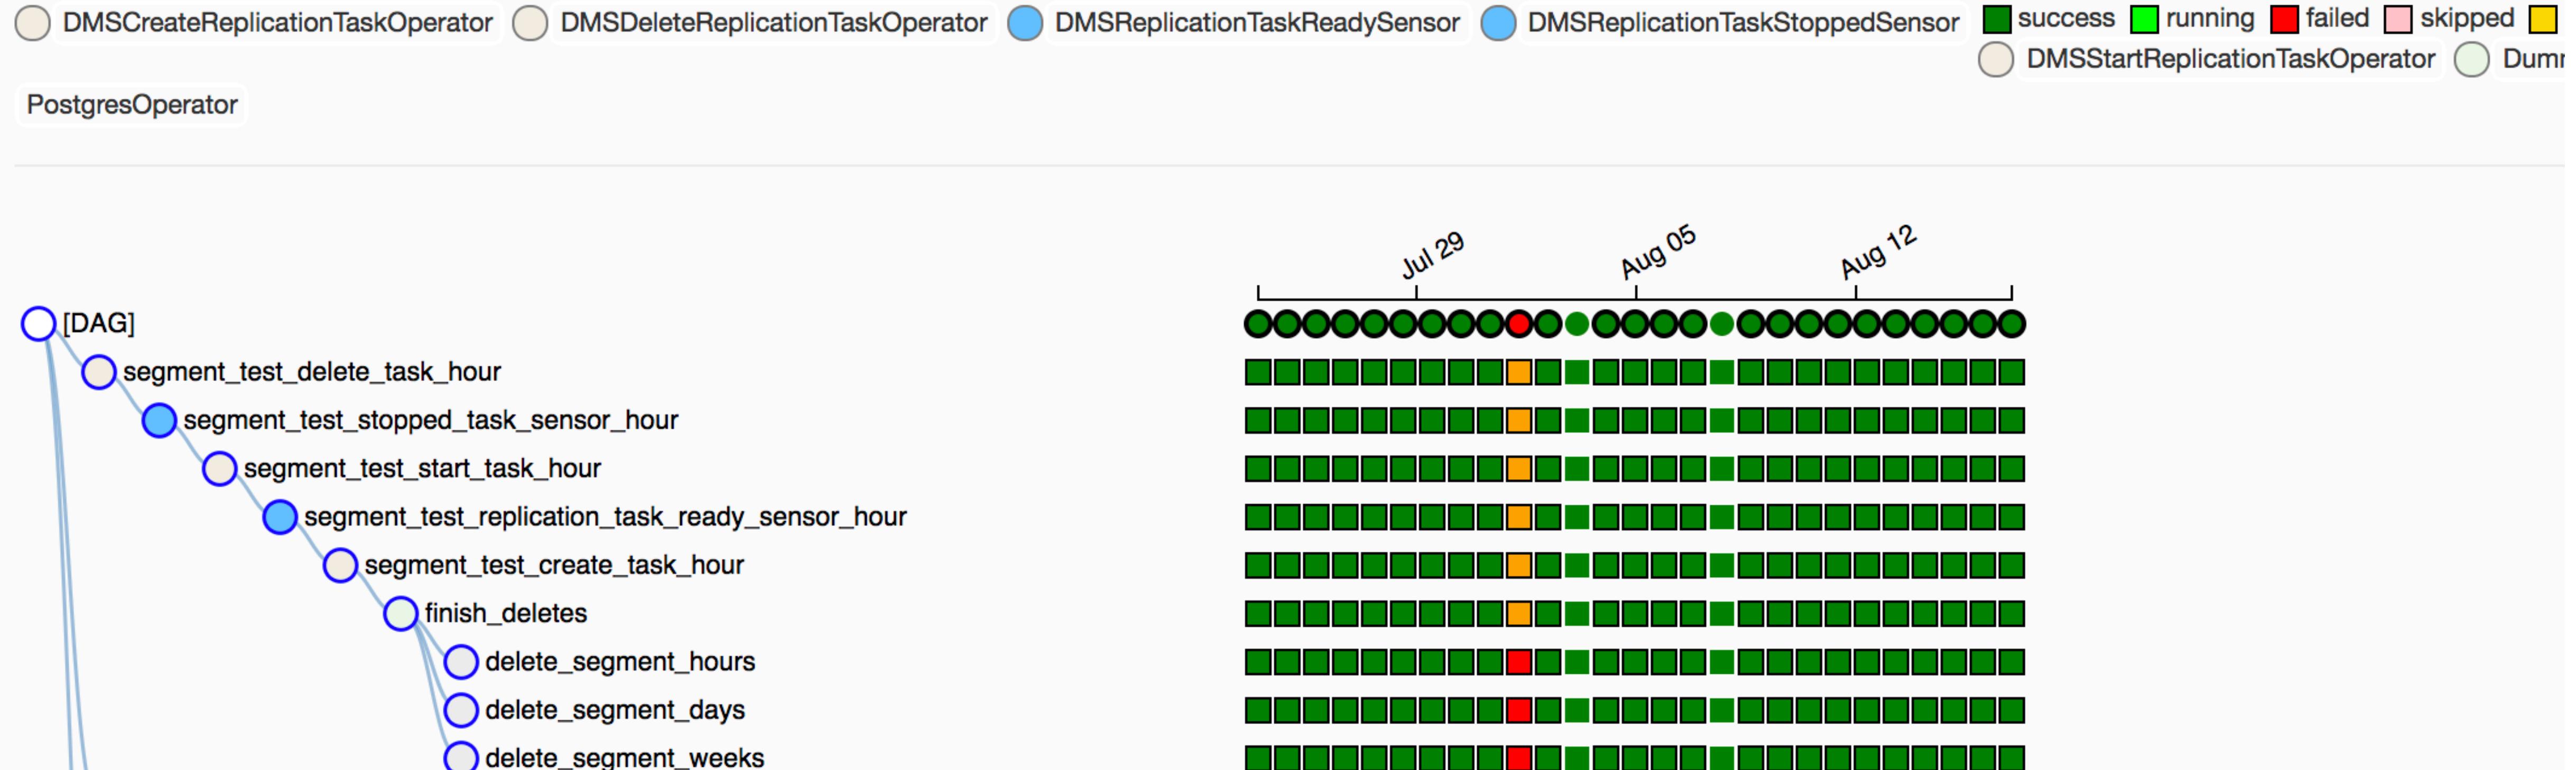


KubeCon

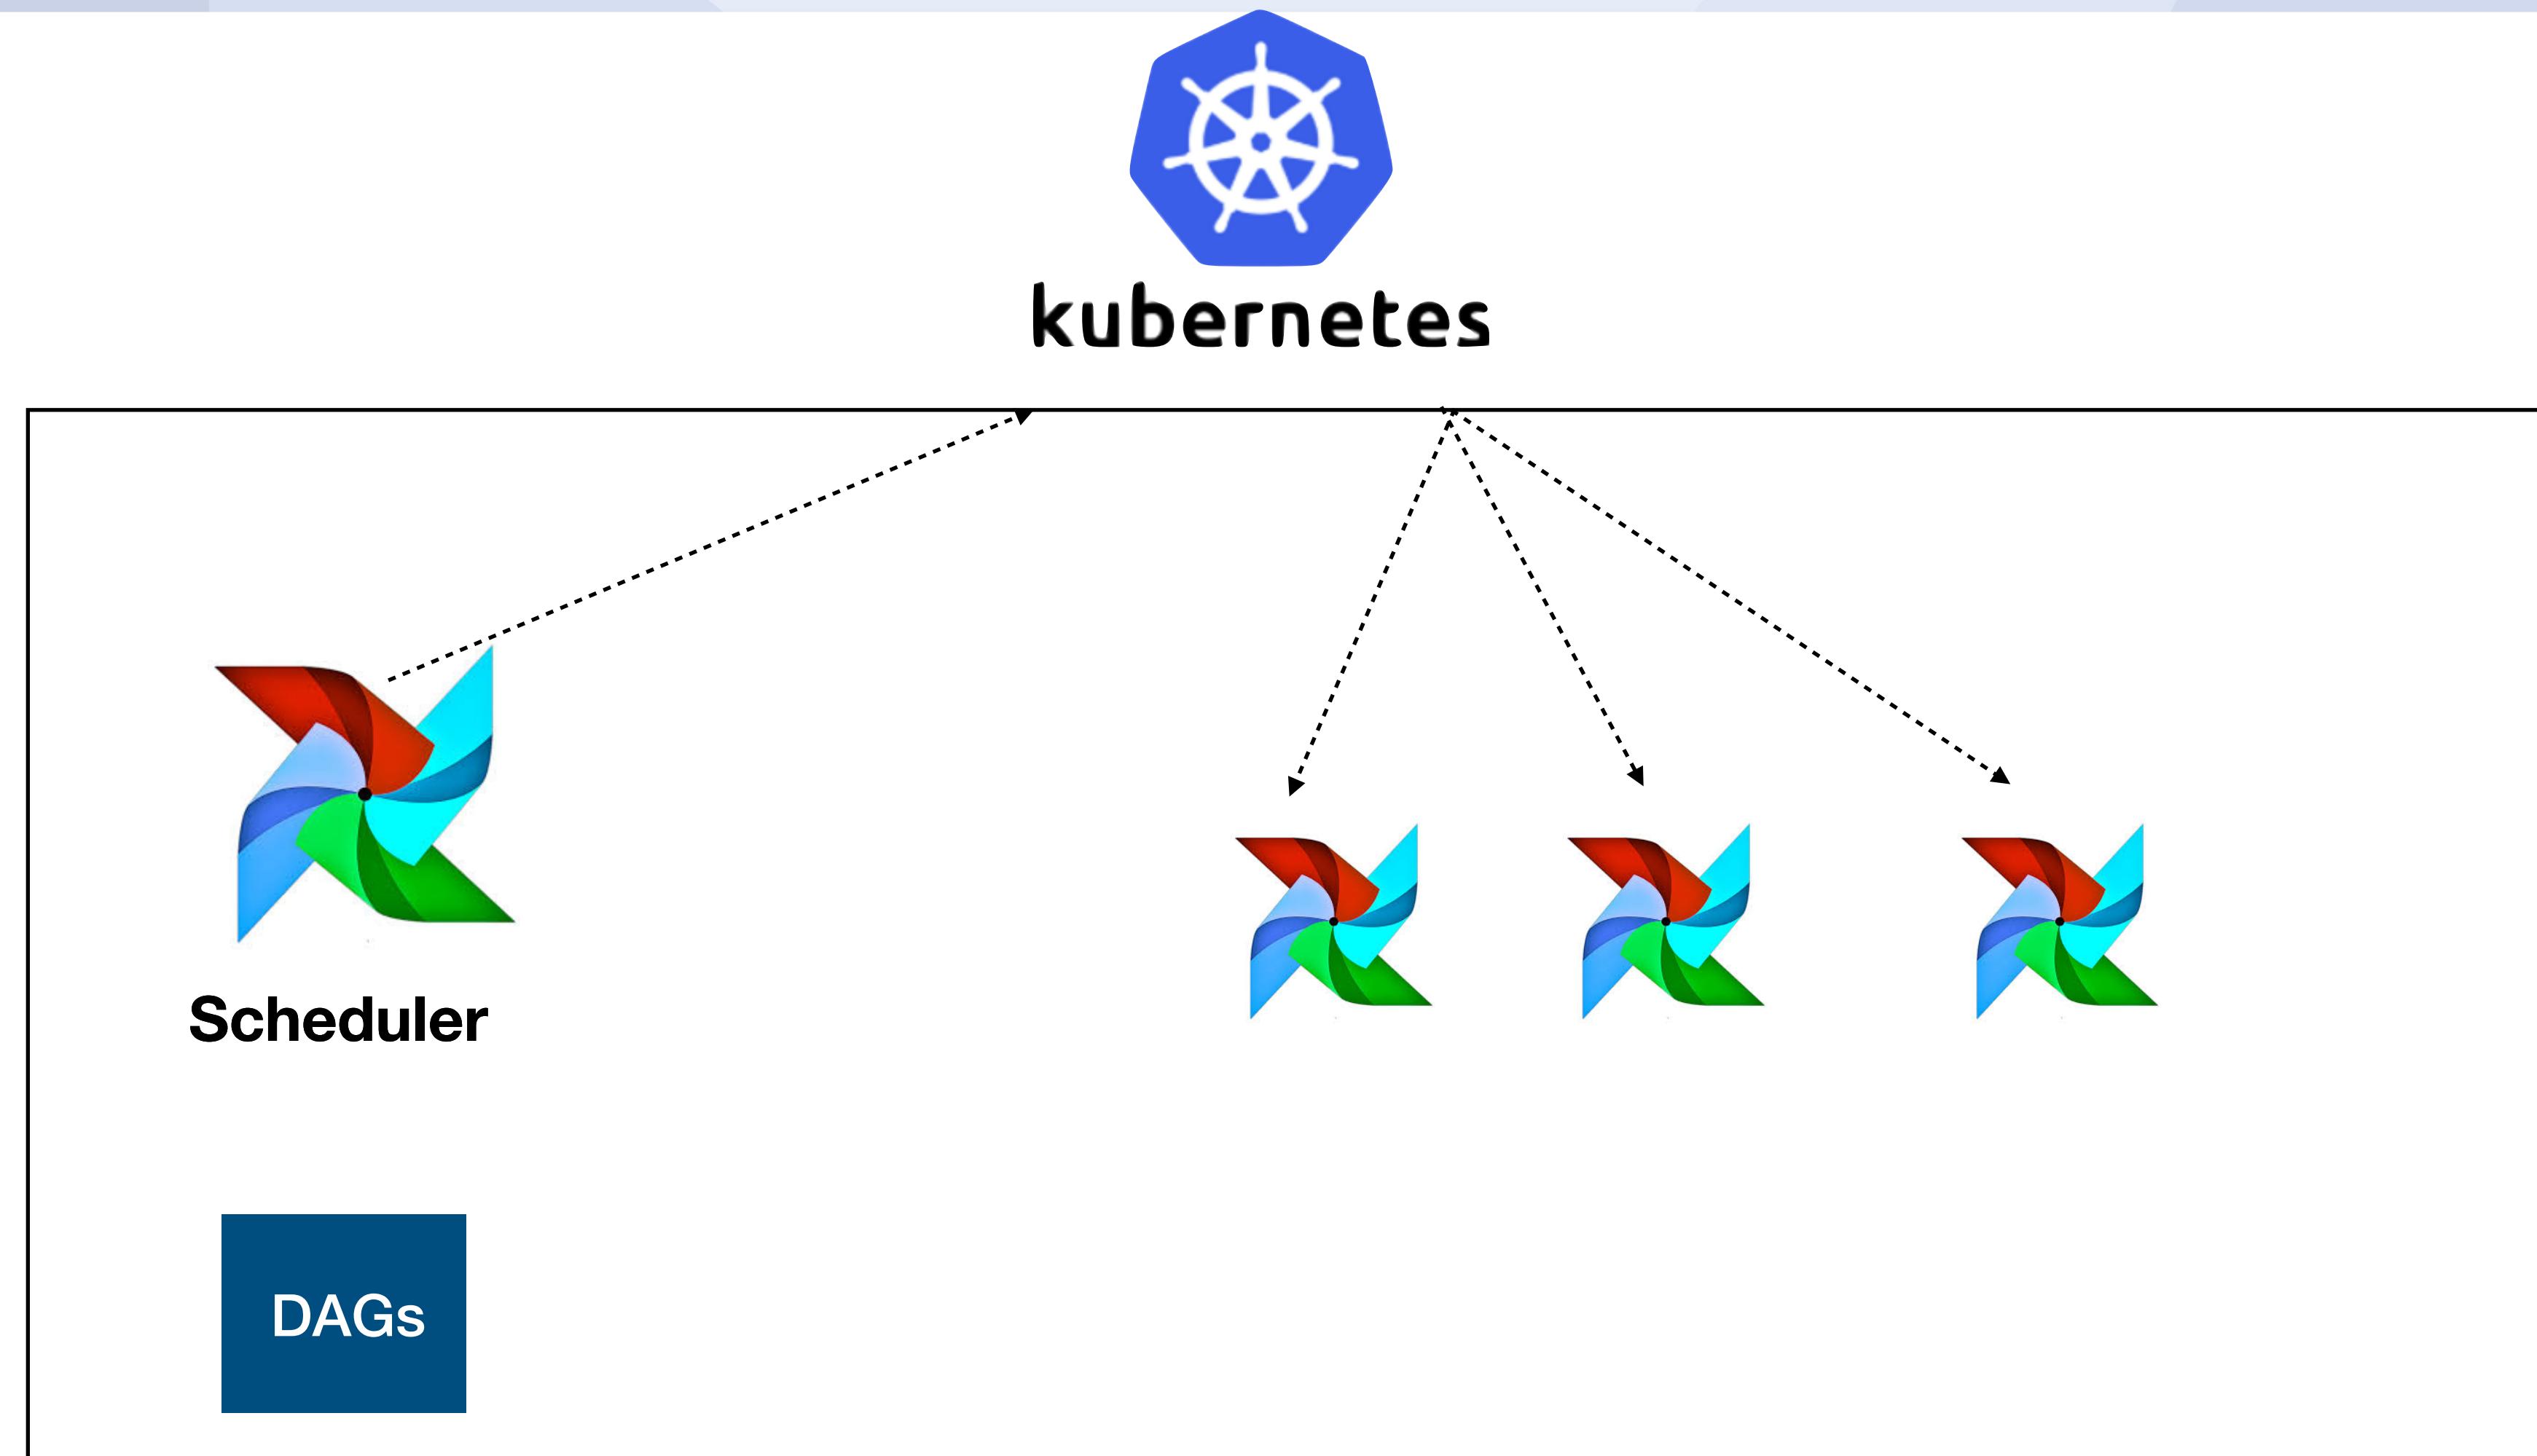


CloudNativeCon

North America 2018

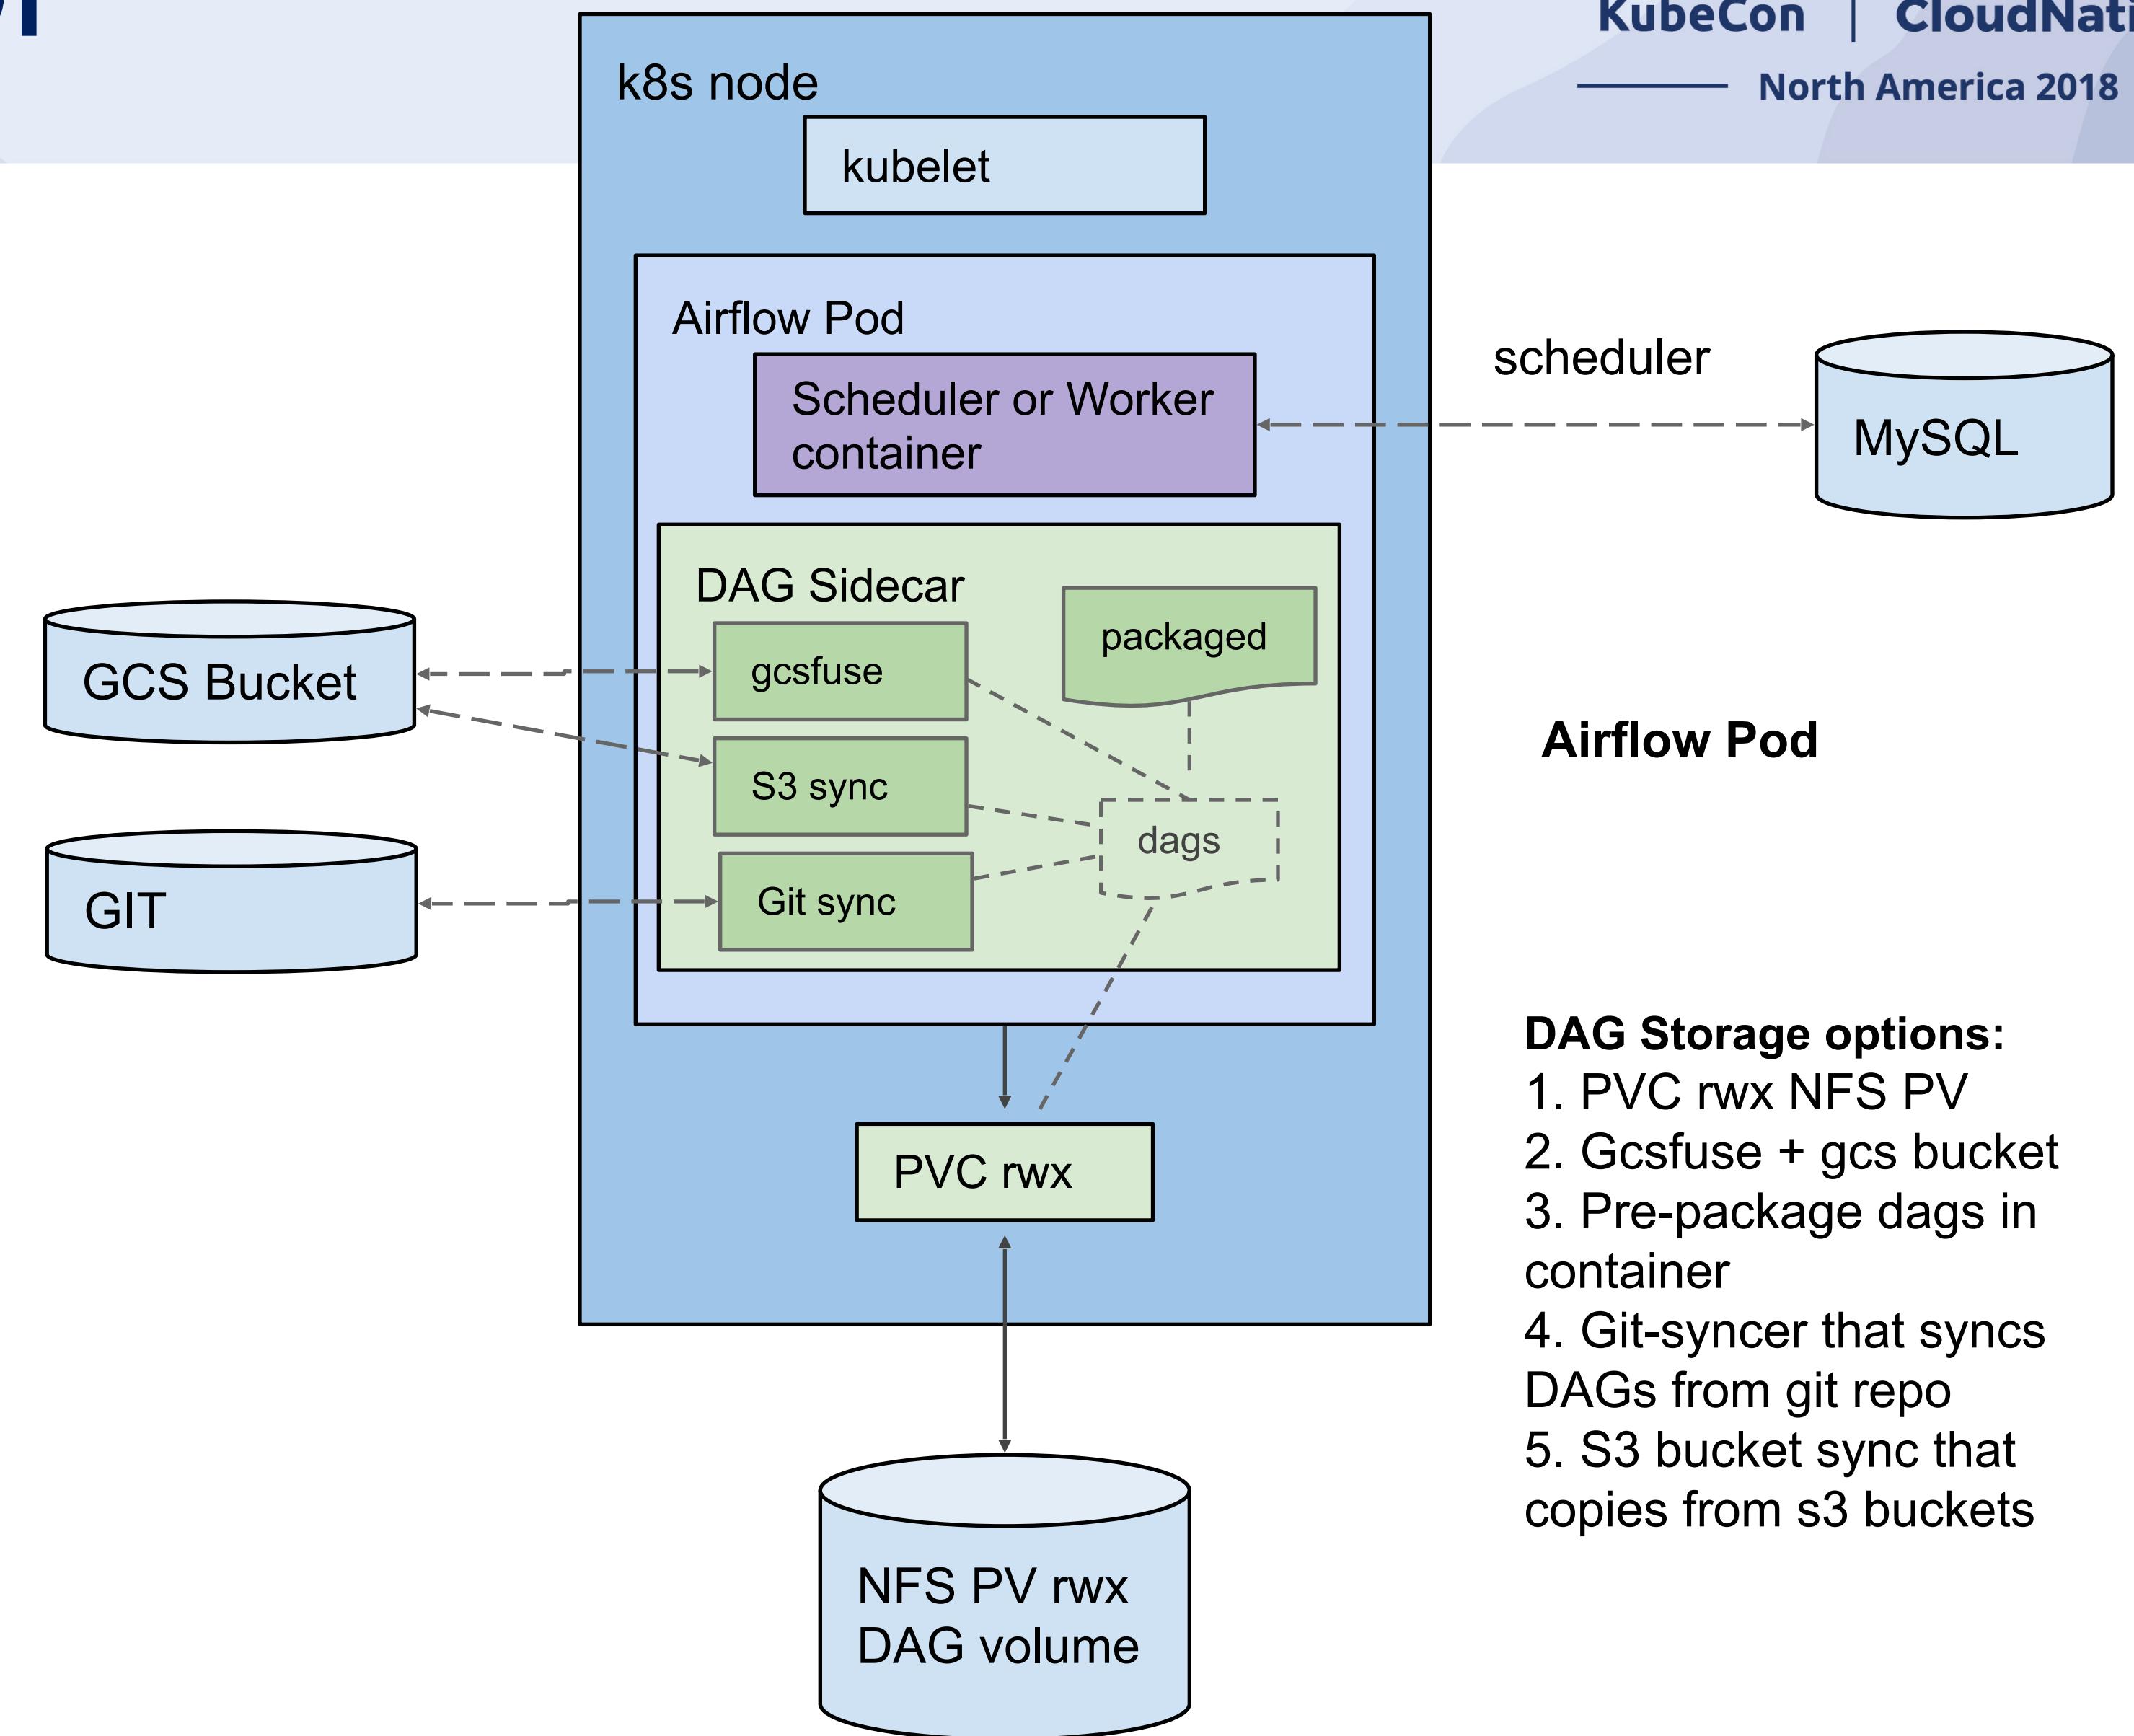


DAG Propagation



AirflowOperator

- Multiple DAG sources are supported via a DAG Sidecar
- Custom Airflow Pod images are supported



DAG Injection

- Three modes: Git-init mode, persistent volume mode, and “pre-baked” mode (1.10.1)
- Git-init mode + pre-baked is recommended for development and small instances of Airflow, because it does not involve any distributed file systems
- Persistent volume mode recommended for large DAG folders

Do we need this slide ?

K8sExecutor Status

- Has been released with Airflow 1.10 in experimental mode
- Multiple companies already using in production
- Helm chart in progress
- Active community in #sig-big-data on kubernetes.slack.com

Airflow Operator Status

- Supports Airflow 1.10.1
- We are currently active with the #sig-big-data and #airflow-operator channels on kubernetes.slack.com
- Available on Kubernetes Cloud Marketplace in GCP

Demo



KubeCon



CloudNativeCon

North America 2018

What's Next?



Throttling

- The Kubernetes Scheduler is naturally greedy (it will take as many tasks as it can, even if it doesn't have available resources)
- Want to ensure that if the cluster is in trouble, that we stop launching tasks until resources are available
- Maintain a “pending set” to ensure that no more than 5 Airflow tasks are pending at once

Scale Testing

- Have been working with members of Kubernetes team to create GKE instances for scale testing
- Writing scripts to generate giant DAGs to test DAG propagation and parsing times
- Seeing what scale of tasks begin to place pressure on PostgreSQL instance

How can you get involved?

- Looking for brave souls interested in experimenting with Executor and AirflowOperator to take it to beta
- Distributed testing using Chaos Monkey/Kubernetes
- Currently active with the #sig-big-data and #airflow-operator channels on kubernetes.slack.com
- Scale testing
- Throttle DAG tasks in Kubernetes executor

Thank You

Learn more:

github.com/apache/incubator-airflow/

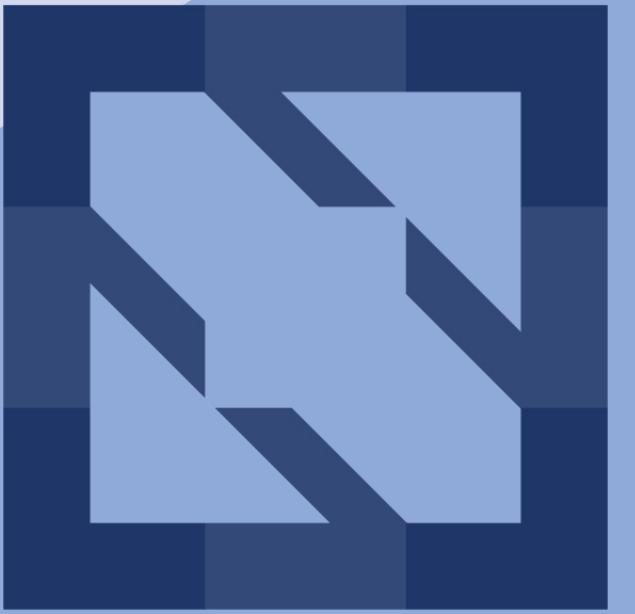
@danimberman

github.com/GoogleCloudPlatform/airflow-operator

@bharanis



KubeCon



CloudNativeCon

North America 2018

