



# Deep Dive: TUF

Trishank Karthik Kuppusamy

*Kubecon North America 2018*



DATADOG

# Repository compromise

## Software updates

- Experts agree: software updates the most security practice (USENIX SOUPS 2015)
- Updates fix security vulns
- However, important problem is often neglected...



### **“...no one can hack my mind”: Comparing Expert and Non-Expert Security Practices**

Julia Ion  
Google  
juliaion@google.com

Rob Reeder  
Google  
reeder@google.com

Sunny Consolvo  
Google  
sconsolvo@google.com

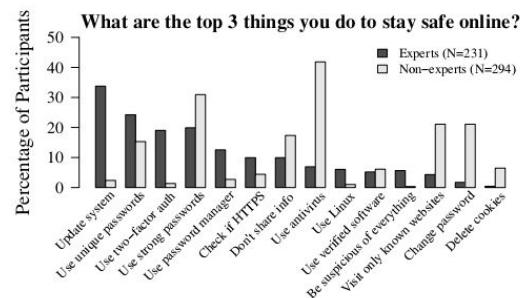


Figure 1: Security measures mentioned by at least 5% of each group. While most experts said they keep their system updated and use two-factor authentication to stay safe online, non-experts emphasized using antivirus software and using strong passwords.

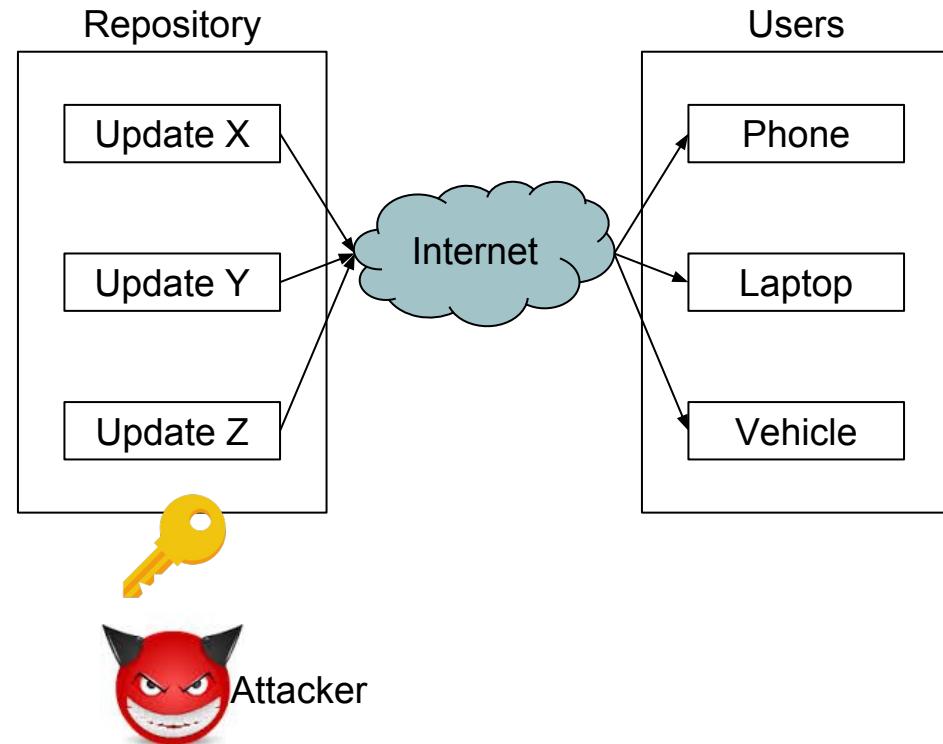
## Repository compromise

- Examples:
  - Microsoft Windows Update (2012): Flame malware targeted Iran nuclear efforts
  - South Korea cyberattack (2013): >\$750M USD in economic damage
  - NotPetya (2017): infected multinational corporations
- Compromise millions of devices



## Goal: compromise-resilience

- Only question of when, not if
- Cannot prevent compromise
- But must severely limit impact





DATADOG

# The Update Framework (TUF)



## What is on a repository?

- Repository contains packages + metadata



## What is on a repository?

- Repository contains packages + metadata
- *Package*
  - Smallest unit of update
  - Software application or library

Package



## What is on a repository?

- Repository contains packages + metadata
- *Package*
  - Smallest unit of update
  - Software application or library
- *Metadata*
  - Cryptographic hashes, file sizes, version numbers, etc.
  - About packages, or other metadata files

```
{  
    "signatures": [  
        {  
            "keyid": "ce3e02e72980b09ca6f5efa68197130b381921e5d0675e2e0c8f3c47e0626bba",  
            "method": "ed25519",  
            "sig": "9095bf34b0cbf9790465c0956810cb3729bc96beed8ee7e42d98997b1e8ec0a6780e575565"  
        }  
    ],  
    "signed": {  
        "_type": "Targets",  
        "expires": "2030-01-01T00:00:00Z",  
        "targets": {  
            "/project/file3.txt": {  
                "hashes": {  
                    "sha256": "141f740f53781d1ca54b8a50af22cbf74e44c21a998fa2a8a05aaac2c002886b"  
                },  
                "length": 28  
            }  
        },  
        "version": 1  
    }  
}
```

↓  
Package

# The Update Framework (TUF): secure software updates

- **Authenticity and integrity**— even if repository compromised
- **Design principles**
  - Separation of duties
  - Threshold signatures
  - Explicit & implicit revocation of keys
  - Minimizing risk using offline keys
  - Selective delegation of trust
  - Diversity of hashing + signing algorithms
- (CCS 2010)
- <https://theupdateframework.com>

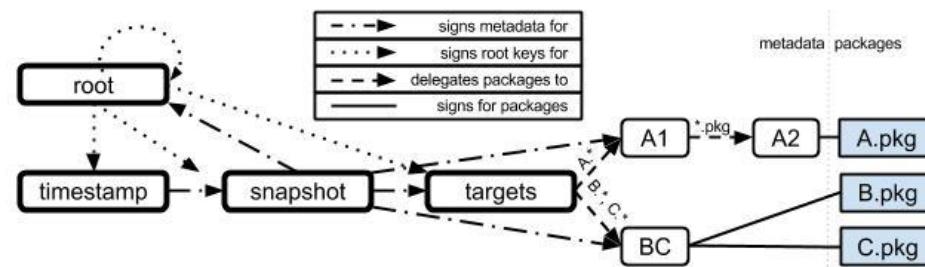
## Survivable Key Compromise in Software Update Systems

Justin Samuel<sup>\*</sup>  
UC Berkeley  
Berkeley, California, USA  
jsamuel@berkeley.edu

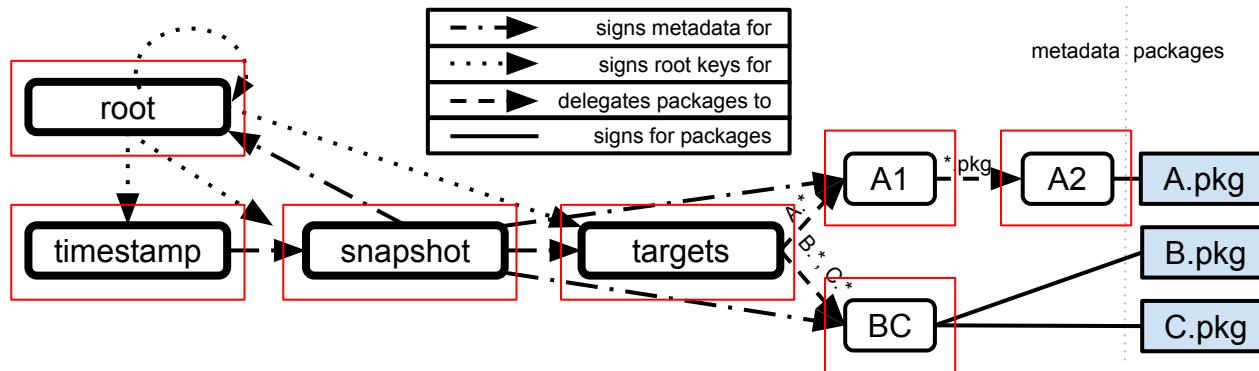
Nick Mathewson  
The Tor Project  
nickm@alum.mit.edu

Justin Cappos  
University of Washington  
Seattle, Washington, USA  
justinc@cs.washington.edu

Roger Dingledine  
The Tor Project  
arma@mit.edu



## Separation of duties

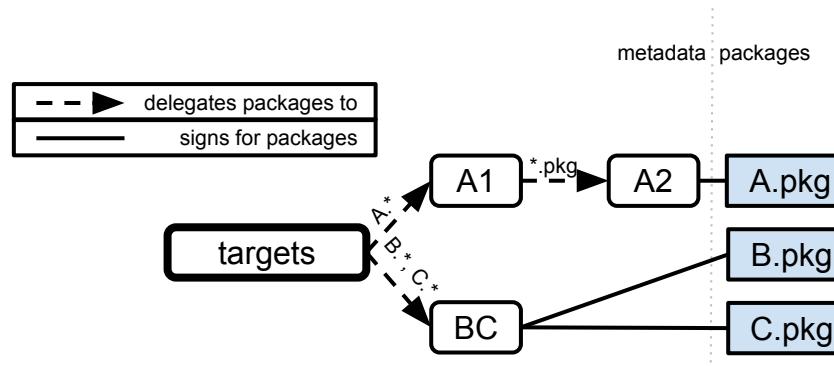


Design principles:

1. **Separation of duties**

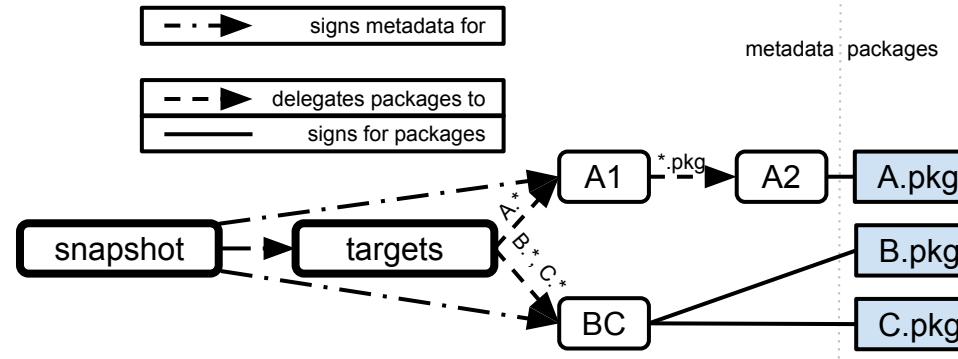
(i.e., don't put all your eggs in one basket).

## The targets role



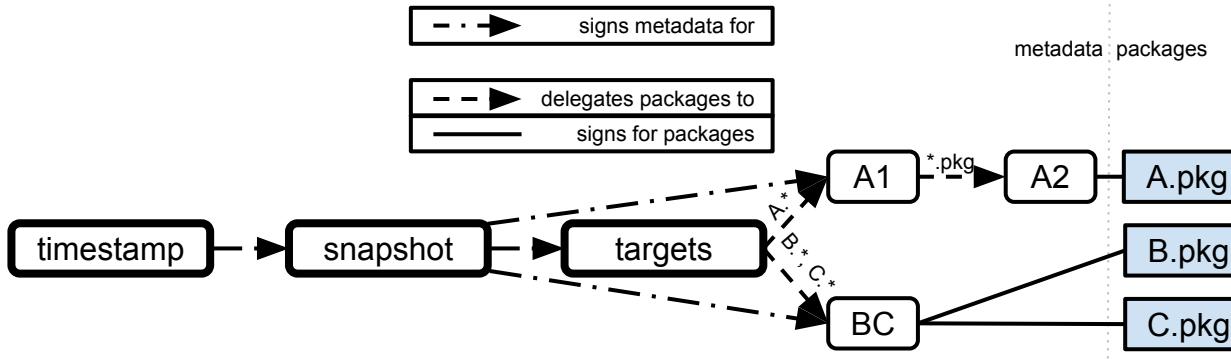
Role	Purpose
<b>targets</b>	Indicates metadata such as the cryptographic hashes and file sizes of packages. May delegate this responsibility to other, custom-made roles.

## The snapshot role



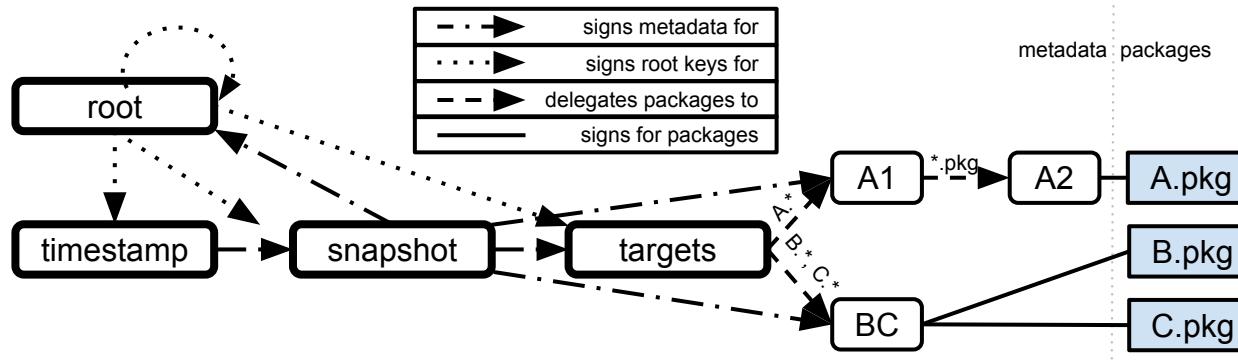
Role	Purpose
<b>targets</b>	Indicates metadata such as the cryptographic hashes and file sizes of packages. May delegate this responsibility to other, custom-made roles.
<b>snapshot</b>	Indicates which packages have been released at the same time by the repository.

## The timestamp role



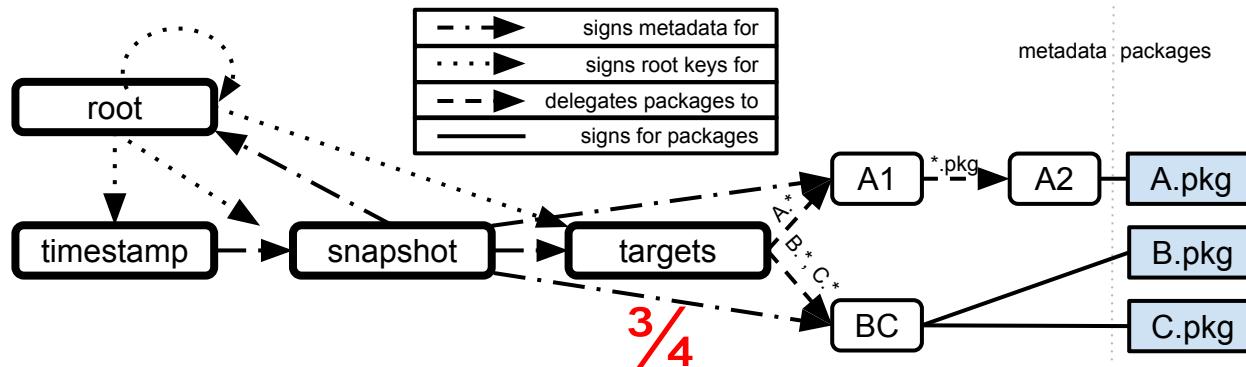
Role	Purpose
<b>targets</b>	Indicates metadata such as the cryptographic hashes and file sizes of packages. May delegate this responsibility to other, custom-made roles.
<b>snapshot</b>	Indicates which packages have been released at the same time by the repository.
<b>timestamp</b>	Indicates whether there is any new metadata or package on the repository.

## The root role



Role	Purpose
<b>targets</b>	Indicates metadata such as the cryptographic hashes and file sizes of packages. May delegate this responsibility to other, custom-made roles.
<b>snapshot</b>	Indicates which packages have been released at the same time by the repository.
<b>timestamp</b>	Indicates whether there is any new metadata or package on the repository.
<b>root</b>	Serves as the certificate authority for the repository. Distributes and revokes the public keys used to verify the root, timestamp, snapshot, and targets role metadata.

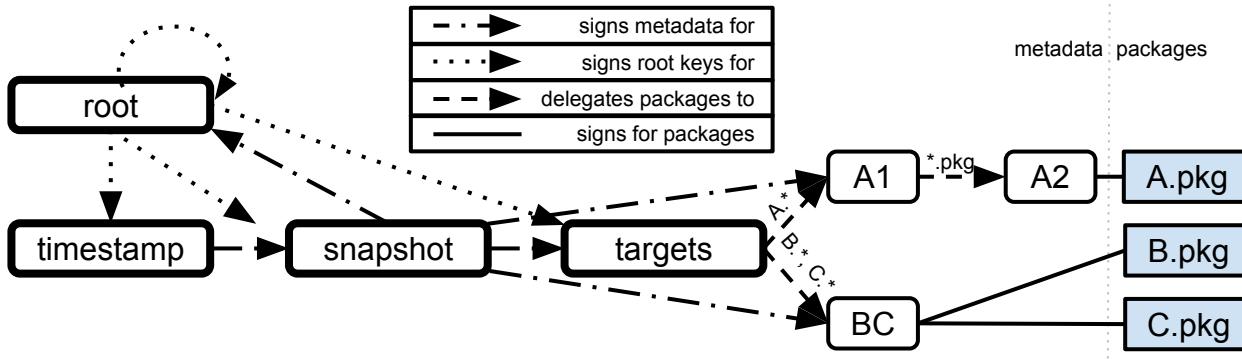
## Threshold signatures



Design principles:

1. Separation of duties.
2. **Threshold signatures**  
(i.e., like the two-man rule to launch nuclear missiles).

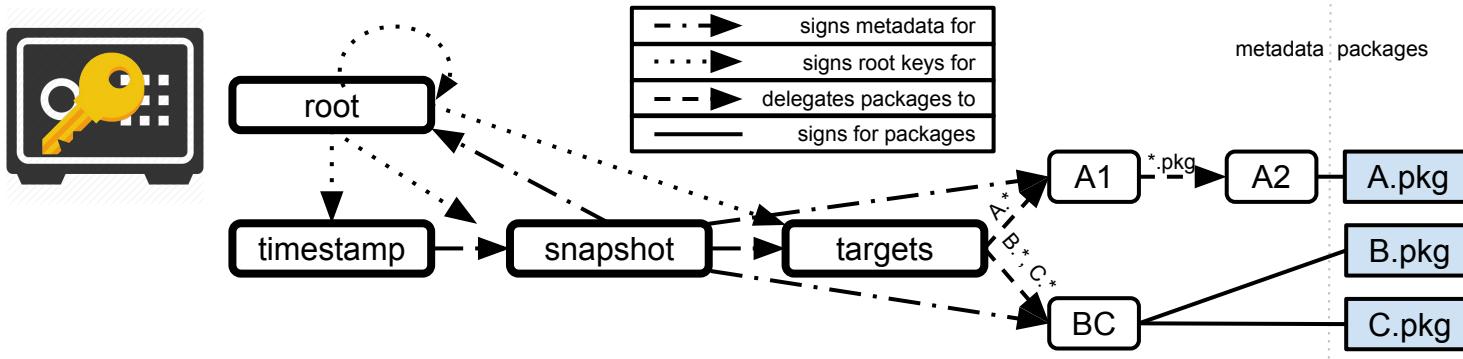
## Explicit & implicit revocation of keys



Design principles:

1. Separation of duties.
2. Threshold signatures.
3. **Explicit and implicit revocation of keys.**

## Minimizing risk with offline keys



### Design principles:

1. Separation of duties.
2. Threshold signatures.
3. Explicit and implicit revocation of keys.
4. **Minimized risk through use of offline keys**  
(i.e., don't put keys to the kingdom under the carpet).

## Diversity of cryptographic algorithms

- **Hedge your bets**
- Can't break TUF unless you **break them all**
- No need to depend on just SHA-2 or SHA-3, RSA or Ed25519
- Can even try **post-quantum crypto** at the same time

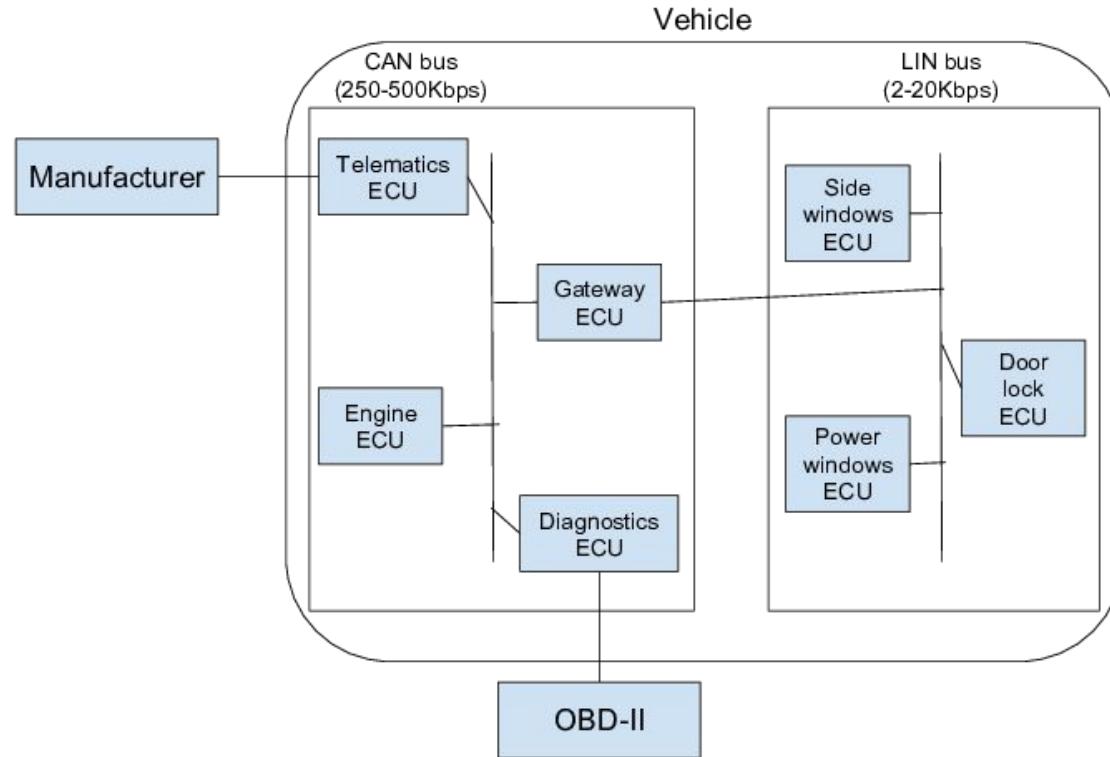




DATADOG

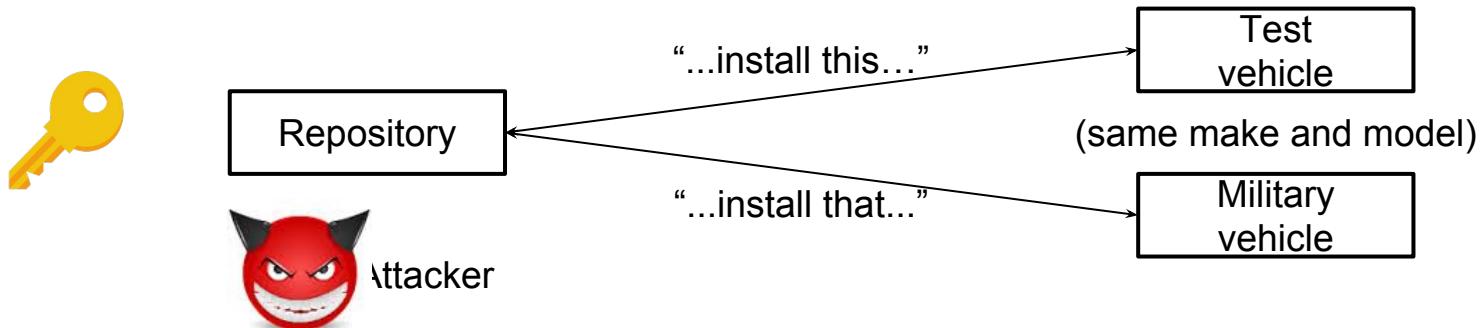
# TUF in practice: Securing Software Updates for Automobiles

## Vehicles and ECUs



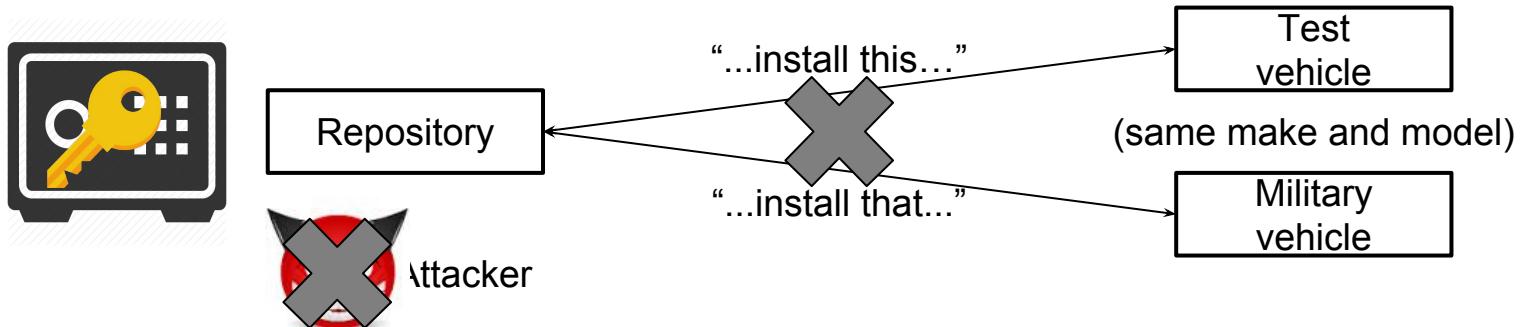
## Customizability...

- Use **online keys** to sign all metadata
- **Pro: on-demand customization**
  - Easy to install different updates on vehicles of same make and model
  - Can instantly blacklist only buggy updates
- **Con: no compromise-resilience**
  - Attackers cannot tamper with metadata without being detected



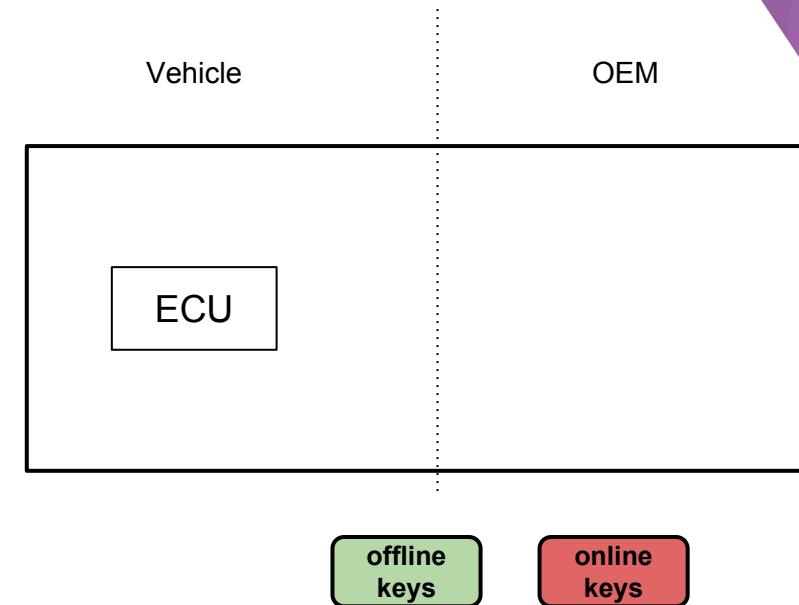
## ...or security?

- Use **offline keys** to sign all metadata
- **Pro: compromise-resilient**
  - Attackers cannot tamper with metadata without being detected
- **Con: no on-demand customization**
  - Difficult to install different updates on vehicles of same make and model
  - Cannot instantly blacklist only buggy updates



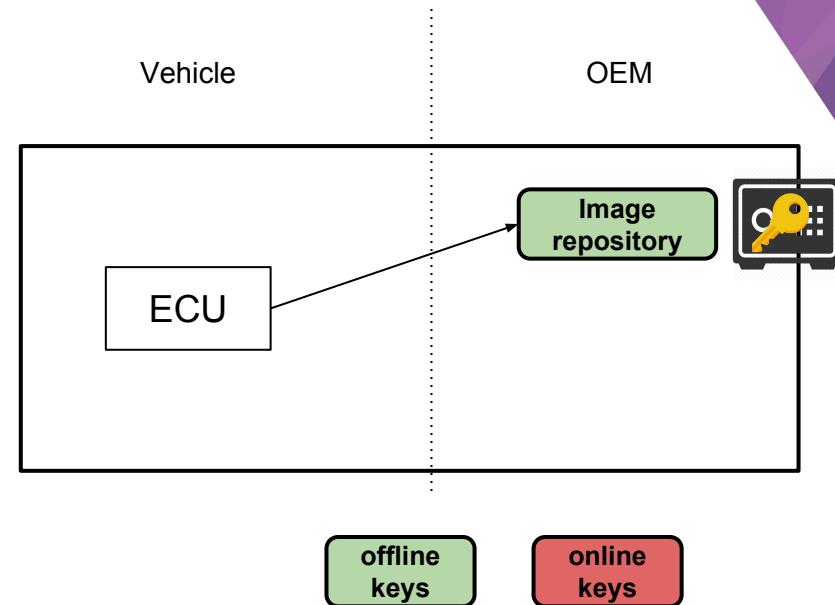
## Key idea

- What if there are **two** repositories?



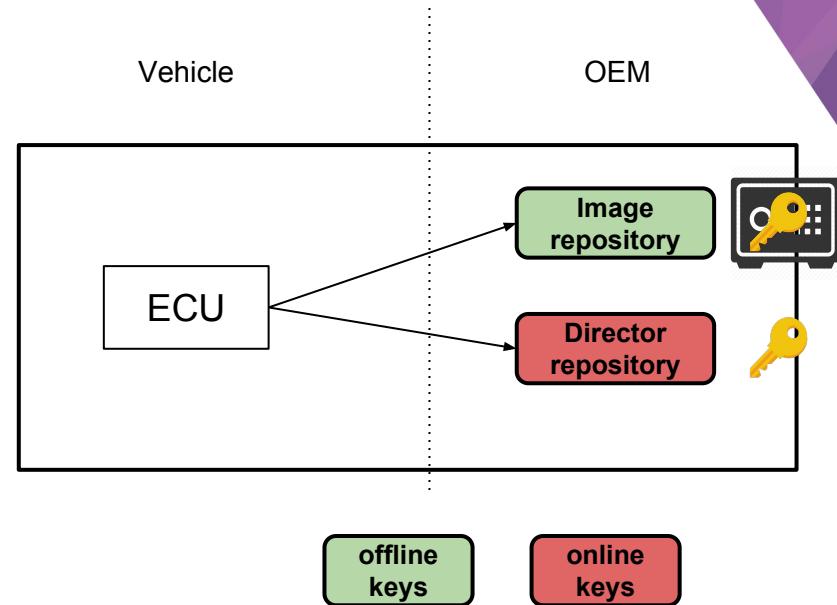
## Key idea: image repository

- What if there are two repositories?
- **Image repository**
  - Uses **offline keys**
  - Provides signed metadata about all available updates for **all ECUs** on **all vehicles**



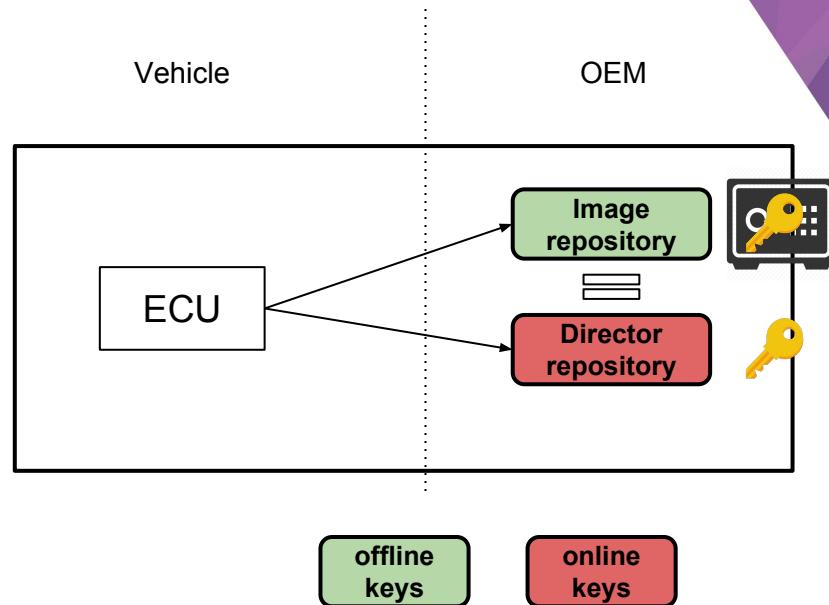
## Key idea: director repository

- What if there are two repositories?
- Image repository
  - Uses offline keys
  - Provides signed metadata about all available updates for all ECUs on all vehicles
- Director repository
  - Uses online keys
  - Signs metadata about **which updates** should be installed on **which ECUs** on a vehicle



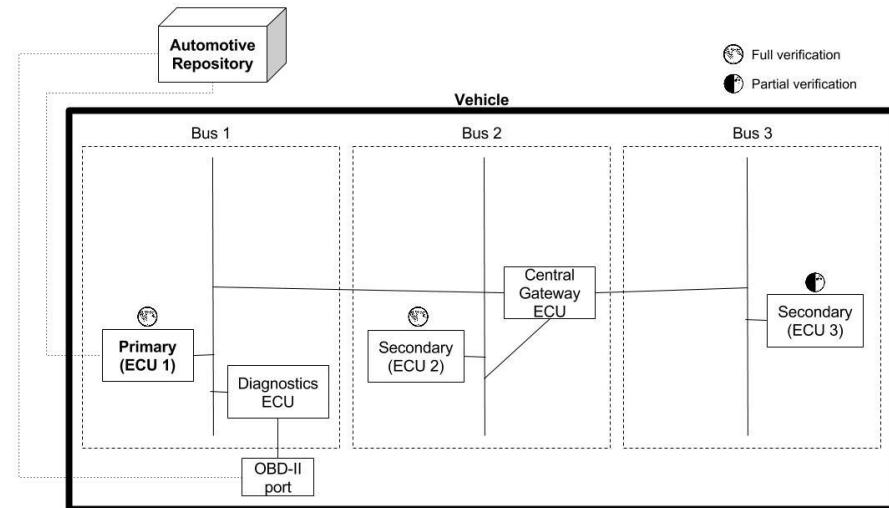
## Key idea: security & customizability

- A vehicle would ensure that installation instructions from the **director** repository **match** updates from the **image** repository.
- Using both repositories provides **both** on-demand **customization** of vehicles & **compromise-resilience**.



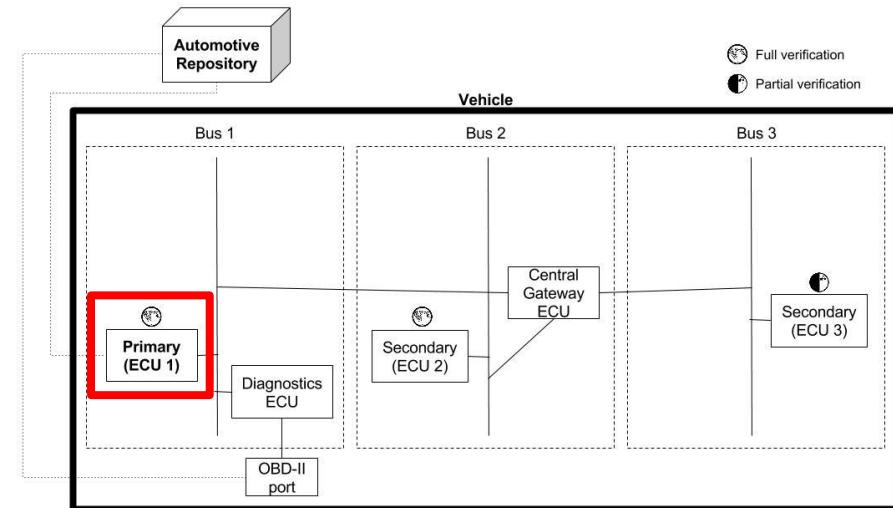
## Primaries and secondaries

- **Two types of ECUs,**  
because:
  - Some ECUs are more / less **powerful** than others.
  - Few ECUs have network **connection** to outside world.
  - ECUs should not **download** metadata independently of each other.



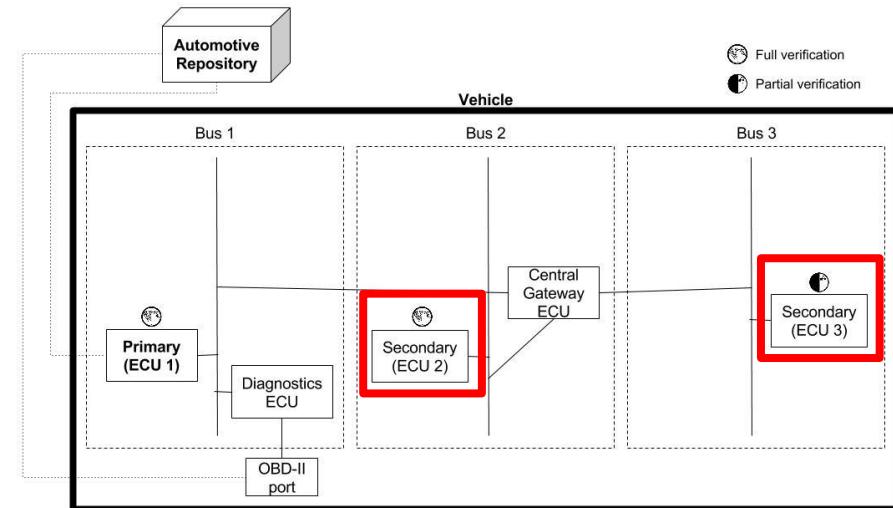
## Primaries

- A **primary** downloads, verifies, distributes metadata + images to secondaries.



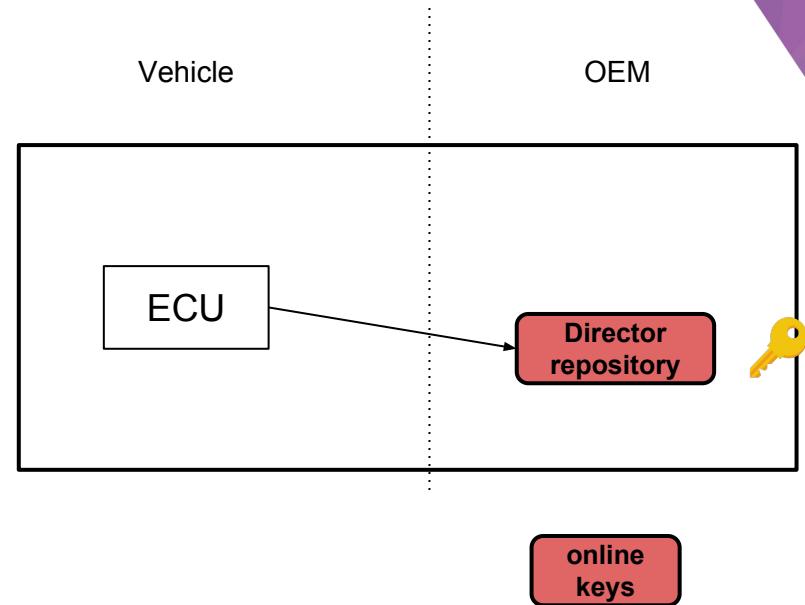
## Secondaries

- A **secondary** verifies both the metadata & image distributed by a primary, before updating to that image.



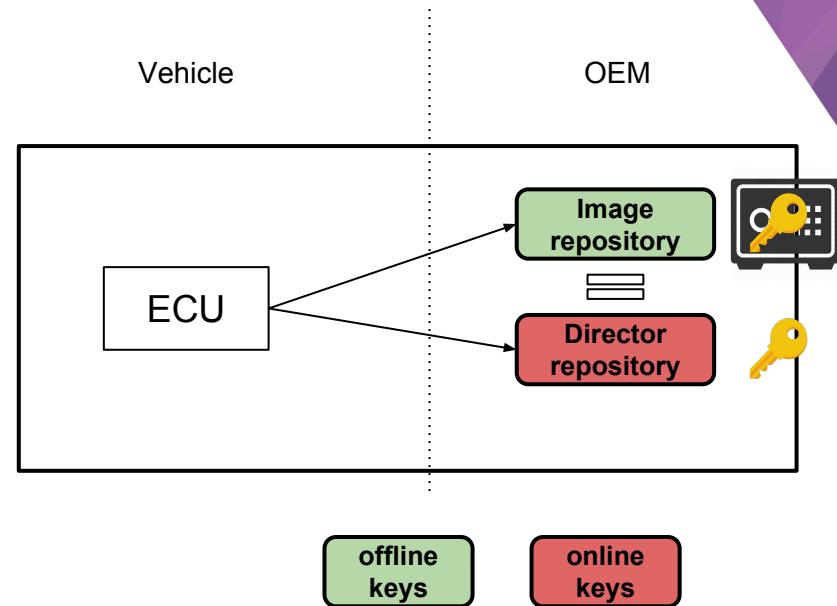
## Partial verification

- Checking **only** metadata from the **director repository**.
- Involves checking **only one signature** on **one metadata file**.



## Full verification

- Checking that metadata about updates chosen by the **director** repository **match** metadata about the same updates on the **image** repository.
- Involves checking **many signatures** on **many metadata files** from both repositories.



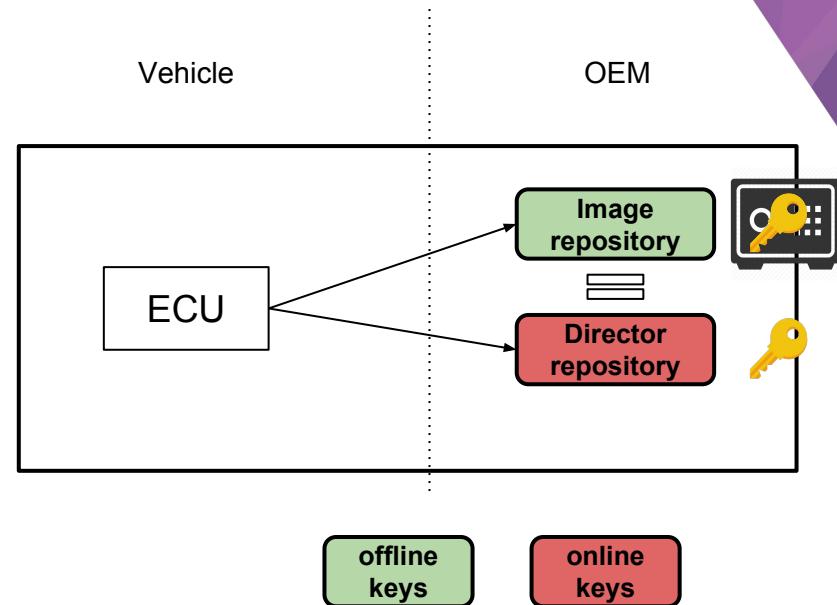
## Partial vs full verification

Increasing difficulty for attackers →

		<b>MitM outside / inside vehicle</b>	<b>MitM + director repository compromise</b>	
			<b>Primaries not compromised</b>	<b>Primaries compromised</b>
Greater compromise-resilience ↓	<b>Partial verification</b>	Cannot cause ECUs to fail to interoperate	Can cause ECUs to fail to interoperate	Can install malware
	<b>Full verification</b>			Cannot install malware
		<span style="background-color: #ccc; padding: 2px 10px; display: inline-block;">Mild</span> <span style="background-color: #fff; border: 1px solid #ccc; padding: 2px 10px; display: inline-block;">Serious</span> <span style="background-color: #cc0000; border: 1px solid #ccc; padding: 2px 10px; display: inline-block;">Critical</span>		

## Takeaway: security & customizability

- Uptane provides **both** on-demand customization of vehicles & compromise-resilience.





DATADOG

# TUF in practice:

## Datadog Agent integrations



# Datadog, Agent, and Agent integrations

- **3 pillars of Datadog monitoring**
  - Infrastructure metrics
  - App performance
  - Logs
- **Agent**
  - Collects events and metrics
- **Agent integrations**
  - Add-ons / plug-ins
  - > 100 and counting

The screenshot shows the Datadog Integrations page. At the top, there's a navigation bar with links for 'CUSTOMERS', 'ABOUT', 'BLOG', 'LOGIN', and a purple 'GET STARTED FREE' button. Below the navigation is a section titled 'Integrations' with the subtext 'More than 200 built-in integrations. See across all your systems, apps, and services.' A large grid of integration icons is displayed in a 4x5 layout. The categories represented by the columns are: All, API, AWS, Azure, Caching, Chaos Engineering, Cloud, Collaboration, Configuration & Deployment, Containers, Cost Management, Data Store, Direct Connect, Exceptions, Google Cloud, Health, Issue Tracking, Languages, Log Collection, Messaging, Monitoring, Network, Notification, Orchestration, OS & System, Processing, Provisioning, Search, Security, Source Control, and Web. Below the grid is a search bar with the placeholder 'Search for an integration...'. The icons for each integration include: Amazon ECS (orange cube), Amazon EKS (blue hexagon with 'K'), CONSUL (red circle with dots), CRI (blue hexagon with a ship wheel), cri-o (blue hexagon with a snowflake), docker (blue cloud with a white dog), AWS Fargate (yellow hexagon with a white 'F'), etcd (blue hexagon with a gear), Google Container Engine (blue hexagon with a white 'G'), and kubernetes (blue hexagon with a white ship wheel).

## Decoupling integrations from Agent release cycle

- **Agent**
  - 6-week release cycle
- **Agent integrations**
  - Latest versions bundled with the Agent every 6 weeks
  - But we also want to publish new versions *independently* of the Agent
  - So customers can beta-test immediately



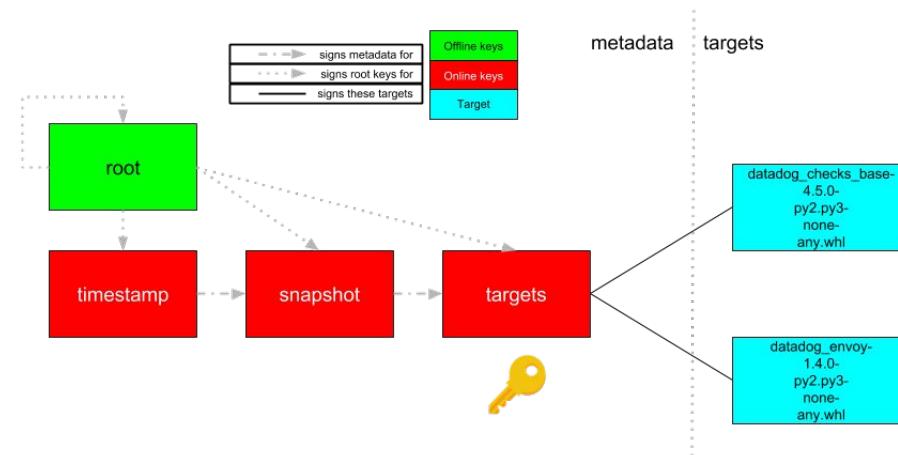
## State-of-the-art: CI/CD

- **CI/CD**

- Continuous integration / continuous deployment

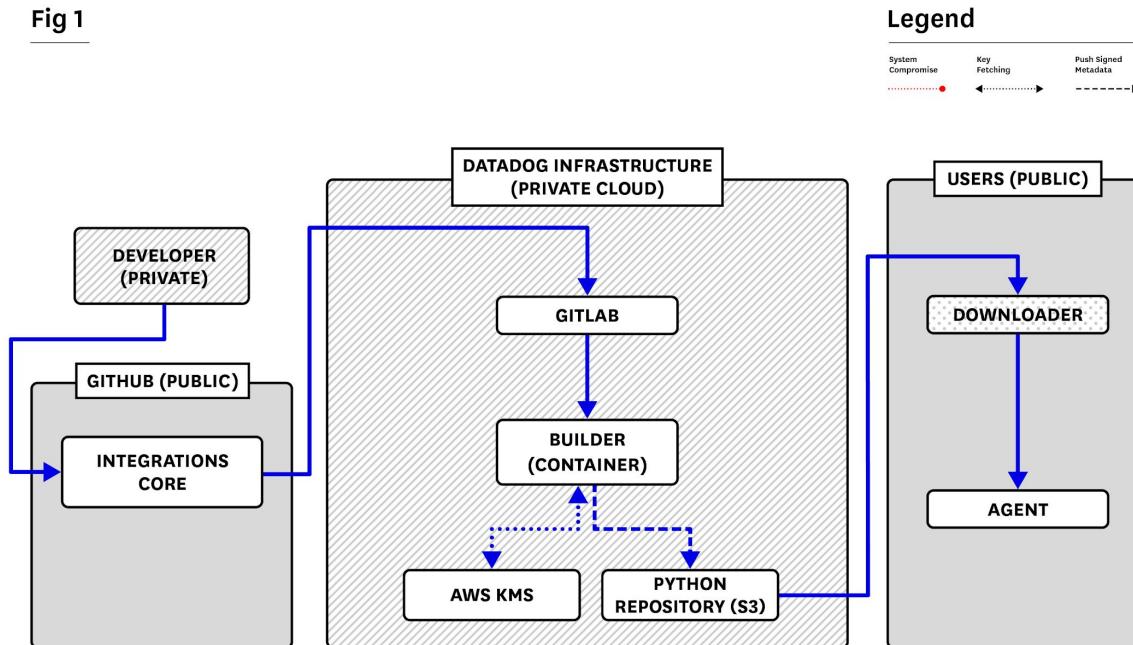
- **Pros**

- Faster deployments
- Clean build environments
- More secure handling of code-signing keys



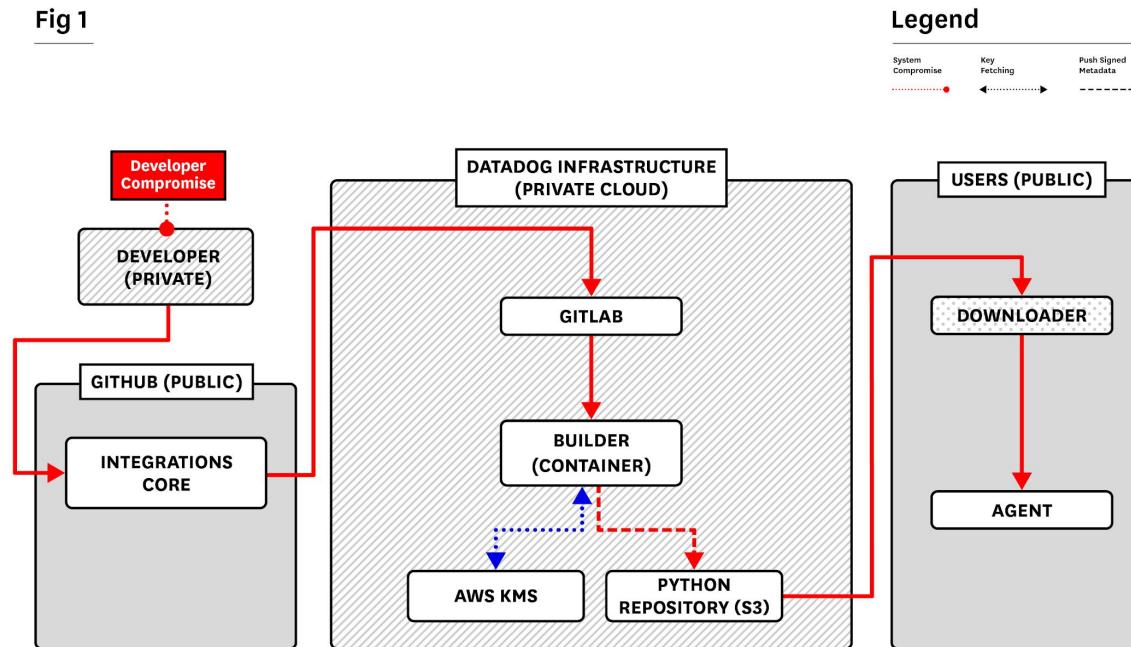
# State-of-the-art: what can go wrong?

**Fig 1**



# State-of-the-art: developer key compromise

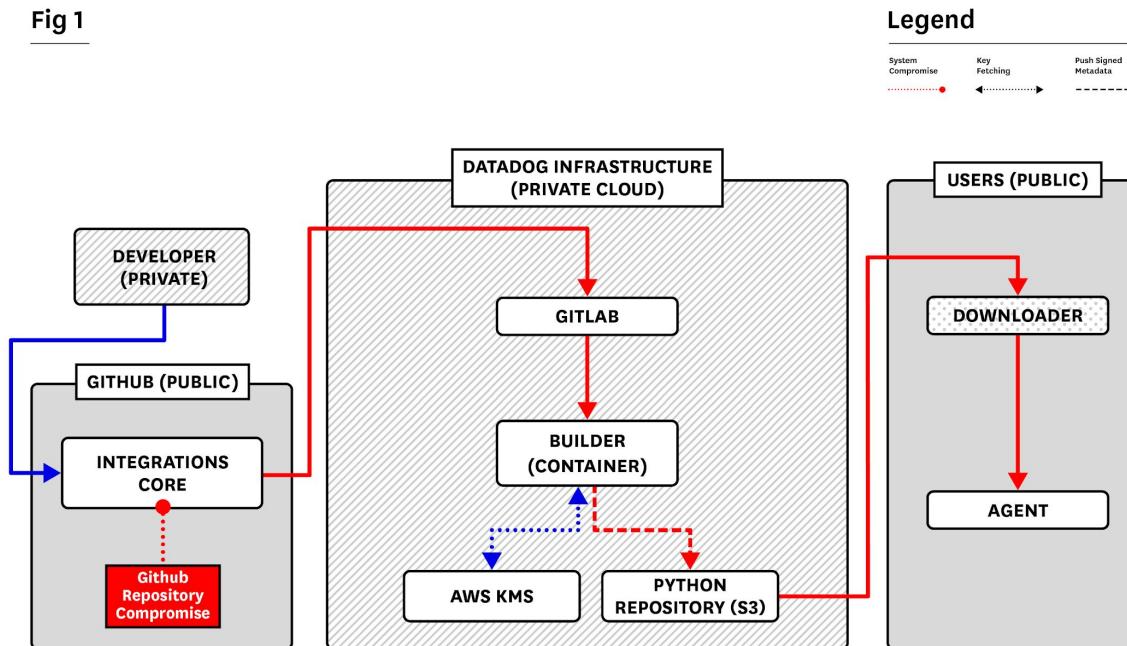
**Fig 1**





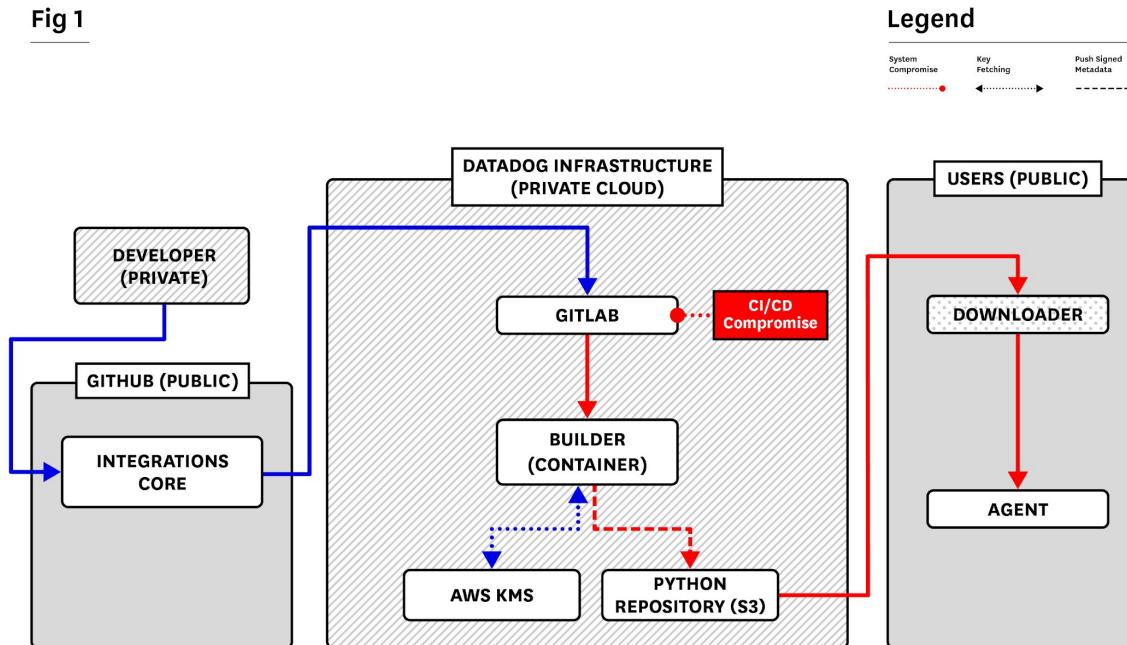
# State-of-the-art: VCS repository compromise

Fig 1



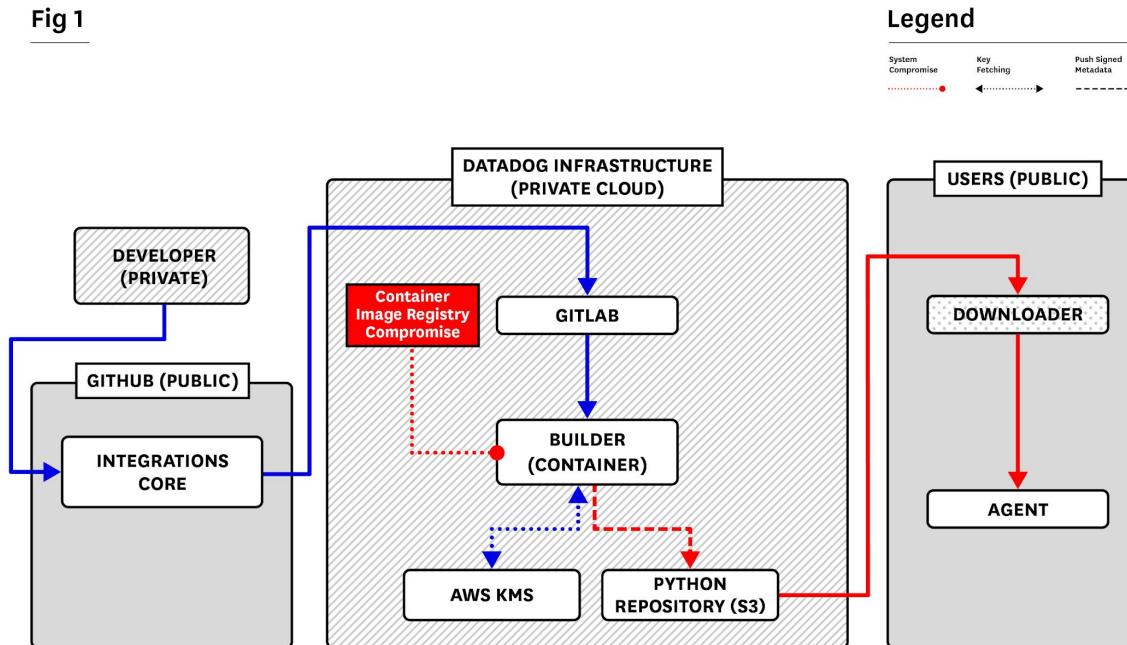
# State-of-the-art: CI/CD system compromise

**Fig 1**



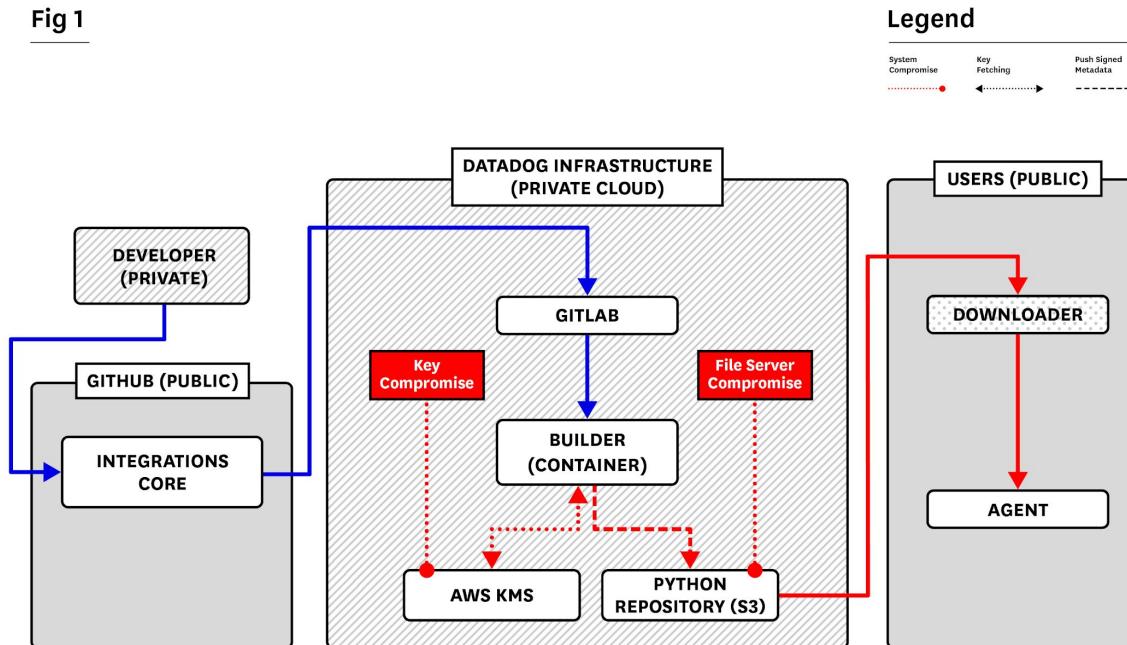
# State-of-the-art: container image registry compromise

**Fig 1**



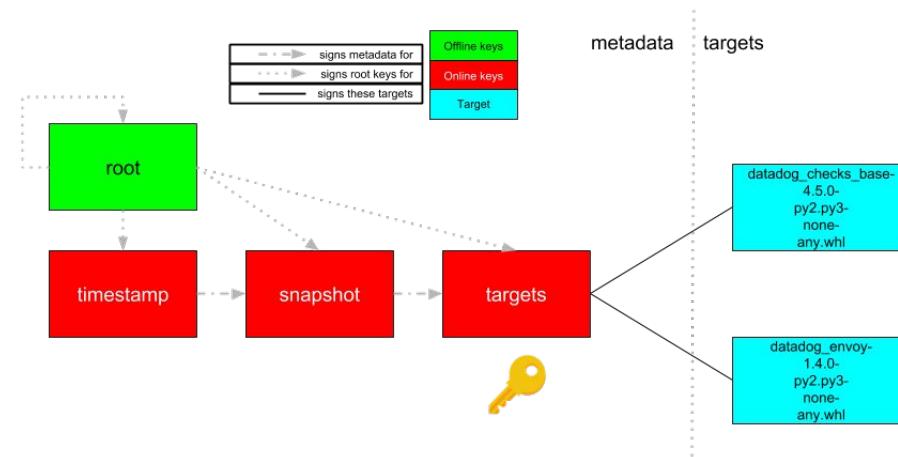
# State-of-the-art: key + file server compromise

**Fig 1**



## State-of-the-art: no compromise-resilience

- CI/CD
  - Continuous integration / continuous deployment
- Pros
  - Faster deployments
  - Clean build environments
  - More secure handling of code-signing keys
- Cons
  - **No compromise-resilience**



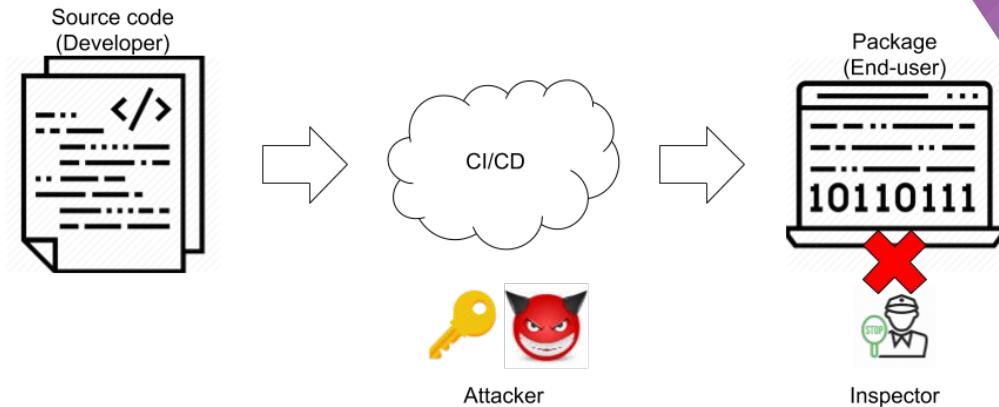
## Key idea: tamper-evident CI/CD

- **Tamper-evident**

- $x \Leftrightarrow$  source code
- $f \Leftrightarrow$  authentic CI/CD pipeline
- $y \Leftrightarrow$  package
- Does  $y = f(x)$ ?

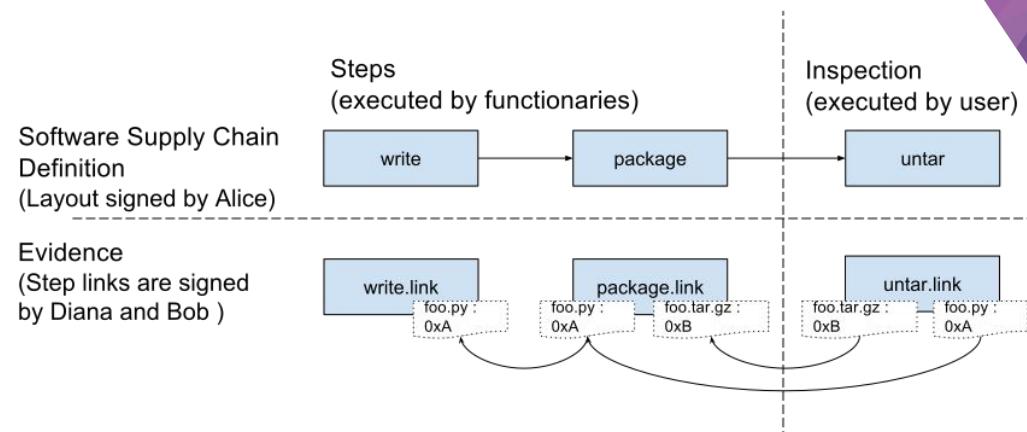
- **Compromise-resilience**

- End-users download  $x, f$ , and  $y$
- If  $y \neq f(x)$ , then reject  $y$



# in-toto: software supply chain integrity

- Pipeline = series of **steps**
  - Every step produces signed link / attestation: “I got this input, and produced that output.”
- **Inspection**
  - Verify whether each step followed pipeline
- Provides **E2E verification** of entire supply chain
- <https://in-toto.io>



# Datadog Agent integrations software supply chain

## 1. tag

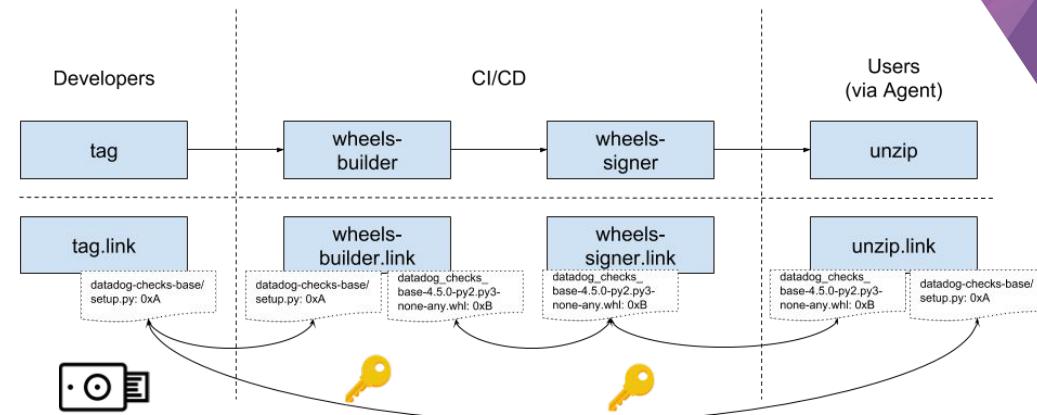
- Developer outputs source code

## 2. wheels-builder

- Container must receive same source code as in “tag”
- (Container builds wheels)
- Container outputs wheels

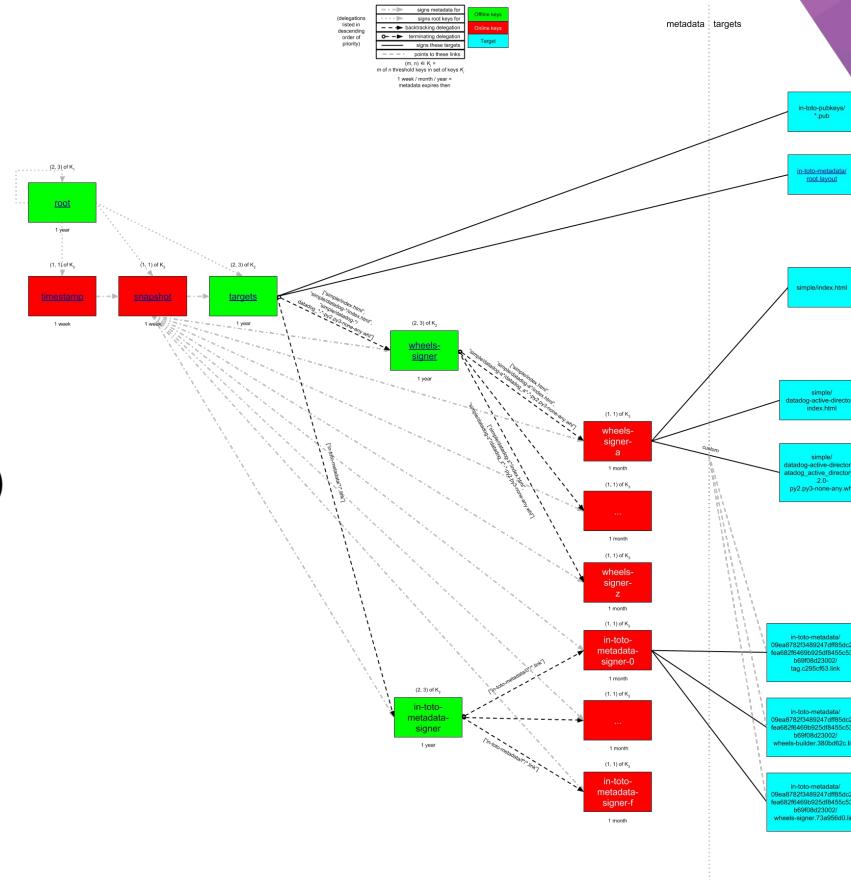
## 3. wheels-signer

- Container must receive same wheels as in “wheels-builder”



# TUF + in-toto = tamper-evident CI/CD

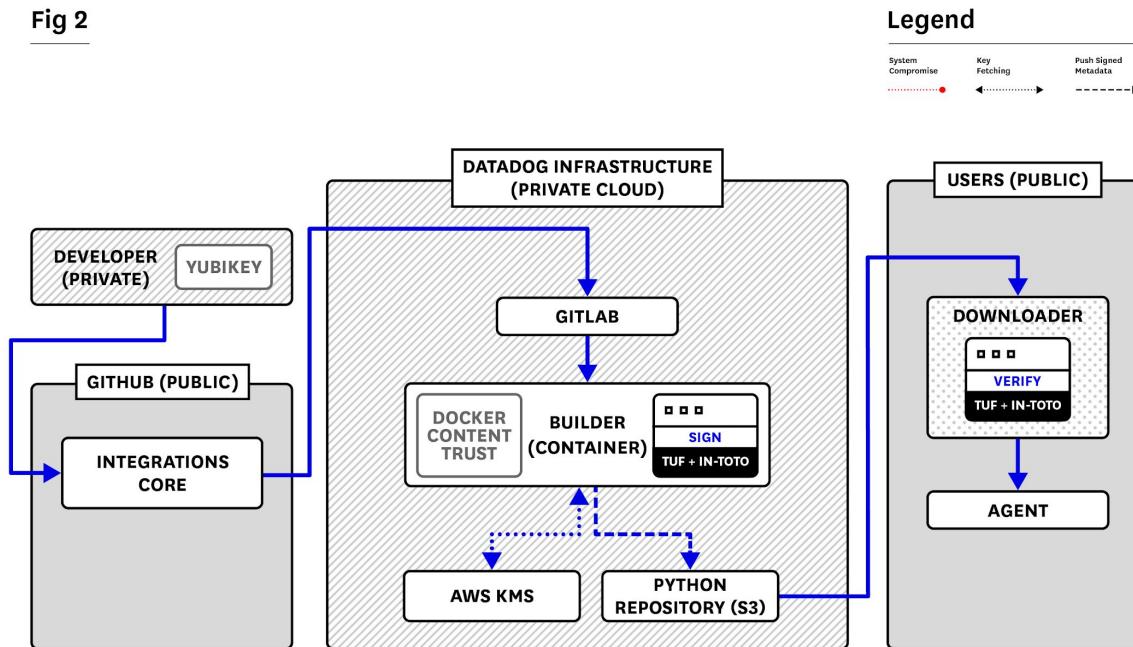
- **Offline keys (administrators)**
  - TUF root of trust
  - in-toto software supply chain
  - Public keys for in-toto software supply chain
- **Semi-offline keys (developers)**
  - Python source code
- **Online keys (CI/CD)**
  - in-toto links
  - Packages (universal Python wheels)





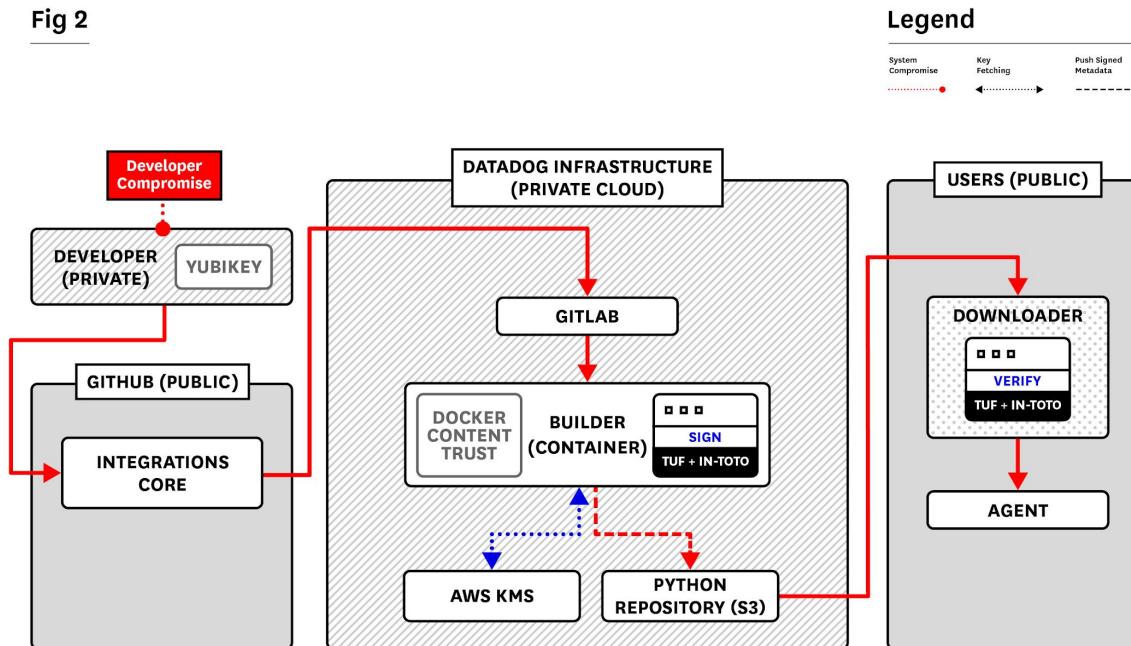
## TUF + in-toto: what can go wrong?

Fig 2



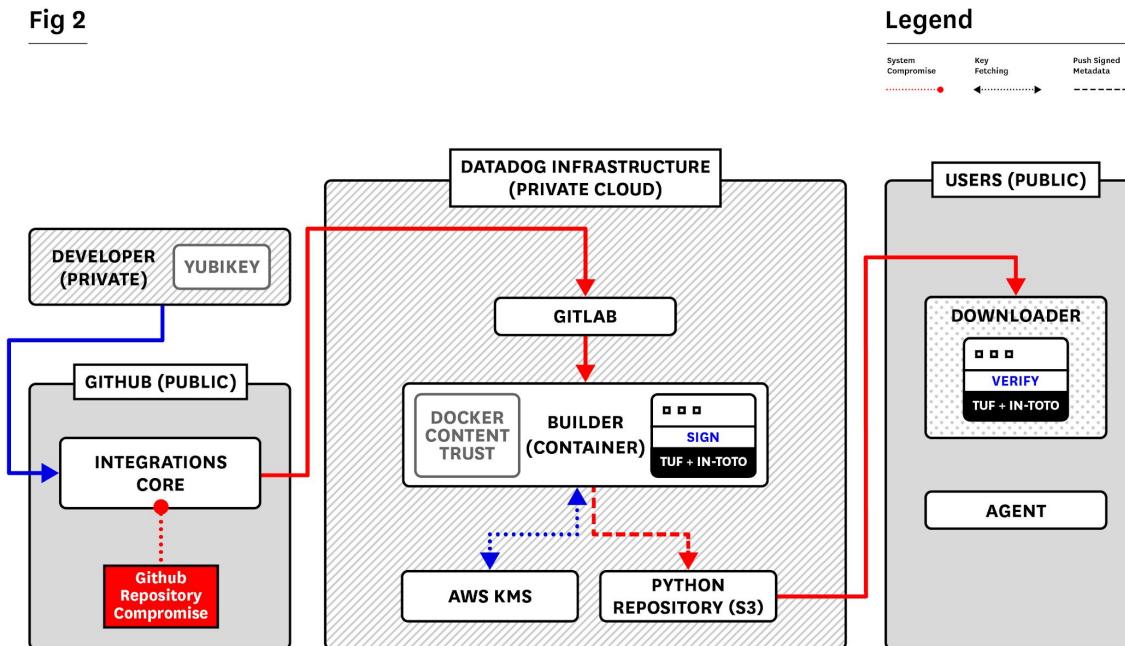
## TUF + in-toto: developer key compromise

**Fig 2**



## TUF + in-toto: VCS repository compromise

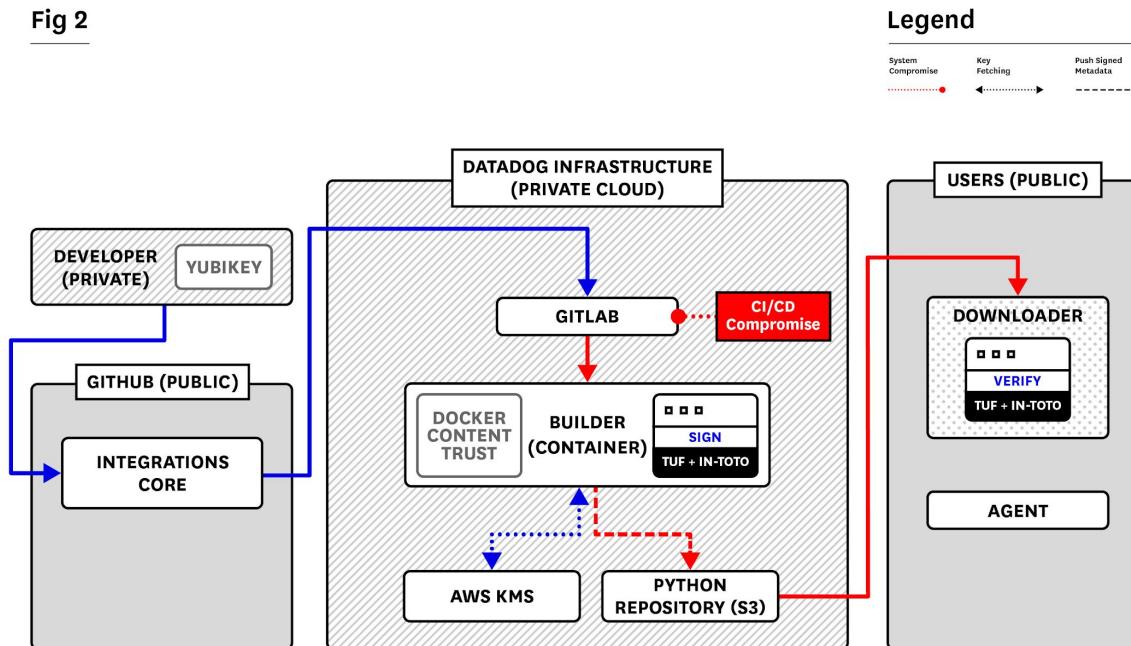
**Fig 2**





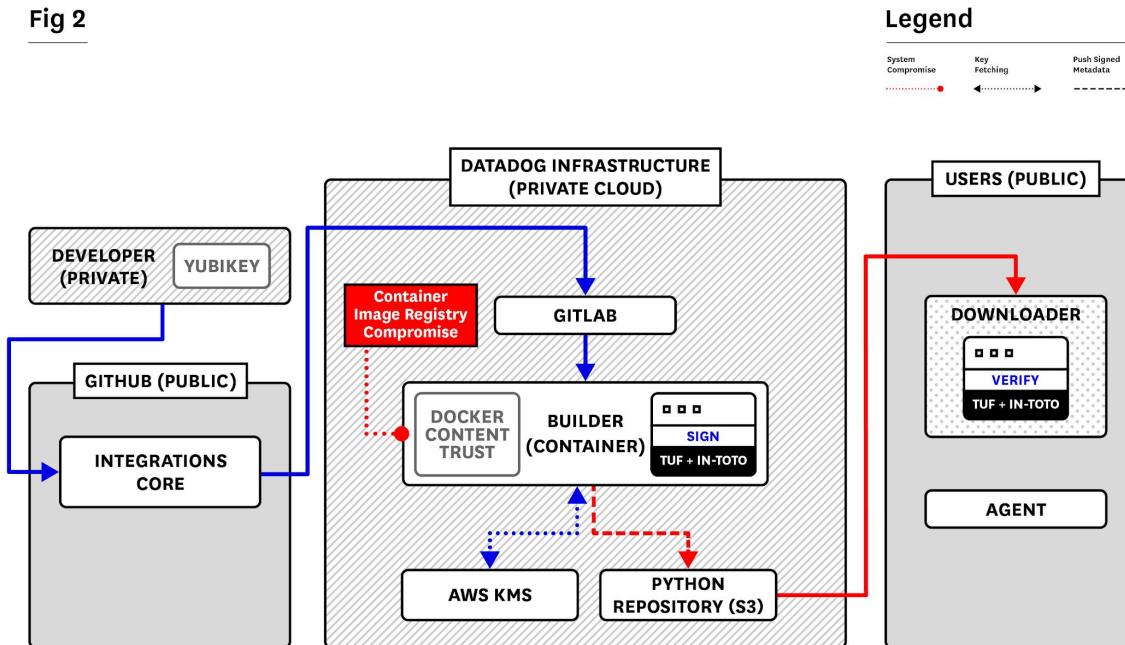
## TUF + in-toto: CI/CD system compromise

Fig 2



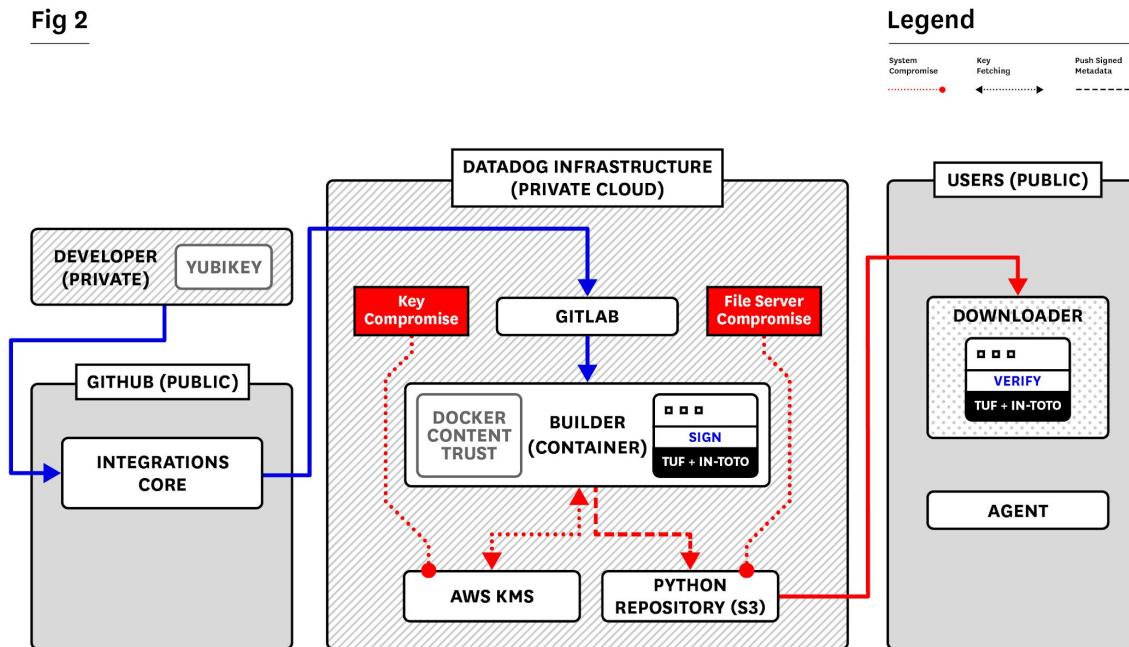
## TUF + in-toto: container image registry compromise

**Fig 2**



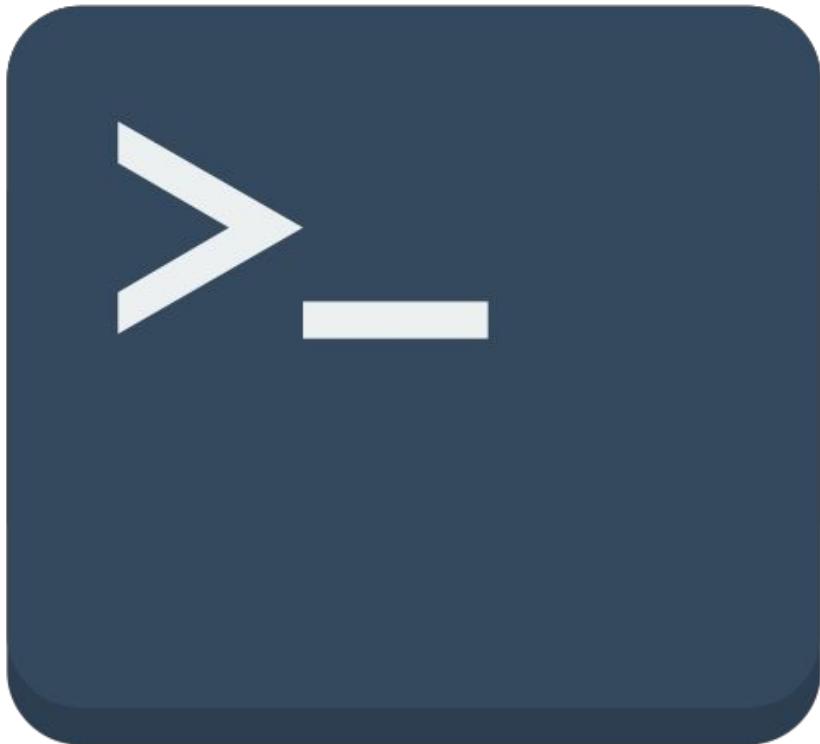
## TUF + in-toto: key + file server compromise

**Fig 2**



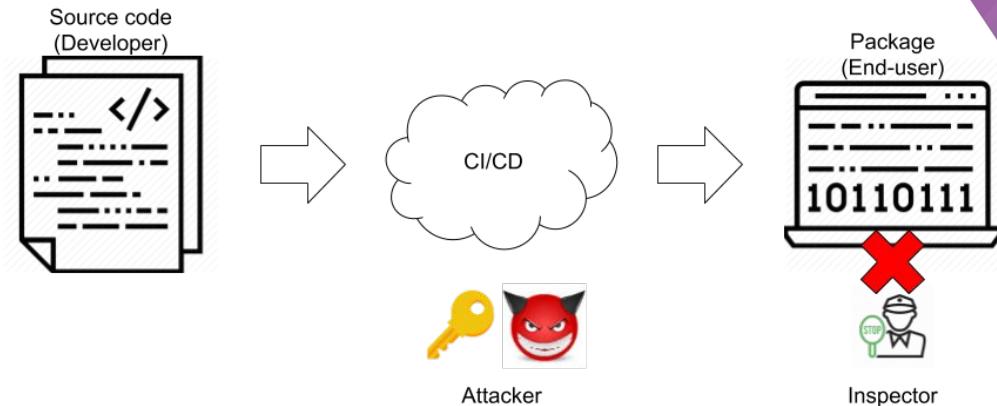


## Live demo of production



## Takeaway: TUF + in-toto = tamper-evident CI/CD

- Tamper-evident
  - $x \Leftrightarrow$  source code
  - $f \Leftrightarrow$  authentic CI/CD pipeline
  - $y \Leftrightarrow$  package
  - Does  $y = f(x)$ ?
- Compromise-resilience
  - End-users download  $x, f$ , and  $y$
  - If  $y \neq f(x)$ , then reject  $y$
- Industry-first
  - Datadog Agent 6.8.0



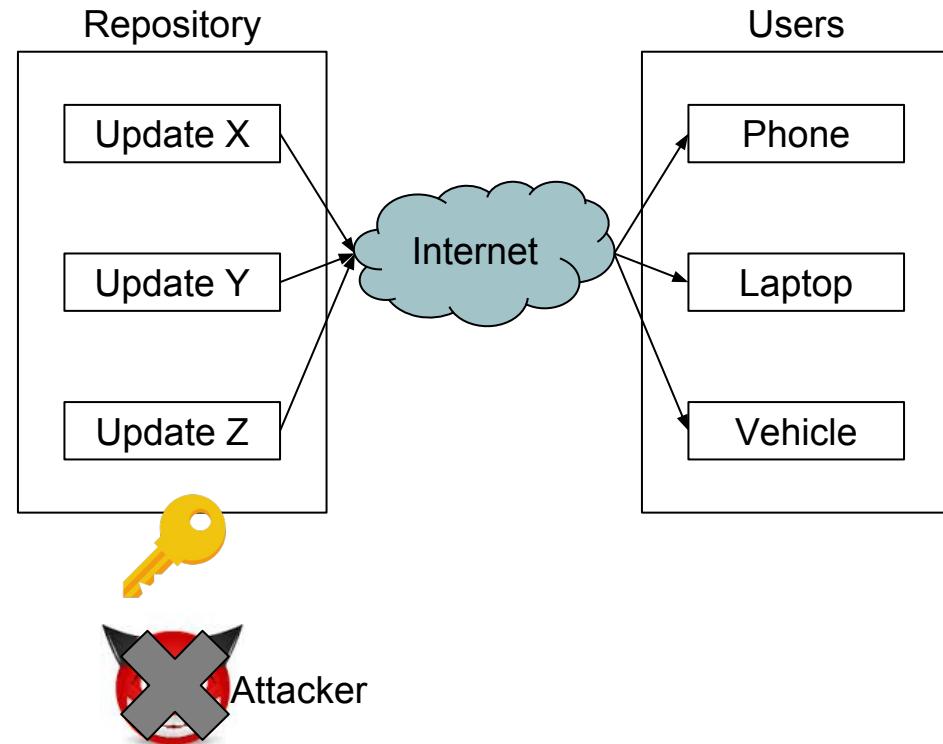


DATADOG

# Conclusions

## Takeaway: TUF = compromise-resilience

- Only question of when, not if
- Cannot prevent compromise
- But must severely limit impact
- Use TUF





## TUF: integrations & deployments



Flynn



IBM



Advanced  
Telematic  
SYSTEMS





## Acknowledgements

- **Datadog**
  - Andrew Becherer, Douglas DePerry, Agent-Integrations, Agent-Core
- **NYU**
  - Sebastien Awwad, Justin Cappos, Lois Anne DeLong, Vladimir Diaz, Lukas Puhringer, Santiago Torres-Arias
- **Docker**
  - Nathan McCauley, Diogo Monica, David Lawrence, Justin Cormack
- **CoreOS**
  - Evan Cordell, Jacob Moshenko
- **Uptane**
  - Russ Bielawski, Akan Brown, Meghan Caiazzo, Sam Lauzon, John Liming, Cameron Mott, André Weimerskirch



## Q & A

- Thanks for your time!
- Email: [trishank@datadog.com](mailto:trishank@datadog.com)
- Yubikey: <https://github.com/DataDog/yubikey>
- TUF: <https://theupdateframework.com>
- in-toto: <https://in-toto.io/>

